

chicago

May 24, 2023

1 Project Overview

The City of Chicago Vehicle Safety Board (CCVSB) interested in reducing traffic accidents and becoming aware of any interesting patterns.

2 Business Problem

The business problem is to build a classifier that can predict the primary contributory cause of car accidents in Chicago city.

3 Defineing the Questions

1. *Are there any specific locations or road segments in Chicago city that have a higher frequency of car accidents?*
2. *What are the contributing factors or characteristics associated with severe car accidents in Chicago city?*
3. *Are there any seasonal or temporal patterns in car accidents in Chicago city?*
4. *Can we build a classification model to predict the primary contributory cause of car accidents?*

Additionally, I will create a classification model to categorize accidents into two main groups for future reference:

1. Accidents caused by unintentional factors: *These accidents occur when drivers are not purposely or knowingly involved in causing the accident. They may result from factors such as driver error, environmental conditions, mechanical failures, or other unforeseen circumstances.*
2. Accidents caused by intentional factors: *These accidents involve drivers who are deliberately or knowingly involved in causing the accident. They may engage in reckless driving, aggressive behavior, or intentionally violate traffic laws, leading to the occurrence of the accident.*

By developing this classification model, we aim to distinguish between accidents that result from unintentional factors and those that involve intentional actions. This categorization will enable us to analyze and understand the different contributing factors and characteristics associated with each category, leading to targeted strategies for accident prevention and improving overall road safety.

4 Data

The dataset was from Chicago city. Their were three datasets that was obtain from Chicago Data Portal:

- * Traffic_Crashes_-_People
- * Traffic_Crashes_-_Vehicles
- * Traffic_Crashes_-_Crashes

The data provides up-to-date information as per now May 2023 from 2015. The two datasets was cleaned and merged to one.

5 Data Grocery

Index	Column Name	Description
1	SEX	Gender of the person involved in the accident
2	AGE	Age of the person involved in the accident
3	DRIVER_ACTION	Action taken by the driver before the accident
4	DRIVER_VISION	Vision condition of the driver during the accident
5	PHYSICAL_CONDITION	Physical condition of the driver at the time of the accident
6	MANEUVER	Maneuver performed by the driver during the accident
7	POSTED_SPEED_LIMIT	Speed limit posted on the road where the accident occurred
8	TRAFFIC_CONTROL_DEVICE	Control device present at the accident location
9	DEVICE_CONDITION	Condition of the traffic control device
10	WEATHER_CONDITION	Weather conditions during the accident
11	LIGHTING_CONDITION	Lighting conditions during the accident
12	TRAFFICWAY_TYPE	Type of the trafficway where the accident occurred
13	ROADWAY_SURFACE_CONDITION	Surface condition of the roadway at the accident location
14	ROAD_DEFECT	Defects present on the road where the accident occurred
15	PRIM_CONTRIBUTORY_CAUSE	Primary contributory cause of the accident
16	CRASH_HOUR	Hour of the day when the accident occurred
17	CRASH_DAY_OF_WEEK	Day of the week when the accident occurred
18	CRASH_MONTH	Month when the accident occurred
19	LATITUDE	Latitude coordinate of the accident location
20	LONGITUDE	Longitude coordinate of the accident location
21	LOCATION	Location description of the accident

6 Recording the Experimental Design

To record the experimental design for building the classifier to predict the primary contributory cause of car accidents in Chicago, I used the following steps:

- Data Collection: I gather crashes, vehicle and people accident data from <https://data.cityofchicago.org/Transportation/Traffic-Crashes-Vehicles/68nd-jvt3> and <https://data.cityofchicago.org/Transportation/Traffic-Crashes-People/u6pd-qa9d>
- Data Preprocessing: Clean the data by handling missing values, inconsistencies, and outliers.

Transform categorical variables into numerical representations suitable for machine learning algorithms. Normalize or standardize numerical features if necessary

- **Exploratory Data Analysis (EDA):** Perform exploratory analysis to understand the characteristics and distributions of variables. Identify patterns, correlations, or any interesting insights within the data. Visualize the data using plots, charts, or graphs to aid in understanding.
- **Feature Engineering:** Extract relevant features from the available data that may contribute to predicting the primary contributory cause of car accidents.
- **Target Variable Binning:** Analyze the distribution of the primary contributory cause categories. Merge or eliminate categories with very few samples to limit the number of target categories.
- **Feature Selection:** Select the most informative features that are likely to have a significant impact on the prediction.
- **Model Selection and Training:** Choose a suitable machine learning algorithm for multi-class classification, considering factors like performance, interpretability, and scalability. Split the preprocessed data into training and testing sets. Train the chosen model on the training data using appropriate algorithms and methodologies.
- **Model Evaluation:** Evaluate the trained model's performance using relevant evaluation metrics for multi-class classification, such as accuracy, precision, recall, F1-score, or confusion matrix. Perform cross-validation techniques like k-fold cross-validation to assess the model's robustness.
- **Model Optimization:** Fine-tune the model by optimizing hyperparameters to improve its performance. Use techniques like grid search, random search.
- **Predictions and Interpretation:** Use the optimized model to predict the primary contributory cause of car accidents for new instances. Analyze the predictions and interpret the results to gain insights into patterns, potential causes, or any interesting findings that can aid accident prevention efforts.
- **Reporting and Recommendations:** Summarize the findings and insights obtained from the classifier. Provide actionable recommendations for the Vehicle Safety Board or the City of Chicago based on the analysis and predictions.

7 Loading the datasets

```
[1]: # importing the libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

```
[2]: # Reading csv file
people=pd.read_csv('C:\\Users\\Admin\\Documents\\Iano\\phase 3_
↳project\\dsc-phase-3-project-v2-3\\.data\\Traffic_Crashes_-_People.csv')
```

```
vehicles=pd.read_csv('C:\\Users\\Admin\\Documents\\Iano\\phase 3_
↳project\\dsc-phase-3-project-v2-3\\.data\\Traffic_Crashes_-_Vehicles.csv')
crashes=pd.read_csv('C:\\Users\\Admin\\Documents\\Iano\\phase 3_
↳project\\dsc-phase-3-project-v2-3\\.data\\Traffic_Crashes_-_Crashes.csv')
```

C:\Users\Admin\AppData\Local\Temp\ipykernel_14600\4169662019.py:2: DtypeWarning: Columns (29) have mixed types. Specify dtype option on import or set low_memory=False.

```
people=pd.read_csv('C:\\Users\\Admin\\Documents\\Iano\\phase 3 project\\dsc-
phase-3-project-v2-3\\.data\\Traffic_Crashes_-_People.csv')
```

C:\Users\Admin\AppData\Local\Temp\ipykernel_14600\4169662019.py:3: DtypeWarning: Columns (2,19,21,40,41,42,44,48,49,50,53,55,58,59,61,71) have mixed types. Specify dtype option on import or set low_memory=False.

```
vehicles=pd.read_csv('C:\\Users\\Admin\\Documents\\Iano\\phase 3 project\\dsc-
phase-3-project-v2-3\\.data\\Traffic_Crashes_-_Vehicles.csv')
```

```
[3]: # A function to print the shape of our datasets
def print_dataset_shape(*datasets):
    """
    Prints the shape of one or more datasets (number of rows and columns).
    Assumes datasets are in a Pandas DataFrame format.
    """
    for idx, dataset in enumerate(datasets):
        print(f"Dataset {idx + 1} - Number of rows: {dataset.shape[0]}")
        print(f"Dataset {idx + 1} - Number of columns: {dataset.shape[1]}")
    # print the shape of our dataset
    print_dataset_shape(people, vehicles,crashes)
```

```
Dataset 1 - Number of rows: 1584616
Dataset 1 - Number of columns: 30
Dataset 2 - Number of rows: 1472816
Dataset 2 - Number of columns: 72
Dataset 3 - Number of rows: 722809
Dataset 3 - Number of columns: 49
```

```
[4]: # A function to get info of our dataset
#data = {
#    'vehicles': [...], # Vehicle information
#    'people': [...],   # People information
#    'crashes': [...]   # Crash information
#}

#vehicles, people, crashes = get_info(data)
```

```
[5]: # Getting the info of our data
people.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 1584616 entries, 0 to 1584615

Data columns (total 30 columns):

#	Column	Non-Null Count	Dtype
0	PERSON_ID	1584616 non-null	object
1	PERSON_TYPE	1584616 non-null	object
2	CRASH_RECORD_ID	1584616 non-null	object
3	RD_NO	1574225 non-null	object
4	VEHICLE_ID	1553626 non-null	float64
5	CRASH_DATE	1584616 non-null	object
6	SEAT_NO	320590 non-null	float64
7	CITY	1156363 non-null	object
8	STATE	1172272 non-null	object
9	ZIPCODE	1057964 non-null	object
10	SEX	1559451 non-null	object
11	AGE	1123239 non-null	float64
12	DRIVERS_LICENSE_STATE	931218 non-null	object
13	DRIVERS_LICENSE_CLASS	785449 non-null	object
14	SAFETY_EQUIPMENT	1580144 non-null	object
15	AIRBAG_DEPLOYED	1554823 non-null	object
16	EJECTION	1565274 non-null	object
17	INJURY_CLASSIFICATION	1583928 non-null	object
18	HOSPITAL	270739 non-null	object
19	EMS_AGENCY	167981 non-null	object
20	EMS_RUN_NO	27702 non-null	object
21	DRIVER_ACTION	1261166 non-null	object
22	DRIVER_VISION	1260713 non-null	object
23	PHYSICAL_CONDITION	1262045 non-null	object
24	PEDPEDAL_ACTION	29391 non-null	object
25	PEDPEDAL_VISIBILITY	29333 non-null	object
26	PEDPEDAL_LOCATION	29390 non-null	object
27	BAC_RESULT	1262189 non-null	object
28	BAC_RESULT VALUE	1859 non-null	float64
29	CELL_PHONE_USE	1158 non-null	object

dtypes: float64(4), object(26)

memory usage: 362.7+ MB

```
[6]: # Getting the info of our data
vehicles.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 1472816 entries, 0 to 1472815

Data columns (total 72 columns):

#	Column	Non-Null Count	Dtype
0	CRASH_UNIT_ID	1472816 non-null	int64
1	CRASH_RECORD_ID	1472816 non-null	object
2	RD_NO	1463413 non-null	object

3	CRASH_DATE	1472816 non-null	object
4	UNIT_NO	1472816 non-null	int64
5	UNIT_TYPE	1470808 non-null	object
6	NUM_PASSENGERS	218014 non-null	float64
7	VEHICLE_ID	1439674 non-null	float64
8	CMRC_VEH_I	27533 non-null	object
9	MAKE	1439669 non-null	object
10	MODEL	1439525 non-null	object
11	LIC_PLATE_STATE	1308626 non-null	object
12	VEHICLE_YEAR	1206500 non-null	float64
13	VEHICLE_DEFECT	1439674 non-null	object
14	VEHICLE_TYPE	1439674 non-null	object
15	VEHICLE_USE	1439674 non-null	object
16	TRAVEL_DIRECTION	1439674 non-null	object
17	MANEUVER	1439674 non-null	object
18	TOWED_I	181112 non-null	object
19	FIRE_I	1188 non-null	object
20	OCCUPANT_CNT	1439674 non-null	float64
21	EXCEED_SPEED_LIMIT_I	2397 non-null	object
22	TOWED_BY	135298 non-null	object
23	TOWED_TO	83272 non-null	object
24	AREA_00_I	51730 non-null	object
25	AREA_01_I	390692 non-null	object
26	AREA_02_I	236041 non-null	object
27	AREA_03_I	140245 non-null	object
28	AREA_04_I	141169 non-null	object
29	AREA_05_I	218874 non-null	object
30	AREA_06_I	228709 non-null	object
31	AREA_07_I	208279 non-null	object
32	AREA_08_I	216804 non-null	object
33	AREA_09_I	91248 non-null	object
34	AREA_10_I	132225 non-null	object
35	AREA_11_I	258240 non-null	object
36	AREA_12_I	253662 non-null	object
37	AREA_99_I	163988 non-null	object
38	FIRST_CONTACT_POINT	1436590 non-null	object
39	CMV_ID	15356 non-null	float64
40	USDOT_NO	8735 non-null	object
41	CCMC_NO	1894 non-null	object
42	ILCC_NO	1317 non-null	object
43	COMMERCIAL_SRC	10326 non-null	object
44	GVWR	8652 non-null	object
45	CARRIER_NAME	14683 non-null	object
46	CARRIER_STATE	13796 non-null	object
47	CARRIER_CITY	13546 non-null	object
48	HAZMAT_PLACARDS_I	301 non-null	object
49	HAZMAT_NAME	56 non-null	object
50	UN_NO	522 non-null	object

51	HAZMAT_PRESENT_I	11210 non-null	object
52	HAZMAT_REPORT_I	10886 non-null	object
53	HAZMAT_REPORT_NO	1 non-null	object
54	MCS_REPORT_I	10934 non-null	object
55	MCS_REPORT_NO	7 non-null	object
56	HAZMAT_VIO_CAUSE_CRASH_I	11051 non-null	object
57	MCS_VIO_CAUSE_CRASH_I	10844 non-null	object
58	IDOT_PERMIT_NO	850 non-null	object
59	WIDE_LOAD_I	128 non-null	object
60	TRAILER1_WIDTH	2735 non-null	object
61	TRAILER2_WIDTH	324 non-null	object
62	TRAILER1_LENGTH	2222 non-null	float64
63	TRAILER2_LENGTH	65 non-null	float64
64	TOTAL_VEHICLE_LENGTH	2699 non-null	float64
65	AXLE_CNT	4008 non-null	float64
66	VEHICLE_CONFIG	12720 non-null	object
67	CARGO_BODY_TYPE	12154 non-null	object
68	LOAD_TYPE	11624 non-null	object
69	HAZMAT_OUT_OF_SERVICE_I	10560 non-null	object
70	MCS_OUT_OF_SERVICE_I	10806 non-null	object
71	HAZMAT_CLASS	1018 non-null	object

dtypes: float64(9), int64(2), object(61)
memory usage: 809.0+ MB

```
[7]: # Getting the info of our data
      crashes.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 722809 entries, 0 to 722808
Data columns (total 49 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CRASH_RECORD_ID                       722809 non-null object
1   RD_NO                                 718253 non-null object
2   CRASH_DATE_EST_I                      54664 non-null  object
3   CRASH_DATE                            722809 non-null object
4   POSTED_SPEED_LIMIT                    722809 non-null int64
5   TRAFFIC_CONTROL_DEVICE                 722809 non-null object
6   DEVICE_CONDITION                       722809 non-null object
7   WEATHER_CONDITION                     722809 non-null object
8   LIGHTING_CONDITION                    722809 non-null object
9   FIRST_CRASH_TYPE                      722809 non-null object
10  TRAFFICWAY_TYPE                        722809 non-null object
11  LANE_CNT                               199002 non-null float64
12  ALIGNMENT                             722809 non-null object
13  ROADWAY_SURFACE_COND                   722809 non-null object
14  ROAD_DEFECT                            722809 non-null object
15  REPORT_TYPE                            702521 non-null object
```

16	CRASH_TYPE	722809	non-null	object
17	INTERSECTION_RELATED_I	165794	non-null	object
18	NOT_RIGHT_OF_WAY_I	33751	non-null	object
19	HIT_AND_RUN_I	225032	non-null	object
20	DAMAGE	722809	non-null	object
21	DATE_POLICE_NOTIFIED	722809	non-null	object
22	PRIM_CONTRIBUTORY_CAUSE	722809	non-null	object
23	SEC_CONTRIBUTORY_CAUSE	722809	non-null	object
24	STREET_NO	722809	non-null	int64
25	STREET_DIRECTION	722805	non-null	object
26	STREET_NAME	722808	non-null	object
27	BEAT_OF_OCCURRENCE	722804	non-null	float64
28	PHOTOS_TAKEN_I	9081	non-null	object
29	STATEMENTS_TAKEN_I	15398	non-null	object
30	DOORING_I	2186	non-null	object
31	WORK_ZONE_I	4209	non-null	object
32	WORK_ZONE_TYPE	3290	non-null	object
33	WORKERS_PRESENT_I	1085	non-null	object
34	NUM_UNITS	722809	non-null	int64
35	MOST_SEVERE_INJURY	721233	non-null	object
36	INJURIES_TOTAL	721244	non-null	float64
37	INJURIES_FATAL	721244	non-null	float64
38	INJURIES_INCAPACITATING	721244	non-null	float64
39	INJURIES_NON_INCAPACITATING	721244	non-null	float64
40	INJURIES_REPORTED_NOT_EVIDENT	721244	non-null	float64
41	INJURIES_NO_INDICATION	721244	non-null	float64
42	INJURIES_UNKNOWN	721244	non-null	float64
43	CRASH_HOUR	722809	non-null	int64
44	CRASH_DAY_OF_WEEK	722809	non-null	int64
45	CRASH_MONTH	722809	non-null	int64
46	LATITUDE	718130	non-null	float64
47	LONGITUDE	718130	non-null	float64
48	LOCATION	718130	non-null	object

dtypes: float64(11), int64(6), object(32)

memory usage: 270.2+ MB

```
[8]: # Function to display the head of our datasets
def display_data_head(people, vehicles, crashes):
    dfs = [people.head(), vehicles.head(), crashes.head()]
    df_names = ["people", "vehicles", "crashes"]
    for df, name in zip(dfs, df_names):
        print(f"\n{name}:\n")
        display(df)
    # Display the head of our datasets
    display_data_head(people, vehicles, crashes)
```

people:

	PERSON_ID	PERSON_TYPE	CRASH_RECORD_ID \
0	01577624	DRIVER	e8d0a18503a3ef7a69ee631eacffd421ea154ea9782131...
1	01577610	DRIVER	0690865a402d40a7eab391f94a658b48dc03abb636a03a...
2	01577611	DRIVER	0690865a402d40a7eab391f94a658b48dc03abb636a03a...
3	01577604	DRIVER	fe7d5f687f472c631a7a3516d0047a3cf7a8ab2cb0b6a4...
4	01577605	DRIVER	fe7d5f687f472c631a7a3516d0047a3cf7a8ab2cb0b6a4...

	RD_NO	VEHICLE_ID	CRASH_DATE	SEAT_NO	CITY	STATE \
0	NaN	1500926.0	05/17/2023 10:20:00 AM	NaN	CHICAGO	IL
1	NaN	1500906.0	05/17/2023 10:10:00 AM	NaN	NaN	NaN
2	NaN	1500913.0	05/17/2023 10:10:00 AM	NaN	VALPARAISO	IN
3	NaN	1500903.0	05/17/2023 10:05:00 AM	NaN	CHICAGO	IL
4	NaN	1500907.0	05/17/2023 10:05:00 AM	NaN	CHICAGO	IL

	ZIPCODE	...	EMS_RUN_NO	DRIVER_ACTION	DRIVER_VISION	PHYSICAL_CONDITION \
0	60637	...	NaN	IMPROPER PARKING	UNKNOWN	UNKNOWN
1	NaN	...	NaN	UNKNOWN	UNKNOWN	UNKNOWN
2	46385	...	NaN	NONE NOT OBSCURED		NORMAL
3	60613	...	NaN	UNKNOWN	UNKNOWN	NORMAL
4	60660	...	NaN	UNKNOWN	UNKNOWN	NORMAL

	PEDPEDAL_ACTION	PEDPEDAL_VISIBILITY	PEDPEDAL_LOCATION	BAC_RESULT \
0	NaN	NaN	NaN	TEST NOT OFFERED
1	NaN	NaN	NaN	TEST NOT OFFERED
2	NaN	NaN	NaN	TEST NOT OFFERED
3	NaN	NaN	NaN	TEST NOT OFFERED
4	NaN	NaN	NaN	TEST NOT OFFERED

	BAC_RESULT	VALUE	CELL_PHONE_USE
0		NaN	NaN
1		NaN	NaN
2		NaN	NaN
3		NaN	NaN
4		NaN	NaN

[5 rows x 30 columns]

vehicles:

	CRASH_UNIT_ID	CRASH_RECORD_ID	RD_NO \
0	1577434	25d92973475a04a93e7fd206fbfce57e8a9a1e25cc85a7...	NaN
1	1577435	25d92973475a04a93e7fd206fbfce57e8a9a1e25cc85a7...	NaN
2	1577450	375ac7f6fcb4ef73d728edc52ed556f23fd465a351833f...	NaN
3	1577451	375ac7f6fcb4ef73d728edc52ed556f23fd465a351833f...	NaN
4	1577452	375ac7f6fcb4ef73d728edc52ed556f23fd465a351833f...	NaN

		CRASH_DATE	UNIT_NO	UNIT_TYPE	NUM_PASSENGERS	VEHICLE_ID	\
0	05/16/2023	11:12:00 PM	1	DRIVER	NaN	1500741.0	
1	05/16/2023	11:12:00 PM	2	DRIVER	1.0	1500742.0	
2	05/16/2023	11:06:00 PM	1	DRIVER	NaN	1500759.0	
3	05/16/2023	11:06:00 PM	2	DRIVER	NaN	1500760.0	
4	05/16/2023	11:06:00 PM	3	DRIVER	NaN	1500761.0	

	CMRC_VEH_I	MAKE	...	TRAILER1_LENGTH	TRAILER2_LENGTH	\
0	NaN	HONDA	...	NaN	NaN	
1	NaN	CHEVROLET	...	NaN	NaN	
2	NaN	DODGE	...	NaN	NaN	
3	NaN	TOYOTA	...	NaN	NaN	
4	NaN	FORD	...	NaN	NaN	

	TOTAL_VEHICLE_LENGTH	AXLE_CNT	VEHICLE_CONFIG	CARGO_BODY_TYPE	LOAD_TYPE	\
0	NaN	NaN	NaN	NaN	NaN	
1	NaN	NaN	NaN	NaN	NaN	
2	NaN	NaN	NaN	NaN	NaN	
3	NaN	NaN	NaN	NaN	NaN	
4	NaN	NaN	NaN	NaN	NaN	

	HAZMAT_OUT_OF_SERVICE_I	MCS_OUT_OF_SERVICE_I	HAZMAT_CLASS
0	NaN		NaN
1	NaN		NaN
2	NaN		NaN
3	NaN		NaN
4	NaN		NaN

[5 rows x 72 columns]

crashes:

	CRASH_RECORD_ID	RD_NO	CRASH_DATE_EST_I	\
0	25d92973475a04a93e7fd206fbfce57e8a9a1e25cc85a7...	NaN	NaN	
1	375ac7f6fcb4ef73d728edc52ed556f23fd465a351833f...	NaN	NaN	
2	246fea010af2010860046c6ef36efb75a8c60244088939...	NaN	NaN	
3	18c220f7eeceb2cf6f9512c9b83382da28d8565fbbaaec...	NaN	NaN	
4	cfecdce601503162eb09337bd6051ea358dca7294d440b...	NaN	NaN	

	CRASH_DATE	POSTED_SPEED_LIMIT	TRAFFIC_CONTROL_DEVICE	\
0	05/16/2023 11:12:00 PM	30	TRAFFIC SIGNAL	
1	05/16/2023 11:06:00 PM	30	NO CONTROLS	
2	05/16/2023 11:05:00 PM	30	NO CONTROLS	
3	05/16/2023 10:20:00 PM	25	NO CONTROLS	
4	05/16/2023 09:45:00 PM	30	UNKNOWN	

	DEVICE_CONDITION	WEATHER_CONDITION	LIGHTING_CONDITION	\
0	FUNCTIONING PROPERLY	CLEAR	DARKNESS, LIGHTED ROAD	
1	NO CONTROLS	CLEAR	DARKNESS, LIGHTED ROAD	
2	NO CONTROLS	CLEAR	DARKNESS, LIGHTED ROAD	
3	NO CONTROLS	CLEAR	DARKNESS	
4	FUNCTIONING PROPERLY	CLEAR	DARKNESS	

	FIRST_CRASH_TYPE	... INJURIES_NON_INCAPACITATING	\
0	REAR END	...	0.0
1	REAR TO FRONT	...	0.0
2	PARKED MOTOR VEHICLE	...	0.0
3	PEDALCYCLIST	...	1.0
4	REAR END	...	0.0

	INJURIES_REPORTED_NOT_EVIDENT	INJURIES_NO_INDICATION	INJURIES_UNKNOWN	\
0	0.0	3.0	0.0	
1	0.0	3.0	0.0	
2	0.0	1.0	0.0	
3	0.0	1.0	0.0	
4	0.0	2.0	0.0	

	CRASH_HOUR	CRASH_DAY_OF_WEEK	CRASH_MONTH	LATITUDE	LONGITUDE	\
0	23	3	5	41.952691	-87.807413	
1	23	3	5	41.997837	-87.688814	
2	23	3	5	42.002331	-87.695032	
3	22	3	5	41.827340	-87.636475	
4	21	3	5	41.808853	-87.640097	

	LOCATION
0	POINT (-87.807413247555 41.952691362649)
1	POINT (-87.688813887189 41.997837266972)
2	POINT (-87.695032165757 42.002331485776)
3	POINT (-87.636475000374 41.827339537397)
4	POINT (-87.640097485203 41.808853153697)

[5 rows x 49 columns]

8 Data cleaning

8.0.1 Cheaking for duplicates

```
[9]: # A function to check for duplicates in our datasets
def check_duplicates(df):
    """
    This function checks for and returns any duplicates in a given dataframe.
    """
    duplicates = df[df.duplicated()]
```

```

if duplicates.shape[0] == 0:
    print("No duplicates found in the dataset")
else:
    print("Duplicates found in the dataset:")
    return duplicates
# Calling for the function to check for duplicates
check_duplicates(people)
check_duplicates(vehicles)
check_duplicates(crashes)

```

No duplicates found in the dataset

No duplicates found in the dataset

No duplicates found in the dataset

The data had no duplicates

Dropping columns that are not relevant

```

[10]: #dropping columns either not relevant
people_drop = people[['PERSON_ID', 'RD_NO', 'CRASH_DATE', 'SEAT_NO', 'CITY',
↳ 'STATE', 'ZIPCODE',
    'DRIVERS_LICENSE_STATE', 'DRIVERS_LICENSE_CLASS', 'SAFETY_EQUIPMENT',
↳ 'AIRBAG_DEPLOYED', 'EJECTION',
    'INJURY_CLASSIFICATION', 'HOSPITAL', 'EMS_AGENCY',
↳ 'EMS_RUN_NO', 'PEDPEDAL_LOCATION']]

```

```

[11]: #dropping columns either not relevant
vehicles_drop = vehicles[['CRASH_UNIT_ID', 'RD_NO', 'CRASH_DATE',
    'NUM_PASSENGERS', 'MAKE', 'MODEL', 'VEHICLE_YEAR', 'CMRC_VEH_I',
    'LIC_PLATE_STATE', 'TOWED_I', 'FIRE_I', 'OCCUPANT_CNT', 'TOWED_BY',
    'TOWED_TO', 'AREA_00_I', 'AREA_01_I', 'AREA_02_I', 'AREA_03_I',
    'AREA_04_I', 'AREA_05_I', 'AREA_06_I', 'AREA_07_I', 'AREA_08_I',
    'AREA_09_I', 'AREA_10_I', 'AREA_11_I', 'AREA_12_I', 'AREA_99_I',
    'CMV_ID', 'USDOT_NO', 'CCMC_NO', 'ILCC_NO',
    'COMMERCIAL_SRC', 'GVWR', 'CARRIER_NAME', 'CARRIER_STATE',
    'CARRIER_CITY', 'HAZMAT_PLACARDS_I', 'HAZMAT_NAME', 'UN_NO',
    'HAZMAT_PRESENT_I', 'HAZMAT_REPORT_I', 'HAZMAT_REPORT_NO',
    'MCS_REPORT_I', 'MCS_REPORT_NO', 'HAZMAT_VIO_CAUSE_CRASH_I',
    'MCS_VIO_CAUSE_CRASH_I', 'IDOT_PERMIT_NO', 'WIDE_LOAD_I',
    'TRAILER1_WIDTH', 'TRAILER2_WIDTH', 'TRAILER1_LENGTH',
    'TRAILER2_LENGTH', 'TOTAL_VEHICLE_LENGTH', 'AXLE_CNT', 'VEHICLE_CONFIG',
    'CARGO_BODY_TYPE', 'LOAD_TYPE', 'HAZMAT_OUT_OF_SERVICE_I',
    'MCS_OUT_OF_SERVICE_I', 'HAZMAT_CLASS']]

```

```

[12]: #dropping columns either not relevant
crashes_drop = crashes[['RD_NO', 'CRASH_DATE_EST_I',
↳ 'CRASH_DATE', 'LANE_CNT', 'REPORT_TYPE', 'DATE_POLICE_NOTIFIED',

```

```
'SEC_CONTRIBUTORY_CAUSE', 'STREET_NO', 'STREET_DIRECTION', □
↳ 'STREET_NAME', 'PHOTOS_TAKEN_I', 'STATEMENTS_TAKEN_I',
    'BEAT_OF_OCCURRENCE', 'MOST_SEVERE_INJURY', □
↳ 'INJURIES_TOTAL', 'INJURIES_FATAL', 'INJURIES_INCAPACITATING',
    'INJURIES_NON_INCAPACITATING', □
↳ 'INJURIES_REPORTED_NOT_EVIDENT', 'INJURIES_NO_INDICATION',
    'INJURIES_UNKNOWN']]
```

```
[13]: #dropping columns that are not relevant
people.drop(columns=people_drop, inplace=True)
vehicles.drop(columns=vehicles_drop, inplace=True)
crashes.drop(columns=crashes_drop, inplace=True)
```

8.0.2 Cheaking missing values

```
[14]: # A function to check for missing values in our dataset
def check_missing_values(data):
    # Count missing values in each column
    missing_values = data.isnull().sum()

    # Convert missing values count to percentage of total rows
    missing_percent = ((missing_values / len(data)) * 100).
    ↳ sort_values(ascending=True)

    # Combine the missing values count and percent into a DataFrame
    missing_df = pd.concat([missing_values, missing_percent], axis=1)
    missing_df.columns = ['Missing Values', '% of Total']

    # Return only columns with missing values
    missing_df = missing_df[missing_df['Missing Values'] > 0]

    return missing_df

# Check missing values in each dataset
display(check_missing_values(people))
display(check_missing_values(vehicles))
display(check_missing_values(crashes))
```

	Missing Values	% of Total
VEHICLE_ID	30990	1.955679
SEX	25165	1.588082
AGE	461377	29.116013
DRIVER_ACTION	323450	20.411885
DRIVER_VISION	323903	20.440473
PHYSICAL_CONDITION	322571	20.356414
PEDPEDAL_ACTION	1555225	98.145229
PEDPEDAL_VISIBILITY	1555283	98.148889

BAC_RESULT	322427	20.347327
BAC_RESULT VALUE	1582757	99.882685
CELL_PHONE_USE	1583458	99.926922

	Missing Values	% of Total
UNIT_TYPE	2008	0.136337
VEHICLE_ID	33142	2.250247
VEHICLE_DEFECT	33142	2.250247
VEHICLE_TYPE	33142	2.250247
VEHICLE_USE	33142	2.250247
TRAVEL_DIRECTION	33142	2.250247
MANEUVER	33142	2.250247
EXCEED_SPEED_LIMIT_I	1470419	99.837251
FIRST_CONTACT_POINT	36226	2.459642

	Missing Values	% of Total
INTERSECTION_RELATED_I	557015	77.062543
NOT_RIGHT_OF_WAY_I	689058	95.330578
HIT_AND_RUN_I	497777	68.867017
DOORING_I	720623	99.697569
WORK_ZONE_I	718600	99.417688
WORK_ZONE_TYPE	719519	99.544831
WORKERS_PRESENT_I	721724	99.849891
LATITUDE	4679	0.647336
LONGITUDE	4679	0.647336
LOCATION	4679	0.647336

- In people data majority values are missing ie: *CELL_PHONE_USE*, *BAC_RESULT VALUE*, *PEDPEDAL_VISIBILITY* and *PEDPEDAL_ACTION* I will remove them
- In vehicles data majority values are missing in the *EXCEED_SPEED_LIMIT_I* so I will remove it and drop the remaining since it has about 2% of the data.
- In crashes data majority of the data is missing ie: *INTERSECTION_RELATED_I*, *NOT_RIGHT_OF_WAY_I*, *HIT_AND_RUN_I*, *DOORING_I*, *WORK_ZONE_I*, *WORK_ZONE_TYPE* and *WORKERS_PRESENT_I* so I will remove them and the remaining I will drop them

```
[15]: # Removing columns that have large amounts of missing values
people.drop(['PEDPEDAL_ACTION', 'PEDPEDAL_VISIBILITY', 'BAC_RESULT VALUE',
↳ 'CELL_PHONE_USE'], axis=1, inplace=True)
vehicles.drop('EXCEED_SPEED_LIMIT_I', axis=1, inplace=True)
crashes.drop(labels=['INTERSECTION_RELATED_I',
↳ 'NOT_RIGHT_OF_WAY_I', 'HIT_AND_RUN_I',
↳ 'DOORING_I', 'WORK_ZONE_I', 'WORK_ZONE_TYPE',
↳ 'WORKERS_PRESENT_I'], axis=1, inplace=True)
```

```
[16]: # Dropping the missing values
vehicles.dropna(axis=0, inplace=True)
crashes.dropna(axis=0, inplace=True)
```

Rechecking missing values have been drop

```
[17]: # A function to check for missing values in our dataset
def check_missing_values(data):
    # Count missing values in each column
    missing_values = data.isnull().sum()

    # Convert missing values count to percentage of total rows
    missing_percent = ((missing_values / len(data)) * 100).
    ↪sort_values(ascending=True)

    # Combine the missing values count and percent into a DataFrame
    missing_df = pd.concat([missing_values, missing_percent], axis=1)
    missing_df.columns = ['Missing Values', '% of Total']

    # Return only columns with missing values
    missing_df = missing_df[missing_df['Missing Values'] > 0]

    return missing_df

# Check missing values in each dataset
display(check_missing_values(people))
display(check_missing_values(vehicles))
display(check_missing_values(crashes))
```

	Missing Values	% of Total
VEHICLE_ID	30990	1.955679
SEX	25165	1.588082
AGE	461377	29.116013
DRIVER_ACTION	323450	20.411885
DRIVER_VISION	323903	20.440473
PHYSICAL_CONDITION	322571	20.356414
BAC_RESULT	322427	20.347327

Empty DataFrame

Columns: [Missing Values, % of Total]

Index: []

Empty DataFrame

Columns: [Missing Values, % of Total]

Index: []

In the vehicles and crashes data do not have missing values in people it has missing values so I will check for the unique categories in each of our remainder variables.

Checking for unique variables

1. people

```
[18]: #provide an overview of the unique values and their frequencies for each column
      ↪ in the DataFrame
for col in people.columns:
    print('\n' + col + '\n')
    print(people[col].value_counts())
```

PERSON_TYPE

```
DRIVER          1233524
PASSENGER       320590
PEDESTRIAN      18190
BICYCLE         10750
NON-MOTOR VEHICLE 1286
NON-CONTACT VEHICLE 276
Name: PERSON_TYPE, dtype: int64
```

CRASH_RECORD_ID

```
31ecf6862c691ff12d3856213b902c146b07337b42a5692e3a176a66d684d221028bb5118ef6d67a
313bcaed9e97bee1855cb1f5e8650f49e8dc17663475a1ee    61
13026c7fb51566d9ca487a093e38c6f5621c2ec25be48c306b6574983b61daeee589524b96bb2bfe
66ddd0f695c8d2bf3ab0297558528e9c7a70363c763d6bd1    50
1829f52c1281a0396ef94692331b3dc530bc4be5a54cd55e94c24a5e5e49b800fbcf9f24dabe4c82
77c8964ad05aad9c89e90fd94021959d6dff5fad55480d595    46
5fd56a31e9c4608adcd9f1d504236f856a72905451941d850fb4ddf1464a44bddfcfac7ed04fee9f
fa4855bfbf07042568fd9033c3f2e48f398f7eb0002a09ab    45
c727dc759107cf17b2e8141149347128bb4bc26b026c7805562206c7c5761c543dd7cc0e47fc1137
9455a2ecbb2847c3d1744d6feb78f276d9a457e9beeb6121    45
..
6e7087f80201ed819e60a0120bba83568f85702fd7dc71972adc264b8e157df6efdf87053e8846eb
2ca927aad1807cb7c56ea09c48bb482dd67bd303bff8a9be    1
246dd3f15c0813afa35029906e027e592f6997a08675f6077194b7ca5c2e90895d8637cabe3bca45
ad12294c463d9243ceeee4a864a8f41e0adb8c1756d074fb    1
1f662ae11036a783afd1a3c5d0a28be233895aaeffd416b4864d23a424fe22d19cb23f8da753f473
f2f150da63960b49dbb3ff5d45f40c78cb59edc0ec8d6e5c    1
70e0f8650d1d08f26f86da60da62deaa040fd5869ab42d0d98ccdface0f1bf093286a29ae73f3c9f
36c7fe970d7922d0c8a36bf4f193af1d82d76ef49f8653ff    1
d0906381565cf7d51d5c0537fc254d957c65328c48367d97f5bfc4e0dd8803d498ce121d4549e437
496276539d09637ab68c86ea530ac75ac4d200a2e32541a1    1
Name: CRASH_RECORD_ID, Length: 721315, dtype: int64
```

VEHICLE_ID

```
332155.0    60
643997.0    47
1481124.0    44
```


366311.0	44
162199.0	44
	..
949494.0	1
950309.0	1
949525.0	1
949550.0	1
739508.0	1

Name: VEHICLE_ID, Length: 1248684, dtype: int64

SEX

M	819613
F	596850
X	142988

Name: SEX, dtype: int64

AGE

25.0	31802
27.0	31709
26.0	31684
28.0	31007
24.0	30716
	...
-49.0	1
-177.0	1
-47.0	1
-40.0	1
106.0	1

Name: AGE, Length: 116, dtype: int64

DRIVER_ACTION

NONE	453928
UNKNOWN	313404
FAILED TO YIELD	114746
OTHER	111805
FOLLOWED TOO CLOSELY	76517
IMPROPER BACKING	38166
IMPROPER TURN	33039
IMPROPER LANE CHANGE	32513
IMPROPER PASSING	28037
DISREGARDED CONTROL DEVICES	21954
TOO FAST FOR CONDITIONS	19502
WRONG WAY/SIDE	4978
IMPROPER PARKING	4706
OVERCORRECTED	2088

EVADING POLICE VEHICLE	2000
CELL PHONE USE OTHER THAN TEXTING	1913
EMERGENCY VEHICLE ON CALL	1149
TEXTING	518
STOPPED SCHOOL BUS	150
LICENSE RESTRICTIONS	53

Name: DRIVER_ACTION, dtype: int64

DRIVER_VISION

NOT OBSCURED	651081
UNKNOWN	578481
OTHER	12766
MOVING VEHICLES	7373
PARKED VEHICLES	4531
WINDSHIELD (WATER/ICE)	3555
BLINDED - SUNLIGHT	1508
TREES, PLANTS	539
BUILDINGS	458
BLINDED - HEADLIGHTS	127
HILLCREST	94
BLOWING MATERIALS	89
EMBANKMENT	78
SIGNBOARD	33

Name: DRIVER_VISION, dtype: int64

PHYSICAL_CONDITION

NORMAL	828169
UNKNOWN	410404
IMPAIRED - ALCOHOL	5620
REMOVED BY EMS	4657
OTHER	3698
FATIGUED/ASLEEP	3368
EMOTIONAL	2826
ILLNESS/FAINTED	1183
HAD BEEN DRINKING	953
IMPAIRED - DRUGS	664
IMPAIRED - ALCOHOL AND DRUGS	341
MEDICATED	162

Name: PHYSICAL_CONDITION, dtype: int64

BAC_RESULT

TEST NOT OFFERED	1243396
TEST REFUSED	13377
TEST PERFORMED, RESULTS UNKNOWN	3100
TEST TAKEN	2316

Name: BAC_RESULT, dtype: int64

Variables 'Unable to determine' and 'Not applicable' does not add any value to our project. We will remove these.

- We are only interested in the DRIVER for PERSON_TYPE. We will slice the rest out of the dataset.
- 'SEX - X' and negative 'AGE' values does not tell much about the person. We will slice them out of the dataset.
- DRIVER_ACTION seems more like a target variable rather than a predictor. We will drop this column.
- BAC_RESULT has mostly TEST NOT OFFERED. the rest of the data will not be able to add much weight for this predictor. We will drop this as well.

```
[19]: people.drop('BAC_RESULT', axis=1, inplace=True)
```

```
[20]: #slice the PERSON_TYPE, AGE and SEX columns.
people_fin = people[(people['PERSON_TYPE'] == 'DRIVER') & (people['SEX'] != 'X') & (people['AGE'] > 0)].copy()
```

```
[21]: people_fin.head()
```

```
[21]:
```

	PERSON_TYPE	CRASH_RECORD_ID	VEHICLE_ID	\
0	DRIVER	e8d0a18503a3ef7a69ee631eacffd421ea154ea9782131...	1500926.0	
2	DRIVER	0690865a402d40a7eab391f94a658b48dc03abb636a03a...	1500913.0	
3	DRIVER	fe7d5f687f472c631a7a3516d0047a3cf7a8ab2cb0b6a4...	1500903.0	
4	DRIVER	fe7d5f687f472c631a7a3516d0047a3cf7a8ab2cb0b6a4...	1500907.0	
7	DRIVER	439723dae3207b29a5fac90ce3efcb21cde0fe63e49350...	1500872.0	

	SEX	AGE	DRIVER_ACTION	DRIVER_VISION	PHYSICAL_CONDITION
0	F	57.0	IMPROPER PARKING	UNKNOWN	UNKNOWN
2	F	57.0	NONE	NOT OBSCURED	NORMAL
3	M	74.0	UNKNOWN	UNKNOWN	NORMAL
4	M	33.0	UNKNOWN	UNKNOWN	NORMAL
7	F	39.0	FOLLOWED TOO CLOSELY	UNKNOWN	NORMAL

```
[22]: people_fin.shape
```

```
[22]: (892110, 8)
```

2. vehicles

```
[23]: #provide an overview of the unique values and their frequencies for each column
      ↪ in the DataFrame
for col in vehicles.columns[1:]:
    print('\n' + col + '\n')
    print(vehicles[col].value_counts())
```

UNIT_NO

1	717525
2	657433
3	47733
4	9880
5	2585
6	820
7	307
8	137
9	65
10	34
11	19
12	12
13	7
14	6
15	5
16	4
0	4
17	3
18	3
3778035	1

Name: UNIT_NO, dtype: int64

UNIT_TYPE

DRIVER	1230227
PARKED	193434
DRIVERLESS	12614
DISABLED VEHICLE	194
NON-CONTACT VEHICLE	114

Name: UNIT_TYPE, dtype: int64

VEHICLE_ID

1500741.0	1
498046.0	1
497916.0	1
498593.0	1
498591.0	1
..	
998498.0	1
998496.0	1
998392.0	1
998377.0	1
567870.0	1

Name: VEHICLE_ID, Length: 1436583, dtype: int64

VEHICLE_DEFECT

NONE	792628
UNKNOWN	629288
OTHER	7301
BRAKES	4682
TIRES	719
STEERING	670
WHEELS	366
SUSPENSION	246
ENGINE/MOTOR	192
FUEL SYSTEM	172
LIGHTS	90
WINDOWS	86
CARGO	46
SIGNALS	39
RESTRAINT SYSTEM	21
TRAILER COUPLING	19
EXHAUST	18

Name: VEHICLE_DEFECT, dtype: int64

VEHICLE_TYPE

PASSENGER	904181
SPORT UTILITY VEHICLE (SUV)	194473
UNKNOWN/NA	135892
VAN/MINI-VAN	69525
PICKUP	45243
TRUCK - SINGLE UNIT	27111
OTHER	16921
BUS OVER 15 PASS.	15135
TRACTOR W/ SEMI-TRAILER	13506
BUS UP TO 15 PASS.	3721
MOTORCYCLE (OVER 150CC)	3271
SINGLE UNIT TRUCK WITH TRAILER	2226
OTHER VEHICLE WITH TRAILER	1867
TRACTOR W/O SEMI-TRAILER	1815
AUTOCYCLE	651
MOPED OR MOTORIZED BICYCLE	406
MOTOR DRIVEN CYCLE	325
ALL-TERRAIN VEHICLE (ATV)	162
FARM EQUIPMENT	75
3-WHEELED MOTORCYCLE (2 REAR WHEELS)	50
RECREATIONAL OFF-HIGHWAY VEHICLE (ROV)	19
SNOWMOBILE	8

Name: VEHICLE_TYPE, dtype: int64

VEHICLE_USE

PERSONAL	928310
UNKNOWN/NA	291176
NOT IN USE	74689
OTHER	44891
TAXI/FOR HIRE	18880
COMMERCIAL - SINGLE UNIT	17517
RIDESHARE SERVICE	11997
CTA	9763
POLICE	9360
CONSTRUCTION/MAINTENANCE	6467
COMMERCIAL - MULTI-UNIT	5819
OTHER TRANSIT	3964
SCHOOL BUS	3872
TOW TRUCK	2752
AMBULANCE	1694
FIRE	1402
STATE OWNED	1200
DRIVER EDUCATION	1140
MASS TRANSIT	832
LAWN CARE/LANDSCAPING	533
AGRICULTURE	151
CAMPER/RV - SINGLE UNIT	78
MILITARY	55
HOUSE TRAILER	23
CAMPER/RV - TOWED/MULTI-UNIT	18

Name: VEHICLE_USE, dtype: int64

TRAVEL_DIRECTION

N	338187
S	330540
W	299337
E	293024
UNKNOWN	113540
SE	18112
NW	16560
SW	13693
NE	13590

Name: TRAVEL_DIRECTION, dtype: int64

MANEUVER

STRAIGHT AHEAD	655794
PARKED	198043
UNKNOWN/NA	112104
SLOW/STOP IN TRAFFIC	109132
TURNING LEFT	84673

BACKING	59023
TURNING RIGHT	47093
PASSING/OVERTAKING	34593
CHANGING LANES	27788
OTHER	23927
ENTERING TRAFFIC LANE FROM PARKING	16955
MERGING	9985
STARTING IN TRAFFIC	8406
U-TURN	8110
LEAVING TRAFFIC LANE TO PARK	6940
AVOIDING VEHICLES/OBJECTS	6202
SKIDDING/CONTROL LOSS	5777
ENTER FROM DRIVE/ALLEY	5534
PARKED IN TRAFFIC LANE	4348
SLOW/STOP - LEFT TURN	3042
DRIVING WRONG WAY	2104
SLOW/STOP - RIGHT TURN	1927
NEGOTIATING A CURVE	1901
SLOW/STOP - LOAD/UNLOAD	1663
TURNING ON RED	558
DRIVERLESS	544
DIVERGING	222
DISABLED	195

Name: MANEUVER, dtype: int64

FIRST_CONTACT_POINT

FRONT	280377
REAR	189471
UNKNOWN	137428
SIDE-LEFT	98025
SIDE-RIGHT	93042
FRONT-LEFT	81521
FRONT-LEFT-CORNER	79328
FRONT-RIGHT-CORNER	77441
FRONT-RIGHT	76901
REAR-LEFT	68076
OTHER	39701
REAR-RIGHT	36487
REAR-LEFT-CORNER	35377
TOTAL (ALL AREAS)	26362
REAR-RIGHT-CORNER	25849
SIDE-LEFT-REAR	20584
SIDE-RIGHT-REAR	15419
SIDE-LEFT-FRONT	13121
ROOF	11947
NONE	11836
SIDE-RIGHT-FRONT	11262

```

UNDER CARRIAGE          5442
TOP                     1586
Name: FIRST_CONTACT_POINT, dtype: int64

```

- We will use 'VEHICLE_ID' to merge this dataset with the People dataset. We will drop 'CRASH_RECORD_ID'.
- UNIT_NO is the number of parties involved. We are not concerned with the count so we will drop this.
- UNIT_TYPE seems to be also irrelevant for our EDA. We will drop this.
- VEHICLE_DEFECT has about 90% of the data either as NONE or UNKNOWN. We will drop this as well.
- VEHICLE_TYPE is spread pretty thinly outside of PASSENGER and SUV. We will drop this column as well.
- VEHICLE_USE has majority PERSONAL and UNKNOWN. We will drop this.
- TRAVEL_DIRECTION does not add much to the target variable. We will drop this.
- FIRST_CONTACT_POINT also does not lead to the target variable. We will drop this.

```

[24]: vehicles_del= ['CRASH_RECORD_ID', 'UNIT_NO', 'UNIT_TYPE', 'VEHICLE_DEFECT',
                    'VEHICLE_TYPE',
                    ↪ 'VEHICLE_USE', 'TRAVEL_DIRECTION', 'FIRST_CONTACT_POINT']

```

```

[25]: vehicles.drop(vehicles_del, axis=1, inplace=True)

```

```

[26]: vehicles.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 1436583 entries, 0 to 1472815
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   VEHICLE_ID  1436583 non-null  float64
1   MANEUVER    1436583 non-null  object
dtypes: float64(1), object(1)
memory usage: 32.9+ MB

```

Manuever seems to be the only column with usable information from the vehicle dataset

3. crashes

predators

```

[27]: # provide an overview of the unique values and their frequencies for each
      ↪ column in the DataFrame
for col in crashes[1:]:
    print('\n' + col + '\n')
    print(crashes[col].value_counts())

```

CRASH_RECORD_ID


```

25d92973475a04a93e7fd206fbfce57e8a9a1e25cc85a7e998bb71e476a95e2cb27abd1cef40a8ef
d9ec4929c34da8f7f5403333b420bf4ca753bf77fd8417fb      1
f6c245fb13de9caad165e3e1b47f7cf983a5d2cd20778c10e27f32bc8397b8b35519328803bc3124
0c1ed821e967f06c7da57a1a6819c7b36ffc0f5edefa89d8      1
27372c090bbfe8d0d75174a1f8bd4f671cd23c7d3adb432f9a23d536ba39acdba613aac4ab1bed63
d6be672c2754c04b274fa9f3b74a816c43c62e4335358c4c      1
1eb85c55fa0a4aa0d786373f71c41ea3018e1ecd7e4cfdc4f17a23b54d6a4f7f60918311726e851b
63a91aa35b062e9f9f58579c4f7da8a73cd1f7a61abdd6b3      1
d25134adc4da9eb829478249bb81e772c2518e6af1259596608e7c383698066f89f0c6cfb2a50069
04581d783c883ba098632cedd9b33a1de4b84c6dddde8b5e      1

```

..

```

fe508a7f777e72abe27dc44523631785c8b154d0c1f64134442d786d1a52c7cce1e53992d8b653b6
538ecce2b3b8d6f68642b1c8c417b984ad21886cd992ad89      1
7487a5867f896b21bc590f951152cb4490ea78b8d6c93a5788b93470a1d1db5be34d0f708c48f6ae
9361b9bfb16252fea1111e29cd96b73c3d85e5c26fd18ce      1
805fc4a1098a75409f274ca9da96d205e5bc46e6fbd9c4344334716257e944371cd8073f3f412146
8ef75cc41bcebf07e5cc65d352b96a6fbd0e1cdc4d7a63d2      1
0cf0ddd3b5f04310fa83ee31efe6e294073182f877882552238c3fd8d84f4fd821aed2cdb50b46da
b2be7059c1decf79fed661b1b91b722d950a3af705750779      1
a802658be15312809c771559e4f81088cfb226830792a50470f4ecf9dbdc4fd83c1e187199279ea5
3e604a6cc30bc0c0fd5ba00b0c0e924746c0f4a23b44edc5      1

```

Name: CRASH_RECORD_ID, Length: 718130, dtype: int64

POSTED_SPEED_LIMIT

```

30    529461
35    48155
25    45024
20    29451
15    25146
10    16404
0      7260
40    6937
45    4524
5     4356
55     580
50    171
3     163
9      95
39    74
99    66
1     39
60    38
24    36
2     26
32    17
34    14
65    14

```

33	13
11	11
6	7
26	6
7	5
36	5
70	4
12	3
14	3
29	3
4	2
38	2
22	2
23	2
31	2
18	2
8	2
44	1
62	1
49	1
16	1
63	1

Name: POSTED_SPEED_LIMIT, dtype: int64

TRAFFIC_CONTROL_DEVICE

NO CONTROLS	410639
TRAFFIC SIGNAL	199033
STOP SIGN/FLASHER	71454
UNKNOWN	26603
OTHER	4625
LANE USE MARKING	1173
YIELD	1014
OTHER REG. SIGN	735
OTHER WARNING SIGN	582
RAILROAD CROSSING GATE	467
PEDESTRIAN CROSSING SIGN	416
DELINEATORS	278
SCHOOL ZONE	278
FLASHING CONTROL SIGNAL	254
POLICE/FLAGMAN	240
OTHER RAILROAD CROSSING	165
RR CROSSING SIGN	113
NO PASSING	41
BICYCLE CROSSING SIGN	20

Name: TRAFFIC_CONTROL_DEVICE, dtype: int64

DEVICE_CONDITION

NO CONTROLS	415554
FUNCTIONING PROPERLY	247015
UNKNOWN	44104
OTHER	5481
FUNCTIONING IMPROPERLY	3447
NOT FUNCTIONING	2185
WORN REFLECTIVE MATERIAL	262
MISSING	82

Name: DEVICE_CONDITION, dtype: int64

WEATHER_CONDITION

CLEAR	565028
RAIN	62200
UNKNOWN	37202
SNOW	26272
CLOUDY/OVERCAST	21315
OTHER	2280
FREEZING RAIN/DRIZZLE	1343
FOG/SMOKE/HAZE	1039
SLEET/HAIL	928
BLOWING SNOW	377
SEVERE CROSS WIND GATE	141
BLOWING SAND, SOIL, DIRT	5

Name: WEATHER_CONDITION, dtype: int64

LIGHTING_CONDITION

DAYLIGHT	461609
DARKNESS, LIGHTED ROAD	158409
DARKNESS	34381
UNKNOWN	30873
DUSK	20897
DAWN	11961

Name: LIGHTING_CONDITION, dtype: int64

FIRST_CRASH_TYPE

PARKED MOTOR VEHICLE	167732
REAR END	162911
SIDESWIPE SAME DIRECTION	108644
TURNING	101973
ANGLE	77877
FIXED OBJECT	33895
PEDESTRIAN	16485
PEDALCYCLIST	10464
SIDESWIPE OPPOSITE DIRECTION	10236

OTHER OBJECT	7108
REAR TO FRONT	6278
HEAD ON	6153
REAR TO SIDE	3736
OTHER NONCOLLISION	2324
REAR TO REAR	1321
ANIMAL	508
OVERTURNED	446
TRAIN	39

Name: FIRST_CRASH_TYPE, dtype: int64

TRAFFICWAY_TYPE

NOT DIVIDED	315998
DIVIDED - W/MEDIAN (NOT RAISED)	117894
ONE-WAY	92191
PARKING LOT	48957
DIVIDED - W/MEDIAN BARRIER	41339
FOUR WAY	39211
OTHER	19457
ALLEY	11941
UNKNOWN	8124
T-INTERSECTION	8007
CENTER TURN LANE	5340
DRIVEWAY	2350
RAMP	2150
UNKNOWN INTERSECTION TYPE	1952
FIVE POINT, OR MORE	882
Y-INTERSECTION	864
TRAFFIC ROUTE	713
NOT REPORTED	423
ROUNDBOUT	205
L-INTERSECTION	132

Name: TRAFFICWAY_TYPE, dtype: int64

ALIGNMENT

STRAIGHT AND LEVEL	700534
STRAIGHT ON GRADE	9185
CURVE, LEVEL	5130
STRAIGHT ON HILLCREST	1958
CURVE ON GRADE	995
CURVE ON HILLCREST	328

Name: ALIGNMENT, dtype: int64

ROADWAY_SURFACE_COND

DRY	531911
-----	--------

WET	95865
UNKNOWN	56779
SNOW OR SLUSH	26301
ICE	5178
OTHER	1815
SAND, MUD, DIRT	281

Name: ROADWAY_SURFACE_COND, dtype: int64

ROAD_DEFECT

NO DEFECTS	584232
UNKNOWN	119266
RUT, HOLES	5751
OTHER	3991
WORN SURFACE	2954
SHOULDER DEFECT	1378
DEBRIS ON ROADWAY	558

Name: ROAD_DEFECT, dtype: int64

CRASH_TYPE

NO INJURY / DRIVE AWAY	528457
INJURY AND / OR TOW DUE TO CRASH	189673

Name: CRASH_TYPE, dtype: int64

DAMAGE

OVER \$1,500	438210
\$501 - \$1,500	195974
\$500 OR LESS	83946

Name: DAMAGE, dtype: int64

PRIM_CONTRIBUTORY_CAUSE

UNABLE TO DETERMINE	277139
FAILING TO YIELD RIGHT-OF-WAY	78508
FOLLOWING TOO CLOSELY	71129
NOT APPLICABLE	37862
IMPROPER OVERTAKING/PASSING	34895
FAILING TO REDUCE SPEED TO AVOID CRASH	30563
IMPROPER BACKING	28941

IMPROPER LANE USAGE
26124
IMPROPER TURNING/NO SIGNAL
23738
DRIVING SKILLS/KNOWLEDGE/EXPERIENCE
23470
DISREGARDING TRAFFIC SIGNALS
13904
WEATHER
11294
OPERATING VEHICLE IN ERRATIC, RECKLESS, CARELESS, NEGLIGENT OR AGGRESSIVE MANNER
9148
DISREGARDING STOP SIGN
7941
DISTRACTION - FROM INSIDE VEHICLE
4981
EQUIPMENT - VEHICLE CONDITION
4534
PHYSICAL CONDITION OF DRIVER
4371
VISION OBSCURED (SIGNS, TREE LIMBS, BUILDINGS, ETC.)
4106
DRIVING ON WRONG SIDE/WRONG WAY
3752
UNDER THE INFLUENCE OF ALCOHOL/DRUGS (USE WHEN ARREST IS EFFECTED)
3508
DISTRACTION - FROM OUTSIDE VEHICLE
3028
EXCEEDING AUTHORIZED SPEED LIMIT
1967
ROAD ENGINEERING/SURFACE/MARKING DEFECTS
1866
EXCEEDING SAFE SPEED FOR CONDITIONS
1674
ROAD CONSTRUCTION/MAINTENANCE
1599
DISREGARDING OTHER TRAFFIC SIGNS
1548
EVASIVE ACTION DUE TO ANIMAL, OBJECT, NONMOTORIST
1347
CELL PHONE USE OTHER THAN TEXTING
973
DISREGARDING ROAD MARKINGS
902
HAD BEEN DRINKING (USE WHEN ARREST IS NOT MADE)
760
ANIMAL
606

TURNING RIGHT ON RED

509

DISTRACTION - OTHER ELECTRONIC DEVICE (NAVIGATION DEVICE, DVD PLAYER, ETC.)

336

RELATED TO BUS STOP

315

TEXTING

291

DISREGARDING YIELD SIGN

251

PASSING STOPPED SCHOOL BUS

88

BICYCLE ADVANCING LEGALLY ON RED LIGHT

75

OBSTRUCTED CROSSWALKS

67

MOTORCYCLE ADVANCING LEGALLY ON RED LIGHT

20

Name: PRIM_CONTRIBUTORY_CAUSE, dtype: int64

NUM_UNITS

2 627995

1 39874

3 39790

4 7746

5 1858

6 540

7 180

8 78

9 33

10 16

11 7

12 5

18 3

14 2

13 1

15 1

16 1

Name: NUM_UNITS, dtype: int64

CRASH_HOUR

15 55274

16 54844

17 53547

14 48270

18 44235

13	43853
12	42376
8	37875
11	36611
9	33040
10	32640
19	32539
7	30265
20	26283
21	23466
22	21487
23	18641
6	15519
0	15474
1	13186
2	11373
5	9806
3	9263
4	8263

Name: CRASH_HOUR, dtype: int64

CRASH_DAY_OF_WEEK

6	116893
7	106525
5	103058
3	102337
4	101550
2	98944
1	88823

Name: CRASH_DAY_OF_WEEK, dtype: int64

CRASH_MONTH

10	66291
9	62380
12	61300
5	60947
8	60757
11	60092
7	59040
3	58478
6	57578
1	57437
2	57065
4	56765

Name: CRASH_MONTH, dtype: int64

LATITUDE

41.976201	1146
41.900959	670
41.791420	521
41.751461	503
41.722257	397

...

41.909120	1
41.776586	1
41.784034	1
41.891171	1
41.868220	1

Name: LATITUDE, Length: 277256, dtype: int64

LONGITUDE

-87.905309	1146
-87.619928	670
-87.580148	521
-87.585972	503
-87.585276	397

...

-87.586279	1
-87.721550	1
-87.765236	1
-87.653917	1
-87.708311	1

Name: LONGITUDE, Length: 277227, dtype: int64

LOCATION

POINT (-87.905309125103 41.976201139024)	1146
POINT (-87.619928173678 41.900958919109)	670
POINT (-87.580147768689 41.791420282098)	521
POINT (-87.585971992965 41.751460603167)	503
POINT (-87.585275565077 41.722257273006)	397

...

POINT (-87.706015088063 41.824268608978)	1
POINT (-87.626643215394 41.767465371623)	1
POINT (-87.586278704646 41.758654278901)	1
POINT (-87.721550137644 41.742757072518)	1
POINT (-87.708310986354 41.86822049587)	1

Name: LOCATION, Length: 277408, dtype: int64

[28]: `crashes.shape`

[28]: (718130, 21)

We will use 'PRIM_CONTRIBUTORY_CAUSE' as our target variables.

Target variable

```
[29]: crashes['PRIM_CONTRIBUTORY_CAUSE'].value_counts()
```

```
[29]: UNABLE TO DETERMINE
277139
FAILING TO YIELD RIGHT-OF-WAY
78508
FOLLOWING TOO CLOSELY
71129
NOT APPLICABLE
37862
IMPROPER OVERTAKING/PASSING
34895
FAILING TO REDUCE SPEED TO AVOID CRASH
30563
IMPROPER BACKING
28941
IMPROPER LANE USAGE
26124
IMPROPER TURNING/NO SIGNAL
23738
DRIVING SKILLS/KNOWLEDGE/EXPERIENCE
23470
DISREGARDING TRAFFIC SIGNALS
13904
WEATHER
11294
OPERATING VEHICLE IN ERRATIC, RECKLESS, CARELESS, NEGLIGENT OR AGGRESSIVE MANNER
9148
DISREGARDING STOP SIGN
7941
DISTRACTION - FROM INSIDE VEHICLE
4981
EQUIPMENT - VEHICLE CONDITION
4534
PHYSICAL CONDITION OF DRIVER
4371
VISION OBSCURED (SIGNS, TREE LIMBS, BUILDINGS, ETC.)
4106
DRIVING ON WRONG SIDE/WRONG WAY
3752
UNDER THE INFLUENCE OF ALCOHOL/DRUGS (USE WHEN ARREST IS EFFECTED)
3508
```

DISTRACTION - FROM OUTSIDE VEHICLE
 3028
 EXCEEDING AUTHORIZED SPEED LIMIT
 1967
 ROAD ENGINEERING/SURFACE/MARKING DEFECTS
 1866
 EXCEEDING SAFE SPEED FOR CONDITIONS
 1674
 ROAD CONSTRUCTION/MAINTENANCE
 1599
 DISREGARDING OTHER TRAFFIC SIGNS
 1548
 EVASIVE ACTION DUE TO ANIMAL, OBJECT, NONMOTORIST
 1347
 CELL PHONE USE OTHER THAN TEXTING
 973
 DISREGARDING ROAD MARKINGS
 902
 HAD BEEN DRINKING (USE WHEN ARREST IS NOT MADE)
 760
 ANIMAL
 606
 TURNING RIGHT ON RED
 509
 DISTRACTION - OTHER ELECTRONIC DEVICE (NAVIGATION DEVICE, DVD PLAYER, ETC.)
 336
 RELATED TO BUS STOP
 315
 TEXTING
 291
 DISREGARDING YIELD SIGN
 251
 PASSING STOPPED SCHOOL BUS
 88
 BICYCLE ADVANCING LEGALLY ON RED LIGHT
 75
 OBSTRUCTED CROSSWALKS
 67
 MOTORCYCLE ADVANCING LEGALLY ON RED LIGHT
 20
 Name: PRIM_CONTRIBUTORY_CAUSE, dtype: int64

```

[30]: ## Filter crashes based on specific conditions
      crashes_fin=crashes[(crashes.PRIM_CONTRIBUTORY_CAUSE != 'UNABLE TO DETERMINE') &
                           (crashes.PRIM_CONTRIBUTORY_CAUSE != 'NOT_
                           ↳APPLICABLE')].copy()
  
```

```
[31]: # cheaking the shape of our dataset
crashes_fin.shape
```

```
[31]: (403129, 21)
```

8.0.3 merging the dataset

```
[32]: #Merging Datasets on the Crash_Record_ID Column
df_1 = pd.merge(people_fin, vehicles, on='VEHICLE_ID')
merg_data = pd.merge(df_1, crashes_fin, on='CRASH_RECORD_ID')
merg_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 577429 entries, 0 to 577428
Data columns (total 29 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   PERSON_TYPE                           577429 non-null object
1   CRASH_RECORD_ID                       577429 non-null object
2   VEHICLE_ID                           577429 non-null float64
3   SEX                                   577429 non-null object
4   AGE                                   577429 non-null float64
5   DRIVER_ACTION                         577429 non-null object
6   DRIVER_VISION                        577429 non-null object
7   PHYSICAL_CONDITION                   577429 non-null object
8   MANEUVER                             577429 non-null object
9   POSTED_SPEED_LIMIT                  577429 non-null int64
10  TRAFFIC_CONTROL_DEVICE               577429 non-null object
11  DEVICE_CONDITION                     577429 non-null object
12  WEATHER_CONDITION                    577429 non-null object
13  LIGHTING_CONDITION                   577429 non-null object
14  FIRST_CRASH_TYPE                     577429 non-null object
15  TRAFFICWAY_TYPE                      577429 non-null object
16  ALIGNMENT                            577429 non-null object
17  ROADWAY_SURFACE_COND                 577429 non-null object
18  ROAD_DEFECT                          577429 non-null object
19  CRASH_TYPE                           577429 non-null object
20  DAMAGE                              577429 non-null object
21  PRIM_CONTRIBUTORY_CAUSE              577429 non-null object
22  NUM_UNITS                            577429 non-null int64
23  CRASH_HOUR                           577429 non-null int64
24  CRASH_DAY_OF_WEEK                    577429 non-null int64
25  CRASH_MONTH                          577429 non-null int64
26  LATITUDE                             577429 non-null float64
27  LONGITUDE                            577429 non-null float64
28  LOCATION                             577429 non-null object
dtypes: float64(4), int64(5), object(20)
```

memory usage: 132.2+ MB

```
[33]: # Cheaking the shape of the data
      merg_data.shape
```

```
[33]: (577429, 29)
```

```
[34]: # cheaking the head of the data
      merg_data.head()
```

```
[34]: PERSON_TYPE                                CRASH_RECORD_ID  VEHICLE_ID  \
0      DRIVER  25d92973475a04a93e7fd206fbfce57e8a9a1e25cc85a7...  1500741.0
1      DRIVER  25d92973475a04a93e7fd206fbfce57e8a9a1e25cc85a7...  1500742.0
2      DRIVER  375ac7f6fcb4ef73d728edc52ed556f23fd465a351833f...  1500759.0
3      DRIVER  375ac7f6fcb4ef73d728edc52ed556f23fd465a351833f...  1500760.0
4      DRIVER  375ac7f6fcb4ef73d728edc52ed556f23fd465a351833f...  1500761.0

      SEX  AGE  DRIVER_ACTION DRIVER_VISION  PHYSICAL_CONDITION  \
0  F  20.0  FOLLOWED TOO CLOSELY      UNKNOWN      NORMAL
1  F  53.0      UNKNOWN      UNKNOWN      NORMAL
2  M  22.0  IMPROPER BACKING  NOT OBSCURED  IMPAIRED - ALCOHOL
3  M  67.0      NONE  NOT OBSCURED      NORMAL
4  M  54.0      NONE  NOT OBSCURED      NORMAL

      MANEUVER  POSTED_SPEED_LIMIT  ...  CRASH_TYPE  \
0  STRAIGHT AHEAD      30  ...  NO INJURY / DRIVE AWAY
1  STRAIGHT AHEAD      30  ...  NO INJURY / DRIVE AWAY
2      BACKING      30  ...  NO INJURY / DRIVE AWAY
3  STRAIGHT AHEAD      30  ...  NO INJURY / DRIVE AWAY
4  STRAIGHT AHEAD      30  ...  NO INJURY / DRIVE AWAY

      DAMAGE  PRIM_CONTRIBUTORY_CAUSE  NUM_UNITS  \
0  $501 - $1,500  FOLLOWING TOO CLOSELY      2
1  $501 - $1,500  FOLLOWING TOO CLOSELY      2
2  $501 - $1,500  UNDER THE INFLUENCE OF ALCOHOL/DRUGS (USE WHEN...  3
3  $501 - $1,500  UNDER THE INFLUENCE OF ALCOHOL/DRUGS (USE WHEN...  3
4  $501 - $1,500  UNDER THE INFLUENCE OF ALCOHOL/DRUGS (USE WHEN...  3

      CRASH_HOUR  CRASH_DAY_OF_WEEK  CRASH_MONTH  LATITUDE  LONGITUDE  \
0      23      3      5  41.952691 -87.807413
1      23      3      5  41.952691 -87.807413
2      23      3      5  41.997837 -87.688814
3      23      3      5  41.997837 -87.688814
4      23      3      5  41.997837 -87.688814

      LOCATION
0  POINT (-87.807413247555 41.952691362649)
```

```

1 POINT (-87.807413247555 41.952691362649)
2 POINT (-87.688813887189 41.997837266972)
3 POINT (-87.688813887189 41.997837266972)
4 POINT (-87.688813887189 41.997837266972)

```

[5 rows x 29 columns]

```

[35]: #checking columns of our merged dataset
      merg_data.columns

```

```

[35]: Index(['PERSON_TYPE', 'CRASH_RECORD_ID', 'VEHICLE_ID', 'SEX', 'AGE',
            'DRIVER_ACTION', 'DRIVER_VISION', 'PHYSICAL_CONDITION', 'MANEUVER',
            'POSTED_SPEED_LIMIT', 'TRAFFIC_CONTROL_DEVICE', 'DEVICE_CONDITION',
            'WEATHER_CONDITION', 'LIGHTING_CONDITION', 'FIRST_CRASH_TYPE',
            'TRAFFICWAY_TYPE', 'ALIGNMENT', 'ROADWAY_SURFACE_COND', 'ROAD_DEFECT',
            'CRASH_TYPE', 'DAMAGE', 'PRIM_CONTRIBUTORY_CAUSE', 'NUM_UNITS',
            'CRASH_HOUR', 'CRASH_DAY_OF_WEEK', 'CRASH_MONTH', 'LATITUDE',
            'LONGITUDE', 'LOCATION'],
           dtype='object')

```

```

[36]: # check the unique variables for each category.
      for col in merg_data.columns:
          print('\n' + col + '\n')
          print(merg_data[col].value_counts())

```

PERSON_TYPE

DRIVER 577429

Name: PERSON_TYPE, dtype: int64

CRASH_RECORD_ID

```

c9d233e31a4f2a07733ef75f0404e75c360b30c7ee9bc45076938dc80c375578c1468bc096ecb773
d2bfc71270d746d95f416a5bd6b15fbcf8707b1748693722      12
196e0d42ec2dd3c503f98ad28d08067091e9801170ce6b264599642baf11c87f2064fdeccff3cd9d
ed1c8d7bf2329640af1e4730adfc3a36127f5245ec7e152        9
7be311dead41c5337cbd12d40bb7be93c505303d6f1cf92e72a2b7c695ae95b472a66d9b3a6b505a
0e4c2279d53acf3b6115320fcab54d8ee1aa3d0c811e3a0        9
2294d7387552dd1804e1eddd6b4ce561209f90ede1fb98805a0658ad6fcb9f8d3f8846416358dbf
8b9e71db2095d505387d2ce4ad894ad9a5a9ea5ae691abaa        9
d3c41d043f9f56c4ab63d2e0e6d229bd1283ffd5e600b0f5518e0493e41415c72b369d05f4bb9804
c784b7e5789a3e84ae0d600300234f496a00c0a43a8e8d75        8
..
45e733513cc4265f69b45a19840c6391430fe641d3bccd344a69b1c78d84978418de9d57e93d16bb
eae4e5aeeb09d3437036741c60da9c446cbd1dc02179c31        1
9174e4112627175c3ef23f38bb488b99282fe07ae0c48ee795a998a7779388ae0f4d18d08b71d258

```

```

7e0035ae2dd6ae4fab6f575500a0e350fc29107d1ec68c5d      1
0601e8309277db928d5579ecab19a9c9856ebeba2024986c91ceb9906487f7722039e60da9734cc5
0b6ae7639ecd4dcfce1f7bd780c6a223ba9aaf039b1730ee      1
17520e7ada1c93b9d5eeda92e555fa10d90afa78b5e0e6f4a03e57179ef5933b2e92c174dfa7f12f
8fa6e89543883f9ae7bc8a018fbd8bd76e8ef69ca815c0b0      1
a802658be15312809c771559e4f81088cfb226830792a50470f4ecf9dbdc4fd83c1e187199279ea5
3e604a6cc30bc0c0fd5ba00b0c0e924746c0f4a23b44edc5      1
Name: CRASH_RECORD_ID, Length: 360447, dtype: int64

```

VEHICLE_ID

```

1500741.0      1
469979.0      1
470005.0      1
470017.0      1
470032.0      1
..
1043136.0      1
1043135.0      1
953192.0      1
953184.0      1
567870.0      1
Name: VEHICLE_ID, Length: 577429, dtype: int64

```

SEX

```

M      345523
F      231906
Name: SEX, dtype: int64

```

AGE

```

27.0      17531
25.0      17469
26.0      17259
28.0      17065
24.0      16647
...
110.0      2
108.0      2
109.0      1
104.0      1
107.0      1
Name: AGE, Length: 108, dtype: int64

```

DRIVER_ACTION

```

NONE      269888

```

FAILED TO YIELD	69467
OTHER	51970
FOLLOWED TOO CLOSELY	47494
UNKNOWN	38207
IMPROPER TURN	19894
IMPROPER BACKING	18826
IMPROPER LANE CHANGE	17623
DISREGARDED CONTROL DEVICES	12344
IMPROPER PASSING	12328
TOO FAST FOR CONDITIONS	11212
WRONG WAY/SIDE	2334
IMPROPER PARKING	1683
CELL PHONE USE OTHER THAN TEXTING	1094
EVADING POLICE VEHICLE	1010
OVERCORRECTED	993
EMERGENCY VEHICLE ON CALL	623
TEXTING	305
STOPPED SCHOOL BUS	98
LICENSE RESTRICTIONS	36

Name: DRIVER_ACTION, dtype: int64

DRIVER_VISION

NOT OBSCURED	402587
UNKNOWN	154308
OTHER	7213
MOVING VEHICLES	5407
PARKED VEHICLES	3337
WINDSHIELD (WATER/ICE)	2459
BLINDED - SUNLIGHT	1157
TREES, PLANTS	374
BUILDINGS	306
BLINDED - HEADLIGHTS	75
HILLCREST	65
EMBANKMENT	63
BLOWING MATERIALS	53
SIGNBOARD	25

Name: DRIVER_VISION, dtype: int64

PHYSICAL_CONDITION

NORMAL	498584
UNKNOWN	64022
IMPAIRED - ALCOHOL	4392
FATIGUED/ASLEEP	2382
REMOVED BY EMS	2067
OTHER	1851
EMOTIONAL	1755

ILLNESS/FAINTED	940
HAD BEEN DRINKING	559
IMPAIRED - DRUGS	503
IMPAIRED - ALCOHOL AND DRUGS	260
MEDICATED	114

Name: PHYSICAL_CONDITION, dtype: int64

MANEUVER

STRAIGHT AHEAD	334081
SLOW/STOP IN TRAFFIC	64937
TURNING LEFT	47831
TURNING RIGHT	24298
BACKING	23304
PASSING/OVERTAKING	14445
CHANGING LANES	13288
OTHER	9067
ENTERING TRAFFIC LANE FROM PARKING	8040
UNKNOWN/NA	5428
STARTING IN TRAFFIC	4890
MERGING	4491
U-TURN	4459
SKIDDING/CONTROL LOSS	3671
AVOIDING VEHICLES/OBJECTS	3501
ENTER FROM DRIVE/ALLEY	2823
LEAVING TRAFFIC LANE TO PARK	2739
SLOW/STOP - LEFT TURN	1867
SLOW/STOP - RIGHT TURN	1168
DRIVING WRONG WAY	951
NEGOTIATING A CURVE	888
SLOW/STOP - LOAD/UNLOAD	883
TURNING ON RED	275
DIVERGING	98
PARKED	5
PARKED IN TRAFFIC LANE	1

Name: MANEUVER, dtype: int64

POSTED_SPEED_LIMIT

30	447733
35	44864
25	28767
20	15932
15	12952
40	7786
10	7127
45	4680
0	4345

5	2104
55	541
50	171
3	136
39	73
9	56
60	35
34	14
32	13
2	13
1	12
33	12
99	10
24	9
7	8
11	7
36	5
65	4
44	2
31	2
63	2
12	2
70	2
23	2
38	2
29	2
49	1
4	1
6	1
26	1

Name: POSTED_SPEED_LIMIT, dtype: int64

TRAFFIC_CONTROL_DEVICE

NO CONTROLS	268293
TRAFFIC SIGNAL	214682
STOP SIGN/FLASHER	74589
UNKNOWN	10019
OTHER	3784
LANE USE MARKING	1415
YIELD	1181
OTHER REG. SIGN	720
OTHER WARNING SIGN	594
RAILROAD CROSSING GATE	467
PEDESTRIAN CROSSING SIGN	386
DELINEATORS	299
POLICE/FLAGMAN	264
FLASHING CONTROL SIGNAL	256

SCHOOL ZONE	190
OTHER RAILROAD CROSSING	155
RR CROSSING SIGN	70
NO PASSING	53
BICYCLE CROSSING SIGN	12

Name: TRAFFIC_CONTROL_DEVICE, dtype: int64

DEVICE_CONDITION

NO CONTROLS	275828
FUNCTIONING PROPERLY	268945
UNKNOWN	21966
OTHER	4683
FUNCTIONING IMPROPERLY	3687
NOT FUNCTIONING	1935
WORN REFLECTIVE MATERIAL	292
MISSING	93

Name: DEVICE_CONDITION, dtype: int64

WEATHER_CONDITION

CLEAR	461336
RAIN	57824
SNOW	23114
CLOUDY/OVERCAST	20434
UNKNOWN	9340
OTHER	1814
FREEZING RAIN/DRIZZLE	1250
SLEET/HAIL	935
FOG/SMOKE/HAZE	921
BLOWING SNOW	343
SEVERE CROSS WIND GATE	114
BLOWING SAND, SOIL, DIRT	4

Name: WEATHER_CONDITION, dtype: int64

LIGHTING_CONDITION

DAYLIGHT	395940
DARKNESS, LIGHTED ROAD	126260
DARKNESS	22850
DUSK	17561
DAWN	9363
UNKNOWN	5455

Name: LIGHTING_CONDITION, dtype: int64

FIRST_CRASH_TYPE

REAR END	177392
----------	--------

TURNING	121054
SIDESWIPE SAME DIRECTION	102062
ANGLE	89605
PARKED MOTOR VEHICLE	33397
FIXED OBJECT	13730
SIDESWIPE OPPOSITE DIRECTION	8748
HEAD ON	6668
REAR TO FRONT	6319
PEDESTRIAN	5705
PEDALCYCLIST	3973
REAR TO SIDE	3901
OTHER OBJECT	2416
REAR TO REAR	1008
OTHER NONCOLLISION	927
ANIMAL	278
OVERTURNED	218
TRAIN	28

Name: FIRST_CRASH_TYPE, dtype: int64

TRAFFICWAY_TYPE

NOT DIVIDED	256091
DIVIDED - W/MEDIAN (NOT RAISED)	114787
ONE-WAY	51042
FOUR WAY	44614
DIVIDED - W/MEDIAN BARRIER	42417
PARKING LOT	22391
OTHER	13789
T-INTERSECTION	8576
CENTER TURN LANE	6297
ALLEY	6148
UNKNOWN	2735
RAMP	1967
UNKNOWN INTERSECTION TYPE	1761
DRIVEWAY	1363
FIVE POINT, OR MORE	979
Y-INTERSECTION	889
TRAFFIC ROUTE	836
NOT REPORTED	477
ROUNDBOUT	140
L-INTERSECTION	130

Name: TRAFFICWAY_TYPE, dtype: int64

ALIGNMENT

STRAIGHT AND LEVEL	560713
STRAIGHT ON GRADE	8877
CURVE, LEVEL	4610

STRAIGHT ON HILLCREST	2012
CURVE ON GRADE	929
CURVE ON HILLCREST	288

Name: ALIGNMENT, dtype: int64

ROADWAY_SURFACE_COND

DRY	437440
WET	89271
UNKNOWN	22479
SNOW OR SLUSH	21727
ICE	5074
OTHER	1266
SAND, MUD, DIRT	172

Name: ROADWAY_SURFACE_COND, dtype: int64

ROAD_DEFECT

NO DEFECTS	494189
UNKNOWN	72406
RUT, HOLES	3453
OTHER	3284
WORN SURFACE	2455
SHOULDER DEFECT	1152
DEBRIS ON ROADWAY	490

Name: ROAD_DEFECT, dtype: int64

CRASH_TYPE

NO INJURY / DRIVE AWAY	386447
INJURY AND / OR TOW DUE TO CRASH	190982

Name: CRASH_TYPE, dtype: int64

DAMAGE

OVER \$1,500	382384
\$501 - \$1,500	141125
\$500 OR LESS	53920

Name: DAMAGE, dtype: int64

PRIM_CONTRIBUTORY_CAUSE

FAILING TO YIELD RIGHT-OF-WAY	124046
FOLLOWING TOO CLOSELY	114808
IMPROPER OVERTAKING/PASSING	46845

FAILING TO REDUCE SPEED TO AVOID CRASH
45147
IMPROPER TURNING/NO SIGNAL
36551
IMPROPER LANE USAGE
36236
IMPROPER BACKING
31879
DRIVING SKILLS/KNOWLEDGE/EXPERIENCE
26422
DISREGARDING TRAFFIC SIGNALS
23027
WEATHER
15666
DISREGARDING STOP SIGN
11590
OPERATING VEHICLE IN ERRATIC, RECKLESS, CARELESS, NEGLIGENT OR AGGRESSIVE MANNER
8158
DISTRACTION - FROM INSIDE VEHICLE
7575
VISION OBSCURED (SIGNS, TREE LIMBS, BUILDINGS, ETC.)
6106
EQUIPMENT - VEHICLE CONDITION
6083
PHYSICAL CONDITION OF DRIVER
5455
UNDER THE INFLUENCE OF ALCOHOL/DRUGS (USE WHEN ARREST IS EFFECTED)
4862
DRIVING ON WRONG SIDE/WRONG WAY
4551
DISTRACTION - FROM OUTSIDE VEHICLE
4087
EXCEEDING SAFE SPEED FOR CONDITIONS
2228
DISREGARDING OTHER TRAFFIC SIGNS
2220
EXCEEDING AUTHORIZED SPEED LIMIT
2061
ROAD CONSTRUCTION/MAINTENANCE
1864
ROAD ENGINEERING/SURFACE/MARKING DEFECTS
1847
EVASIVE ACTION DUE TO ANIMAL, OBJECT, NONMOTORIST
1642
CELL PHONE USE OTHER THAN TEXTING
1372
DISREGARDING ROAD MARKINGS
1212

TURNING RIGHT ON RED
755
HAD BEEN DRINKING (USE WHEN ARREST IS NOT MADE)
653
ANIMAL
576
DISTRACTION - OTHER ELECTRONIC DEVICE (NAVIGATION DEVICE, DVD PLAYER, ETC.)
510
RELATED TO BUS STOP
404
TEXTING
398
DISREGARDING YIELD SIGN
327
PASSING STOPPED SCHOOL BUS
105
OBSTRUCTED CROSSWALKS
77
BICYCLE ADVANCING LEGALLY ON RED LIGHT
61
MOTORCYCLE ADVANCING LEGALLY ON RED LIGHT
23
Name: PRIM_CONTRIBUTORY_CAUSE, dtype: int64

NUM_UNITS

2	502239
3	47233
1	15406
4	9390
5	2023
6	699
7	208
8	135
9	49
10	22
12	11
18	6
11	6
15	1
16	1

Name: NUM_UNITS, dtype: int64

CRASH_HOUR

16	47029
15	46175
17	46149

14	39356
18	37032
13	35432
12	34016
8	32932
11	30545
9	27520
10	26900
7	26500
19	26095
20	20380
21	17939
22	16095
23	13594
6	12104
0	10137
1	8180
2	6865
5	6437
3	5322
4	4695

Name: CRASH_HOUR, dtype: int64

CRASH_DAY_OF_WEEK

6	96922
5	85436
4	83736
7	83672
3	83247
2	78489
1	65927

Name: CRASH_DAY_OF_WEEK, dtype: int64

CRASH_MONTH

10	53673
12	50276
9	49760
5	49063
11	48996
8	47978
1	47108
7	46961
3	46683
6	45740
4	45731
2	45460

Name: CRASH_MONTH, dtype: int64

LATITUDE

41.976201	969
41.900959	600
41.791420	458
41.751461	456
41.880856	324

...

41.889348	1
41.780923	1
41.740613	1
41.860918	1
41.835886	1

Name: LATITUDE, Length: 155350, dtype: int64

LONGITUDE

-87.905309	969
-87.619928	600
-87.580148	458
-87.585972	456
-87.617636	324

...

-87.661895	1
-87.570930	1
-87.665346	1
-87.548734	1
-87.724474	1

Name: LONGITUDE, Length: 155334, dtype: int64

LOCATION

POINT (-87.905309125103 41.976201139024)	969
POINT (-87.619928173678 41.900958919109)	600
POINT (-87.580147768689 41.791420282098)	458
POINT (-87.585971992965 41.751460603167)	456
POINT (-87.617635891755 41.880856047671)	324

...

POINT (-87.663909046208 41.896235950067)	1
POINT (-87.711181809431 41.891937228592)	1
POINT (-87.765246938384 41.778133615169)	1
POINT (-87.618487458568 41.896927006997)	1
POINT (-87.724474013253 41.835886103363)	1

Name: LOCATION, Length: 155401, dtype: int64

- CRASH_RECORD_ID, VEHICLE_ID can be dropped as we already used it to merge all

datasets.

- PERSON_TYPE has only one variable. We can drop this.
- FIRST_CRASH_TYPE seems more like a resultant variable rather than an action leading to the crash. We will not use this for columns for our prediction.
- Majority of ALIGNMENT is STRAIGHT AND LEVEL. We will drop this column.
- CRASH_TYPE does not present any insightful information. We will drop this.
- We are not interested in DAMAGE for this EDA. We will drop this.
- NUM_UNITS involved in car crash is also irrelevant for our business case. We will drop this.

```
[37]: # Dropping irrelevant columns
final_data=merg_data.drop(['CRASH_RECORD_ID', 'VEHICLE_ID',
↪ 'PERSON_TYPE', 'FIRST_CRASH_TYPE',
                             'ALIGNMENT', 'CRASH_TYPE', 'DAMAGE', 'NUM_UNITS'],
↪ axis=1).copy()
```

```
[38]: # checking the head of our dataset
final_data.head()
```

```
[38]:  SEX    AGE      DRIVER_ACTION DRIVER_VISION  PHYSICAL_CONDITION  \
0   F   20.0  FOLLOWED TOO CLOSELY      UNKNOWN             NORMAL
1   F   53.0                UNKNOWN      UNKNOWN             NORMAL
2   M   22.0    IMPROPER BACKING  NOT OBSCURED  IMPAIRED - ALCOHOL
3   M   67.0                NONE  NOT OBSCURED             NORMAL
4   M   54.0                NONE  NOT OBSCURED             NORMAL

      MANEUVER  POSTED_SPEED_LIMIT  TRAFFIC_CONTROL_DEVICE  \
0  STRAIGHT AHEAD                30      TRAFFIC SIGNAL
1  STRAIGHT AHEAD                30      TRAFFIC SIGNAL
2          BACKING                30          NO CONTROLS
3  STRAIGHT AHEAD                30          NO CONTROLS
4  STRAIGHT AHEAD                30          NO CONTROLS

      DEVICE_CONDITION  WEATHER_CONDITION  ...  TRAFFICWAY_TYPE  \
0  FUNCTIONING PROPERLY          CLEAR  ...    NOT DIVIDED
1  FUNCTIONING PROPERLY          CLEAR  ...    NOT DIVIDED
2          NO CONTROLS          CLEAR  ...     ONE-WAY
3          NO CONTROLS          CLEAR  ...     ONE-WAY
4          NO CONTROLS          CLEAR  ...     ONE-WAY

  ROADWAY_SURFACE_COND  ROAD_DEFECT  \
0          DRY  NO DEFECTS
1          DRY  NO DEFECTS
2          DRY  NO DEFECTS
3          DRY  NO DEFECTS
4          DRY  NO DEFECTS

      PRIM_CONTRIBUTORY_CAUSE  CRASH_HOUR  \
```

0		FOLLOWING TOO CLOSELY	23
1		FOLLOWING TOO CLOSELY	23
2	UNDER THE INFLUENCE OF ALCOHOL/DRUGS (USE WHEN...		23
3	UNDER THE INFLUENCE OF ALCOHOL/DRUGS (USE WHEN...		23
4	UNDER THE INFLUENCE OF ALCOHOL/DRUGS (USE WHEN...		23

	CRASH_DAY_OF_WEEK	CRASH_MONTH	LATITUDE	LONGITUDE	\
0	3	5	41.952691	-87.807413	
1	3	5	41.952691	-87.807413	
2	3	5	41.997837	-87.688814	
3	3	5	41.997837	-87.688814	
4	3	5	41.997837	-87.688814	

	LOCATION
0	POINT (-87.807413247555 41.952691362649)
1	POINT (-87.807413247555 41.952691362649)
2	POINT (-87.688813887189 41.997837266972)
3	POINT (-87.688813887189 41.997837266972)
4	POINT (-87.688813887189 41.997837266972)

[5 rows x 21 columns]

```
[39]: # checking the shape of our dataset
final_data.shape
```

```
[39]: (577429, 21)
```

our final data has 577429 rows with 21 columns and most columns are categorical

8.0.4 Feature Engineering

The feature engineering includes creating date features such as day of the week, handling the high cardinality of weather conditions, contributing cause, etc, and perhaps most importantly, downsampling to account for the class imbalance (injuries are more rare than non-injury-causing crashes). I will look at both Target and Predictor variables to see if we can change some variable to make more meaningful interpretation

Target Variable

```
[40]: # List the primary contributory causes as our target
target_list = list(final_data.PRIM_CONTRIBUTORY_CAUSE.unique())
target_list
```

```
[40]: ['FOLLOWING TOO CLOSELY',
'UNDER THE INFLUENCE OF ALCOHOL/DRUGS (USE WHEN ARREST IS EFFECTED)',
'FAILING TO REDUCE SPEED TO AVOID CRASH',
'FAILING TO YIELD RIGHT-OF-WAY',
'OPERATING VEHICLE IN ERRATIC, RECKLESS, CARELESS, NEGLIGENT OR AGGRESSIVE']
```

```

MANNER',
'IMPROPER LANE USAGE',
'DRIVING SKILLS/KNOWLEDGE/EXPERIENCE',
'IMPROPER OVERTAKING/PASSING',
'IMPROPER TURNING/NO SIGNAL',
'DISTRACTION - FROM OUTSIDE VEHICLE',
'ANIMAL',
'DISREGARDING TRAFFIC SIGNALS',
'EQUIPMENT - VEHICLE CONDITION',
'PHYSICAL CONDITION OF DRIVER',
'IMPROPER BACKING',
'DISREGARDING STOP SIGN',
'DRIVING ON WRONG SIDE/WRONG WAY',
'DISTRACTION - FROM INSIDE VEHICLE',
'CELL PHONE USE OTHER THAN TEXTING',
'TURNING RIGHT ON RED',
'VISION OBSCURED (SIGNS, TREE LIMBS, BUILDINGS, ETC.)',
'RELATED TO BUS STOP',
'WEATHER',
'EVASIVE ACTION DUE TO ANIMAL, OBJECT, NONMOTORIST',
'DISTRACTION - OTHER ELECTRONIC DEVICE (NAVIGATION DEVICE, DVD PLAYER, ETC.)',
'DISREGARDING ROAD MARKINGS',
'ROAD CONSTRUCTION/MAINTENANCE',
'DISREGARDING OTHER TRAFFIC SIGNS',
'ROAD ENGINEERING/SURFACE/MARKING DEFECTS',
'TEXTING',
'BICYCLE ADVANCING LEGALLY ON RED LIGHT',
'HAD BEEN DRINKING (USE WHEN ARREST IS NOT MADE)',
'DISREGARDING YIELD SIGN',
'OBSTRUCTED CROSSWALKS',
'PASSING STOPPED SCHOOL BUS',
'MOTORCYCLE ADVANCING LEGALLY ON RED LIGHT',
'EXCEEDING SAFE SPEED FOR CONDITIONS',
'EXCEEDING AUTHORIZED SPEED LIMIT']

```

```

[41]: #checking the length of the data
      len(target_list)

```

[41]: 38

I have 38 target variables which are a bit too many to classify and predict. I will group them into two categories: UNINTENTIONAL (0) and INTENTIONAL (1) causes. UNINTENTIONAL causes are typically related to errors or mistakes made by drivers or external factors that are beyond their control INTENTIONAL causes refer to accidents that occur due to intentional actions or choices made by drivers

I will group them as follows:

UNINTENTIONAL (0):

- IMPROPER OVERTAKING/PASSING
- DISREGARDING TRAFFIC SIGNALS
- DRIVING SKILLS/KNOWLEDGE/EXPERIENCE
- IMPROPER TURNING/NO SIGNAL
- FAILING TO REDUCE SPEED TO AVOID CRASH
- FOLLOWING TOO CLOSELY
- IMPROPER BACKING
- IMPROPER LANE USAGE
- FAILING TO YIELD RIGHT-OF-WAY
- DISREGARDING STOP SIGN
- VISION OBSCURED (SIGNS, TREE LIMBS, BUILDINGS, ETC.)
- EQUIPMENT - VEHICLE CONDITION
- DISREGARDING OTHER TRAFFIC SIGNS
- DRIVING ON WRONG SIDE/WRONG WAY
- WEATHER
- PHYSICAL CONDITION OF DRIVER
- ROAD ENGINEERING/SURFACE/MARKING DEFECTS
- OBSTRUCTED CROSSWALKS
- EXCEEDING AUTHORIZED SPEED LIMIT
- EXCEEDING SAFE SPEED FOR CONDITIONS
- ROAD CONSTRUCTION/MAINTENANCE
- DISREGARDING ROAD MARKINGS
- DISREGARDING YIELD SIGN
- EVASIVE ACTION DUE TO ANIMAL, OBJECT, NONMOTORIST
- ANIMAL
- RELATED TO BUS STOP
- TURNING RIGHT ON RED
- PASSING STOPPED SCHOOL BUS
- BICYCLE ADVANCING LEGALLY ON RED LIGHT
- MOTORCYCLE ADVANCING LEGALLY ON RED LIGHT

INTENTIONAL (1):

- DISTRACTION - FROM INSIDE VEHICLE
- UNDER THE INFLUENCE OF ALCOHOL/DRUGS (USE WHEN ARREST IS EFFECTED)
- DISTRACTION - FROM OUTSIDE VEHICLE
- TEXTING
- DISREGARDING ROAD MARKINGS
- CELL PHONE USE OTHER THAN TEXTING
- DISTRACTION - OTHER ELECTRONIC DEVICE (NAVIGATION DEVICE, DVD PLAYER, ETC.)
- HAD BEEN DRINKING (USE WHEN ARREST IS NOT MADE)

```
[42]: # convert the categorical variable 'PRIM_CONTRIBUTORY_CAUSE' into numerical
      ↪ values that represent the two categories you defined: UNINTENTIONAL (0) and
      ↪ INTENTIONAL (1).
```

```

final_data.PRIM_CONTRIBUTORY_CAUSE = final_data.PRIM_CONTRIBUTORY_CAUSE.map({
  'DRIVING SKILLS/KNOWLEDGE/EXPERIENCE':0,
  'DISTRACTION - FROM INSIDE VEHICLE':0,
  'VISION OBSCURED (SIGNS, TREE LIMBS, BUILDINGS, ETC.)':0,
  'WEATHER':0,
  'DISTRACTION - FROM OUTSIDE VEHICLE':0,
  'ROAD ENGINEERING/SURFACE/MARKING DEFECTS':0,
  'OBSTRUCTED CROSSWALKS':0,
  'BICYCLE ADVANCING LEGALLY ON RED LIGHT':0,
  'MOTORCYCLE ADVANCING LEGALLY ON RED LIGHT':0,
  'FAILING TO YIELD RIGHT-OF-WAY':0,
  'FAILING TO REDUCE SPEED TO AVOID CRASH':0,
  'PHYSICAL CONDITION OF DRIVER':0,
  'TEXTING':0,
  'ROAD CONSTRUCTION/MAINTENANCE':0,
  'EVASIVE ACTION DUE TO ANIMAL, OBJECT, NONMOTORIST':0,
  'ANIMAL':0,
  'CELL PHONE USE OTHER THAN TEXTING':0,
  'RELATED TO BUS STOP':0,
  'TURNING RIGHT ON RED':0,
  'DISTRACTION - OTHER ELECTRONIC DEVICE (NAVIGATION DEVICE, DVD PLAYER, ETC.
↪)':0,
  'PASSING STOPPED SCHOOL BUS':0,
  'EQUIPMENT - VEHICLE CONDITION':0,
  'IMPROPER OVERTAKING/PASSING':1,
  'DISREGARDING TRAFFIC SIGNALS':1,
  'IMPROPER TURNING/NO SIGNAL':1,
  'FOLLOWING TOO CLOSELY':1,
  'IMPROPER BACKING':1,
  'IMPROPER LANE USAGE':1,
  'DISREGARDING STOP SIGN':1,
  'DISREGARDING OTHER TRAFFIC SIGNS':1,
  'DRIVING ON WRONG SIDE/WRONG WAY':1,
  'OPERATING VEHICLE IN ERRATIC, RECKLESS, CARELESS, NEGLIGENT OR AGGRESSIVE_
↪MANNER':1,
  'DISREGARDING ROAD MARKINGS':1,
  'DISREGARDING YIELD SIGN':1,
  'EXCEEDING AUTHORIZED SPEED LIMIT':1,
  'EXCEEDING SAFE SPEED FOR CONDITIONS':1,
  'UNDER THE INFLUENCE OF ALCOHOL/DRUGS (USE WHEN ARREST IS EFFECTED)':1,
  'HAD BEEN DRINKING (USE WHEN ARREST IS NOT MADE)':1
})

```

```

[43]: # counts of primary contributor
final_data.PRIM_CONTRIBUTORY_CAUSE.value_counts()

```

```
[43]: 1    327208
      0    250221
      Name: PRIM_CONTRIBUTORY_CAUSE, dtype: int64
```

There were two categories that will help in training the model:

UNINTENTIONAL values is 250221

INTENTIONAL values is 327208

Preditors

```
[44]: #provide an overview of the unique values and their frequencies for each column
      ↪ in the DataFrame
      for col in final_data.columns:
          print('\n' + col + '\n')
          print(final_data[col].value_counts())
```

SEX

```
M    345523
F    231906
      Name: SEX, dtype: int64
```

AGE

```
27.0    17531
25.0    17469
26.0    17259
28.0    17065
24.0    16647
...
110.0      2
108.0      2
109.0      1
104.0      1
107.0      1
      Name: AGE, Length: 108, dtype: int64
```

DRIVER_ACTION

NONE	269888
FAILED TO YIELD	69467
OTHER	51970
FOLLOWED TOO CLOSELY	47494
UNKNOWN	38207
IMPROPER TURN	19894
IMPROPER BACKING	18826

IMPROPER LANE CHANGE	17623
DISREGARDED CONTROL DEVICES	12344
IMPROPER PASSING	12328
TOO FAST FOR CONDITIONS	11212
WRONG WAY/SIDE	2334
IMPROPER PARKING	1683
CELL PHONE USE OTHER THAN TEXTING	1094
EVADING POLICE VEHICLE	1010
OVERCORRECTED	993
EMERGENCY VEHICLE ON CALL	623
TEXTING	305
STOPPED SCHOOL BUS	98
LICENSE RESTRICTIONS	36

Name: DRIVER_ACTION, dtype: int64

DRIVER_VISION

NOT OBSCURED	402587
UNKNOWN	154308
OTHER	7213
MOVING VEHICLES	5407
PARKED VEHICLES	3337
WINDSHIELD (WATER/ICE)	2459
BLINDED - SUNLIGHT	1157
TREES, PLANTS	374
BUILDINGS	306
BLINDED - HEADLIGHTS	75
HILLCREST	65
EMBANKMENT	63
BLOWING MATERIALS	53
SIGNBOARD	25

Name: DRIVER_VISION, dtype: int64

PHYSICAL_CONDITION

NORMAL	498584
UNKNOWN	64022
IMPAIRED - ALCOHOL	4392
FATIGUED/ASLEEP	2382
REMOVED BY EMS	2067
OTHER	1851
EMOTIONAL	1755
ILLNESS/FAINTED	940
HAD BEEN DRINKING	559
IMPAIRED - DRUGS	503
IMPAIRED - ALCOHOL AND DRUGS	260
MEDICATED	114

Name: PHYSICAL_CONDITION, dtype: int64

MANEUVER

STRAIGHT AHEAD	334081
SLOW/STOP IN TRAFFIC	64937
TURNING LEFT	47831
TURNING RIGHT	24298
BACKING	23304
PASSING/OVERTAKING	14445
CHANGING LANES	13288
OTHER	9067
ENTERING TRAFFIC LANE FROM PARKING	8040
UNKNOWN/NA	5428
STARTING IN TRAFFIC	4890
MERGING	4491
U-TURN	4459
SKIDDING/CONTROL LOSS	3671
AVOIDING VEHICLES/OBJECTS	3501
ENTER FROM DRIVE/ALLEY	2823
LEAVING TRAFFIC LANE TO PARK	2739
SLOW/STOP - LEFT TURN	1867
SLOW/STOP - RIGHT TURN	1168
DRIVING WRONG WAY	951
NEGOTIATING A CURVE	888
SLOW/STOP - LOAD/UNLOAD	883
TURNING ON RED	275
DIVERGING	98
PARKED	5
PARKED IN TRAFFIC LANE	1

Name: MANEUVER, dtype: int64

POSTED_SPEED_LIMIT

30	447733
35	44864
25	28767
20	15932
15	12952
40	7786
10	7127
45	4680
0	4345
5	2104
55	541
50	171
3	136
39	73
9	56

60	35
34	14
32	13
2	13
1	12
33	12
99	10
24	9
7	8
11	7
36	5
65	4
44	2
31	2
63	2
12	2
70	2
23	2
38	2
29	2
49	1
4	1
6	1
26	1

Name: POSTED_SPEED_LIMIT, dtype: int64

TRAFFIC_CONTROL_DEVICE

NO CONTROLS	268293
TRAFFIC SIGNAL	214682
STOP SIGN/FLASHER	74589
UNKNOWN	10019
OTHER	3784
LANE USE MARKING	1415
YIELD	1181
OTHER REG. SIGN	720
OTHER WARNING SIGN	594
RAILROAD CROSSING GATE	467
PEDESTRIAN CROSSING SIGN	386
DELINEATORS	299
POLICE/FLAGMAN	264
FLASHING CONTROL SIGNAL	256
SCHOOL ZONE	190
OTHER RAILROAD CROSSING	155
RR CROSSING SIGN	70
NO PASSING	53
BICYCLE CROSSING SIGN	12

Name: TRAFFIC_CONTROL_DEVICE, dtype: int64

DEVICE_CONDITION

NO CONTROLS	275828
FUNCTIONING PROPERLY	268945
UNKNOWN	21966
OTHER	4683
FUNCTIONING IMPROPERLY	3687
NOT FUNCTIONING	1935
WORN REFLECTIVE MATERIAL	292
MISSING	93

Name: DEVICE_CONDITION, dtype: int64

WEATHER_CONDITION

CLEAR	461336
RAIN	57824
SNOW	23114
CLOUDY/OVERCAST	20434
UNKNOWN	9340
OTHER	1814
FREEZING RAIN/DRIZZLE	1250
SLEET/HAIL	935
FOG/SMOKE/HAZE	921
BLOWING SNOW	343
SEVERE CROSS WIND GATE	114
BLOWING SAND, SOIL, DIRT	4

Name: WEATHER_CONDITION, dtype: int64

LIGHTING_CONDITION

DAYLIGHT	395940
DARKNESS, LIGHTED ROAD	126260
DARKNESS	22850
DUSK	17561
DAWN	9363
UNKNOWN	5455

Name: LIGHTING_CONDITION, dtype: int64

TRAFFICWAY_TYPE

NOT DIVIDED	256091
DIVIDED - W/MEDIAN (NOT RAISED)	114787
ONE-WAY	51042
FOUR WAY	44614
DIVIDED - W/MEDIAN BARRIER	42417
PARKING LOT	22391
OTHER	13789

T-INTERSECTION	8576
CENTER TURN LANE	6297
ALLEY	6148
UNKNOWN	2735
RAMP	1967
UNKNOWN INTERSECTION TYPE	1761
DRIVEWAY	1363
FIVE POINT, OR MORE	979
Y-INTERSECTION	889
TRAFFIC ROUTE	836
NOT REPORTED	477
ROUNDAABOUT	140
L-INTERSECTION	130

Name: TRAFFICWAY_TYPE, dtype: int64

ROADWAY_SURFACE_COND

DRY	437440
WET	89271
UNKNOWN	22479
SNOW OR SLUSH	21727
ICE	5074
OTHER	1266
SAND, MUD, DIRT	172

Name: ROADWAY_SURFACE_COND, dtype: int64

ROAD_DEFECT

NO DEFECTS	494189
UNKNOWN	72406
RUT, HOLES	3453
OTHER	3284
WORN SURFACE	2455
SHOULDER DEFECT	1152
DEBRIS ON ROADWAY	490

Name: ROAD_DEFECT, dtype: int64

PRIM_CONTRIBUTORY_CAUSE

1	327208
0	250221

Name: PRIM_CONTRIBUTORY_CAUSE, dtype: int64

CRASH_HOUR

16	47029
15	46175
17	46149

14	39356
18	37032
13	35432
12	34016
8	32932
11	30545
9	27520
10	26900
7	26500
19	26095
20	20380
21	17939
22	16095
23	13594
6	12104
0	10137
1	8180
2	6865
5	6437
3	5322
4	4695

Name: CRASH_HOUR, dtype: int64

CRASH_DAY_OF_WEEK

6	96922
5	85436
4	83736
7	83672
3	83247
2	78489
1	65927

Name: CRASH_DAY_OF_WEEK, dtype: int64

CRASH_MONTH

10	53673
12	50276
9	49760
5	49063
11	48996
8	47978
1	47108
7	46961
3	46683
6	45740
4	45731
2	45460

Name: CRASH_MONTH, dtype: int64

LATITUDE

41.976201	969
41.900959	600
41.791420	458
41.751461	456
41.880856	324

...

41.889348	1
41.780923	1
41.740613	1
41.860918	1
41.835886	1

Name: LATITUDE, Length: 155350, dtype: int64

LONGITUDE

-87.905309	969
-87.619928	600
-87.580148	458
-87.585972	456
-87.617636	324

...

-87.661895	1
-87.570930	1
-87.665346	1
-87.548734	1
-87.724474	1

Name: LONGITUDE, Length: 155334, dtype: int64

LOCATION

POINT (-87.905309125103 41.976201139024)	969
POINT (-87.619928173678 41.900958919109)	600
POINT (-87.580147768689 41.791420282098)	458
POINT (-87.585971992965 41.751460603167)	456
POINT (-87.617635891755 41.880856047671)	324

...

POINT (-87.663909046208 41.896235950067)	1
POINT (-87.711181809431 41.891937228592)	1
POINT (-87.765246938384 41.778133615169)	1
POINT (-87.618487458568 41.896927006997)	1
POINT (-87.724474013253 41.835886103363)	1

Name: LOCATION, Length: 155401, dtype: int64

- SEX: The 'SEX' column already contains two categories, 'M' and 'F', which represent male

and female. There is no need for cleaning

- AGE: The ‘AGE’ column represents different age values. It appears that the ages are already binned into specific values.
- DRIVER_ACTION: The ‘DRIVER_ACTION’ column contains various driver actions. It seems to have already been cleaned,
- MANEUVER: The ‘MANEUVER’ column represents different driving maneuvers. It appears to be relatively clean
- POSTED_SPEED_LIMIT: The ‘POSTED_SPEED_LIMIT’ column contains numeric values representing posted speed limits
- CRASH HOUR, DAY, and MONTH are all numerical variables
- LONGITUDES and LATITUDES are continuous numerical variables

The following should be reduced to smaller classes for better classification:

- * Driver Vision
- * Physical Condition
- * Device Condition
- * Weather Condition
- * Lighting Condition
- * Trafficway Type
- * Roadway Surface Condition
- * Road Defect
- *

Driver Vision

```
[45]: # check frequency of each unique value in the "DRIVER_VISION"
final_data.DRIVER_VISION.value_counts()
```

```
[45]: NOT OBSCURED          402587
      UNKNOWN             154308
      OTHER                7213
      MOVING VEHICLES      5407
      PARKED VEHICLES      3337
      WINDSHIELD (WATER/ICE) 2459
      BLINDED - SUNLIGHT   1157
      TREES, PLANTS        374
      BUILDINGS            306
      BLINDED - HEADLIGHTS 75
      HILLCREST            65
      EMBANKMENT           63
      BLOWING MATERIALS    53
      SIGNBOARD            25
      Name: DRIVER_VISION, dtype: int64
```

“DRIVER_VISION” column, all categories except “NOT OBSCURED” and “UNKNOWN” can

be considered as different forms of “OBSCURED” vision. Therefore, we will combine all these variations into a single category called “OBSCURED.”

```
[46]: # Define the categories to be grouped as "OBSCURED"
obsured_categories = ['MOVING VEHICLES', 'PARKED VEHICLES', 'WINDSHIELD (WATER/
↳ICE)',

                    'BLINDED - SUNLIGHT', 'TREES, PLANTS', 'BUILDINGS',
                    'BLINDED - HEADLIGHTS', 'HILLCREST', 'EMBANKMENT',
                    'BLOWING MATERIALS', 'SIGNBOARD', 'OTHER']

# Group the categories into "OBSCURED"
final_data.loc[~final_data['DRIVER_VISION'].isin(['NOT OBSCURED', 'UNKNOWN']),
↳'DRIVER_VISION'] = 'OBSCURED'

# Count the occurrences of each category
final_data['DRIVER_VISION'].value_counts()
```

```
[46]: NOT OBSCURED      402587
      UNKNOWN          154308
      OBSCURED          20534
      Name: DRIVER_VISION, dtype: int64
```

Physical Condition

```
[47]: # Count the occurrences of each category
final_data.PHYSICAL_CONDITION.value_counts()
```

```
[47]: NORMAL              498584
      UNKNOWN             64022
      IMPAIRED - ALCOHOL    4392
      FATIGUED/ASLEEP       2382
      REMOVED BY EMS        2067
      OTHER                 1851
      EMOTIONAL             1755
      ILLNESS/FAINTED       940
      HAD BEEN DRINKING     559
      IMPAIRED - DRUGS      503
      IMPAIRED - ALCOHOL AND DRUGS 260
      MEDICATED             114
      Name: PHYSICAL_CONDITION, dtype: int64
```

I will categorize all conditions other than “NORMAL” and “UNKNOWN” as “IMPAIRED” in the “PHYSICAL_CONDITION” column.

```
[48]: # Define the categories to be renamed as "IMPAIRED"
impaired_categories = ['FATIGUED/ASLEEP', 'EMOTIONAL', 'ILLNESS/FAINTED',
↳'ALCOHOL/DRUGS']

# Rename the categories as "IMPAIRED"
```



```
final_data.loc[~final_data['PHYSICAL_CONDITION'].isin(['NORMAL', 'UNKNOWN']),  
↳ 'PHYSICAL_CONDITION'] = 'IMPAIRED'  
# Count the occurrences of each category  
final_data['PHYSICAL_CONDITION'].value_counts()
```

```
[48]: NORMAL          498584  
      UNKNOWN         64022  
      IMPAIRED        14823  
      Name: PHYSICAL_CONDITION, dtype: int64
```

Device Condition

```
[49]: # Count the occurrences of each category  
final_data.DEVICE_CONDITION.value_counts()
```

```
[49]: NO CONTROLS          275828  
      FUNCTIONING PROPERLY  268945  
      UNKNOWN              21966  
      OTHER                4683  
      FUNCTIONING IMPROPERLY 3687  
      NOT FUNCTIONING       1935  
      WORN REFLECTIVE MATERIAL 292  
      MISSING               93  
      Name: DEVICE_CONDITION, dtype: int64
```

I will group DEVICE_CONDITION into NO_CONTROLS, FUNCTIONING, UNKNOWN, and NOT FUNCTIONING.

- We do not know what OTHER means. We will put them as UNKNOWN.
- WORN_RELECTIVE_MATERIAL basically means that the device was functioning properly. We will re-classify them as FUNCTIONING together with FUNCTIONING PROPERLY.
- MISSING means no controls present. We will group them into NO CONTROLS.
- FUNCTIONING IMPROPERLY is as good as NOT FUNCTIONING at all. We will group them together.

```
[50]: # Define the categories to be grouped  
unknown_categories = ['OTHER']  
functioning_categories = ['WORN REFLECTIVE MATERIAL', 'FUNCTIONING PROPERLY']  
no_controls_categories = ['MISSING']  
not_functioning_categories = ['FUNCTIONING IMPROPERLY']  
  
# Group the categories accordingly  
final_data.loc[final_data['DEVICE_CONDITION'].isin(unknown_categories),  
↳ 'DEVICE_CONDITION'] = 'UNKNOWN'  
final_data.loc[final_data['DEVICE_CONDITION'].isin(functioning_categories),  
↳ 'DEVICE_CONDITION'] = 'FUNCTIONING'
```

```

final_data.loc[final_data['DEVICE_CONDITION'].isin(no_controls_categories),
↳ 'DEVICE_CONDITION'] = 'NO CONTROLS'
final_data.loc[final_data['DEVICE_CONDITION'].isin(not_functioning_categories),
↳ 'DEVICE_CONDITION'] = 'NOT FUNCTIONING'

# Count the occurrences of each category
final_data['DEVICE_CONDITION'].value_counts()

```

```

[50]: NO CONTROLS          275921
      FUNCTIONING        269237
      UNKNOWN           26649
      NOT FUNCTIONING     5622
      Name: DEVICE_CONDITION, dtype: int64

```

Weather Condition

```

[51]: # Count the occurrences of each category
      final_data.WEATHER_CONDITION.value_counts()

```

```

[51]: CLEAR                461336
      RAIN                 57824
      SNOW                 23114
      CLOUDY/OVERCAST     20434
      UNKNOWN             9340
      OTHER               1814
      FREEZING RAIN/DRIZZLE 1250
      SLEET/HAIL           935
      FOG/SMOKE/HAZE       921
      BLOWING SNOW         343
      SEVERE CROSS WIND GATE 114
      BLOWING SAND, SOIL, DIRT 4
      Name: WEATHER_CONDITION, dtype: int64

```

```

[52]: # Replacing various categories in the 'WEATHER_CONDITION' column using the apply
      ↳ method and lambda functions.
final_data.WEATHER_CONDITION = final_data.WEATHER_CONDITION.apply(lambda x:
↳ 'SNOW' if x == 'BLOWING SNOW' else x)
final_data.WEATHER_CONDITION = final_data.WEATHER_CONDITION.apply(lambda x:
↳ 'RAIN' if x in ['FREEZING RAIN/DRIZZLE',
↳
↳ 'SLEET/HAIL'] else x)
final_data.WEATHER_CONDITION = final_data.WEATHER_CONDITION.apply(lambda x:
↳ 'OTHER' if x in ['FOG/SMOKE/HAZE',
↳
↳ 'SEVERE CROSS WIND GATE',
↳
↳ 'BLOWING SAND, SOIL, DIRT'] else x)

```

```
# Count the occurrences of each category
final_data.WEATHER_CONDITION.value_counts()
```

```
[52]: CLEAR                461336
      RAIN                 60009
      SNOW                23457
      CLOUDY/OVERCAST     20434
      UNKNOWN             9340
      OTHER               2853
      Name: WEATHER_CONDITION, dtype: int64
```

Lighting Condition

```
[53]: # Count the occurrences of each category
      final_data.LIGHTING_CONDITION.value_counts()
```

```
[53]: DAYLIGHT            395940
      DARKNESS, LIGHTED ROAD 126260
      DARKNESS            22850
      DUSK                17561
      DAWN                9363
      UNKNOWN            5455
      Name: LIGHTING_CONDITION, dtype: int64
```

There is two different instances for DARKNESS. We will combine them to be one.

```
[54]: #combining DARKNESS, LIGHTED ROAD to one
      final_data.LIGHTING_CONDITION = final_data.LIGHTING_CONDITION.apply(lambda x:
      ↪ 'DARKNESS' if x == 'DARKNESS, LIGHTED ROAD' else x)
      # Count the occurrences of each category
      final_data.LIGHTING_CONDITION.value_counts()
```

```
[54]: DAYLIGHT            395940
      DARKNESS            149110
      DUSK                17561
      DAWN                9363
      UNKNOWN            5455
      Name: LIGHTING_CONDITION, dtype: int64
```

Trafficway Type

```
[55]: # Count the occurrences of each category
      final_data.TRAFFICWAY_TYPE.value_counts()
```

```
[55]: NOT DIVIDED          256091
      DIVIDED - W/MEDIAN (NOT RAISED) 114787
      ONE-WAY            51042
      FOUR WAY           44614
```

DIVIDED - W/MEDIAN BARRIER	42417
PARKING LOT	22391
OTHER	13789
T-INTERSECTION	8576
CENTER TURN LANE	6297
ALLEY	6148
UNKNOWN	2735
RAMP	1967
UNKNOWN INTERSECTION TYPE	1761
DRIVEWAY	1363
FIVE POINT, OR MORE	979
Y-INTERSECTION	889
TRAFFIC ROUTE	836
NOT REPORTED	477
ROUNDAABOUT	140
L-INTERSECTION	130

Name: TRAFFICWAY_TYPE, dtype: int64

All classes related to INTERSECTION can be grouped into one. Combine the two DIVIDED variations into one. The rest is all unique.

```
[56]: # Grouping classes and combining them to one
final_data.TRAFFICWAY_TYPE = final_data.TRAFFICWAY_TYPE.apply(lambda x:
    ↪ 'INTERSECTION' if x in ['T-INTERSECTION',
    ↪
    ↪ 'UNKNOWN INTERSECTION TYPE',
    ↪
    ↪ 'Y-INTERSECTION',
    ↪
    ↪ 'L-INTERSECTION'] else x)
final_data.TRAFFICWAY_TYPE = final_data.TRAFFICWAY_TYPE.apply(lambda x:
    ↪ 'DIVIDED' if x in ['DIVIDED - W/MEDIAN (NOT RAISED)',
    ↪
    ↪ 'DIVIDED - W/MEDIAN BARRIER'] else x)
# Count the occurrences of each category
final_data.TRAFFICWAY_TYPE.value_counts()
```

```
[56]: NOT DIVIDED          256091
DIVIDED                  157204
ONE-WAY                  51042
FOUR WAY                 44614
PARKING LOT              22391
OTHER                   13789
INTERSECTION             11356
CENTER TURN LANE         6297
ALLEY                    6148
UNKNOWN                  2735
```

RAMP	1967
DRIVEWAY	1363
FIVE POINT, OR MORE	979
TRAFFIC ROUTE	836
NOT REPORTED	477
ROUNDABOUT	140

Name: TRAFFICWAY_TYPE, dtype: int64

Roadway Surface Condition

```
[57]: # Count the occurrences of each category
final_data.ROADWAY_SURFACE_COND.value_counts()
```

```
[57]: DRY                437440
      WET                89271
      UNKNOWN           22479
      SNOW OR SLUSH     21727
      ICE                5074
      OTHER             1266
      SAND, MUD, DIRT    172
      Name: ROADWAY_SURFACE_COND, dtype: int64
```

I will combine ICE with SNOW OR SLUSH creating SNOW/SLUSH/ICE since they all occur during a snow. SAND, MUD, DIRT can also be comined to OTHER since they all represent a minority group.

```
[58]: # count the occurrences of each category
final_data.ROADWAY_SURFACE_COND.value_counts()
```

```
[58]: DRY                437440
      WET                89271
      UNKNOWN           22479
      SNOW OR SLUSH     21727
      ICE                5074
      OTHER             1266
      SAND, MUD, DIRT    172
      Name: ROADWAY_SURFACE_COND, dtype: int64
```

Road Defect

```
[59]: # count the occurrences of each category
final_data.ROAD_DEFECT.value_counts()
```

```
[59]: NO DEFECTS         494189
      UNKNOWN           72406
      RUT, HOLES        3453
      OTHER              3284
      WORN SURFACE       2455
```

```
SHOULDER DEFECT      1152
DEBRIS ON ROADWAY    490
Name: ROAD_DEFECT, dtype: int64
```

Other than NO DEFECT and UNKNOWN, all others seem to be a variation of DEFECTS. We will group them together as DEFECTS.

```
[60]: # Grouping classes and combining them to one
final_data.ROAD_DEFECT = final_data.ROAD_DEFECT.apply(lambda x: 'DEFECTS' if x
↳not in ['NO DEFECTS', 'UNKNOWN'] else x)
# count the occurrences of each category
final_data.ROAD_DEFECT.value_counts()
```

```
[60]: NO DEFECTS      494189
UNKNOWN      72406
DEFECTS      10834
Name: ROAD_DEFECT, dtype: int64
```

final list of predictor classes.

```
[61]: # provide an overview of the unique values and their frequencies for each
↳column in the DataFrame
for col in final_data.columns:
    print('\n' + col + '\n')
    print(final_data[col].value_counts())
```

SEX

```
M      345523
F      231906
Name: SEX, dtype: int64
```

AGE

```
27.0      17531
25.0      17469
26.0      17259
28.0      17065
24.0      16647
...
110.0       2
108.0       2
109.0       1
104.0       1
107.0       1
Name: AGE, Length: 108, dtype: int64
```

DRIVER_ACTION

NONE	269888
FAILED TO YIELD	69467
OTHER	51970
FOLLOWED TOO CLOSELY	47494
UNKNOWN	38207
IMPROPER TURN	19894
IMPROPER BACKING	18826
IMPROPER LANE CHANGE	17623
DISREGARDED CONTROL DEVICES	12344
IMPROPER PASSING	12328
TOO FAST FOR CONDITIONS	11212
WRONG WAY/SIDE	2334
IMPROPER PARKING	1683
CELL PHONE USE OTHER THAN TEXTING	1094
EVADING POLICE VEHICLE	1010
OVERCORRECTED	993
EMERGENCY VEHICLE ON CALL	623
TEXTING	305
STOPPED SCHOOL BUS	98
LICENSE RESTRICTIONS	36

Name: DRIVER_ACTION, dtype: int64

DRIVER_VISION

NOT OBSCURED	402587
UNKNOWN	154308
OBSCURED	20534

Name: DRIVER_VISION, dtype: int64

PHYSICAL_CONDITION

NORMAL	498584
UNKNOWN	64022
IMPAIRED	14823

Name: PHYSICAL_CONDITION, dtype: int64

MANEUVER

STRAIGHT AHEAD	334081
SLOW/STOP IN TRAFFIC	64937
TURNING LEFT	47831
TURNING RIGHT	24298
BACKING	23304
PASSING/OVERTAKING	14445
CHANGING LANES	13288
OTHER	9067

ENTERING TRAFFIC LANE FROM PARKING	8040
UNKNOWN/NA	5428
STARTING IN TRAFFIC	4890
MERGING	4491
U-TURN	4459
SKIDDING/CONTROL LOSS	3671
AVOIDING VEHICLES/OBJECTS	3501
ENTER FROM DRIVE/ALLEY	2823
LEAVING TRAFFIC LANE TO PARK	2739
SLOW/STOP - LEFT TURN	1867
SLOW/STOP - RIGHT TURN	1168
DRIVING WRONG WAY	951
NEGOTIATING A CURVE	888
SLOW/STOP - LOAD/UNLOAD	883
TURNING ON RED	275
DIVERGING	98
PARKED	5
PARKED IN TRAFFIC LANE	1

Name: MANEUVER, dtype: int64

POSTED_SPEED_LIMIT

30	447733
35	44864
25	28767
20	15932
15	12952
40	7786
10	7127
45	4680
0	4345
5	2104
55	541
50	171
3	136
39	73
9	56
60	35
34	14
32	13
2	13
1	12
33	12
99	10
24	9
7	8
11	7
36	5

65	4
44	2
31	2
63	2
12	2
70	2
23	2
38	2
29	2
49	1
4	1
6	1
26	1

Name: POSTED_SPEED_LIMIT, dtype: int64

TRAFFIC_CONTROL_DEVICE

NO CONTROLS	268293
TRAFFIC SIGNAL	214682
STOP SIGN/FLASHER	74589
UNKNOWN	10019
OTHER	3784
LANE USE MARKING	1415
YIELD	1181
OTHER REG. SIGN	720
OTHER WARNING SIGN	594
RAILROAD CROSSING GATE	467
PEDESTRIAN CROSSING SIGN	386
DELINEATORS	299
POLICE/FLAGMAN	264
FLASHING CONTROL SIGNAL	256
SCHOOL ZONE	190
OTHER RAILROAD CROSSING	155
RR CROSSING SIGN	70
NO PASSING	53
BICYCLE CROSSING SIGN	12

Name: TRAFFIC_CONTROL_DEVICE, dtype: int64

DEVICE_CONDITION

NO CONTROLS	275921
FUNCTIONING	269237
UNKNOWN	26649
NOT FUNCTIONING	5622

Name: DEVICE_CONDITION, dtype: int64

WEATHER_CONDITION

CLEAR	461336
RAIN	60009
SNOW	23457
CLOUDY/OVERCAST	20434
UNKNOWN	9340
OTHER	2853

Name: WEATHER_CONDITION, dtype: int64

LIGHTING_CONDITION

DAYLIGHT	395940
DARKNESS	149110
DUSK	17561
DAWN	9363
UNKNOWN	5455

Name: LIGHTING_CONDITION, dtype: int64

TRAFFICWAY_TYPE

NOT DIVIDED	256091
DIVIDED	157204
ONE-WAY	51042
FOUR WAY	44614
PARKING LOT	22391
OTHER	13789
INTERSECTION	11356
CENTER TURN LANE	6297
ALLEY	6148
UNKNOWN	2735
RAMP	1967
DRIVEWAY	1363
FIVE POINT, OR MORE	979
TRAFFIC ROUTE	836
NOT REPORTED	477
ROUNDAABOUT	140

Name: TRAFFICWAY_TYPE, dtype: int64

ROADWAY_SURFACE_COND

DRY	437440
WET	89271
UNKNOWN	22479
SNOW OR SLUSH	21727
ICE	5074
OTHER	1266
SAND, MUD, DIRT	172

Name: ROADWAY_SURFACE_COND, dtype: int64

ROAD_DEFECT

NO DEFECTS 494189

UNKNOWN 72406

DEFECTS 10834

Name: ROAD_DEFECT, dtype: int64

PRIM_CONTRIBUTORY_CAUSE

1 327208

0 250221

Name: PRIM_CONTRIBUTORY_CAUSE, dtype: int64

CRASH_HOUR

16 47029

15 46175

17 46149

14 39356

18 37032

13 35432

12 34016

8 32932

11 30545

9 27520

10 26900

7 26500

19 26095

20 20380

21 17939

22 16095

23 13594

6 12104

0 10137

1 8180

2 6865

5 6437

3 5322

4 4695

Name: CRASH_HOUR, dtype: int64

CRASH_DAY_OF_WEEK

6 96922

5 85436

4 83736

7 83672

3 83247

```
2    78489
1    65927
Name: CRASH_DAY_OF_WEEK, dtype: int64
```

CRASH_MONTH

```
10    53673
12    50276
9     49760
5     49063
11    48996
8     47978
1     47108
7     46961
3     46683
6     45740
4     45731
2     45460
Name: CRASH_MONTH, dtype: int64
```

LATITUDE

```
41.976201    969
41.900959    600
41.791420    458
41.751461    456
41.880856    324
...
41.889348      1
41.780923      1
41.740613      1
41.860918      1
41.835886      1
Name: LATITUDE, Length: 155350, dtype: int64
```

LONGITUDE

```
-87.905309    969
-87.619928    600
-87.580148    458
-87.585972    456
-87.617636    324
...
-87.661895      1
-87.570930      1
-87.665346      1
-87.548734      1
-87.724474      1
```

Name: LONGITUDE, Length: 155334, dtype: int64

LOCATION

```
POINT (-87.905309125103 41.976201139024)    969
POINT (-87.619928173678 41.900958919109)    600
POINT (-87.580147768689 41.791420282098)    458
POINT (-87.585971992965 41.751460603167)    456
POINT (-87.617635891755 41.880856047671)    324
```

...

```
POINT (-87.663909046208 41.896235950067)    1
POINT (-87.711181809431 41.891937228592)    1
POINT (-87.765246938384 41.778133615169)    1
POINT (-87.618487458568 41.896927006997)    1
POINT (-87.724474013253 41.835886103363)    1
```

Name: LOCATION, Length: 155401, dtype: int64

By grouping similar classes together, I have improved the precision and relevance of the classification process, enabling a more meaningful analysis of the data.

```
[62]: #Checking the shape of our final data
      final_data.shape
```

```
[62]: (577429, 21)
```

```
[63]: # checking the head of our data
      final_data.head()
```

```
[63]:  SEX    AGE      DRIVER_ACTION DRIVER_VISION PHYSICAL_CONDITION \
0    F   20.0  FOLLOWED TOO CLOSELY      UNKNOWN      NORMAL
1    F   53.0              UNKNOWN      UNKNOWN      NORMAL
2    M   22.0    IMPROPER BACKING  NOT OBSCURED    IMPAIRED
3    M   67.0              NONE  NOT OBSCURED    NORMAL
4    M   54.0              NONE  NOT OBSCURED    NORMAL

      MANEUVER  POSTED_SPEED_LIMIT TRAFFIC_CONTROL_DEVICE DEVICE_CONDITION \
0  STRAIGHT AHEAD              30    TRAFFIC SIGNAL    FUNCTIONING
1  STRAIGHT AHEAD              30    TRAFFIC SIGNAL    FUNCTIONING
2        BACKING              30      NO CONTROLS    NO CONTROLS
3  STRAIGHT AHEAD              30      NO CONTROLS    NO CONTROLS
4  STRAIGHT AHEAD              30      NO CONTROLS    NO CONTROLS

      WEATHER_CONDITION  ... TRAFFICWAY_TYPE ROADWAY_SURFACE_COND ROAD_DEFECT \
0          CLEAR  ...    NOT DIVIDED      DRY  NO DEFECTS
1          CLEAR  ...    NOT DIVIDED      DRY  NO DEFECTS
2          CLEAR  ...      ONE-WAY      DRY  NO DEFECTS
3          CLEAR  ...      ONE-WAY      DRY  NO DEFECTS
4          CLEAR  ...      ONE-WAY      DRY  NO DEFECTS
```

	PRIM_CONTRIBUTORY_CAUSE	CRASH_HOUR	CRASH_DAY_OF_WEEK	CRASH_MONTH	\
0	1	23	3	5	
1	1	23	3	5	
2	1	23	3	5	
3	1	23	3	5	
4	1	23	3	5	

	LATITUDE	LONGITUDE	LOCATION
0	41.952691	-87.807413	POINT (-87.807413247555 41.952691362649)
1	41.952691	-87.807413	POINT (-87.807413247555 41.952691362649)
2	41.997837	-87.688814	POINT (-87.688813887189 41.997837266972)
3	41.997837	-87.688814	POINT (-87.688813887189 41.997837266972)
4	41.997837	-87.688814	POINT (-87.688813887189 41.997837266972)

[5 rows x 21 columns]

```
[64]: #convert the values in the 'AGE' column of the DataFrame to strings
final_data.AGE = final_data.AGE.map(lambda x : str(x))
final_data
```

```
[64]:
```

	SEX	AGE	DRIVER_ACTION	DRIVER_VISION	PHYSICAL_CONDITION	\
0	F	20.0	FOLLOWED TOO CLOSELY	UNKNOWN	NORMAL	
1	F	53.0	UNKNOWN	UNKNOWN	NORMAL	
2	M	22.0	IMPROPER BACKING	NOT OBSCURED	IMPAIRED	
3	M	67.0	NONE	NOT OBSCURED	NORMAL	
4	M	54.0	NONE	NOT OBSCURED	NORMAL	
...	
577424	F	39.0	NONE	NOT OBSCURED	NORMAL	
577425	F	24.0	NONE	UNKNOWN	UNKNOWN	
577426	F	36.0	FAILED TO YIELD	NOT OBSCURED	NORMAL	
577427	F	41.0	NONE	NOT OBSCURED	NORMAL	
577428	M	63.0	NONE	NOT OBSCURED	NORMAL	

	MANEUVER	POSTED_SPEED_LIMIT	TRAFFIC_CONTROL_DEVICE	\
0	STRAIGHT AHEAD	30	TRAFFIC SIGNAL	
1	STRAIGHT AHEAD	30	TRAFFIC SIGNAL	
2	BACKING	30	NO CONTROLS	
3	STRAIGHT AHEAD	30	NO CONTROLS	
4	STRAIGHT AHEAD	30	NO CONTROLS	
...	
577424	BACKING	30	NO CONTROLS	
577425	STRAIGHT AHEAD	30	NO CONTROLS	
577426	STRAIGHT AHEAD	30	YIELD	
577427	STRAIGHT AHEAD	30	YIELD	
577428	STRAIGHT AHEAD	30	TRAFFIC SIGNAL	

	DEVICE_CONDITION	WEATHER_CONDITION	...	TRAFFICWAY_TYPE	\
0	FUNCTIONING	CLEAR	...	NOT DIVIDED	
1	FUNCTIONING	CLEAR	...	NOT DIVIDED	
2	NO CONTROLS	CLEAR	...	ONE-WAY	
3	NO CONTROLS	CLEAR	...	ONE-WAY	
4	NO CONTROLS	CLEAR	...	ONE-WAY	
...	
577424	NO CONTROLS	CLEAR	...	NOT DIVIDED	
577425	NO CONTROLS	RAIN	...	NOT DIVIDED	
577426	NO CONTROLS	CLEAR	...	DIVIDED	
577427	NO CONTROLS	CLEAR	...	DIVIDED	
577428	FUNCTIONING	CLEAR	...	NOT DIVIDED	

	ROADWAY_SURFACE_COND	ROAD_DEFECT	PRIM_CONTRIBUTORY_CAUSE	CRASH_HOUR	\
0	DRY	NO DEFECTS		1	23
1	DRY	NO DEFECTS		1	23
2	DRY	NO DEFECTS		1	23
3	DRY	NO DEFECTS		1	23
4	DRY	NO DEFECTS		1	23
...	
577424	DRY	NO DEFECTS		1	17
577425	WET	UNKNOWN		1	19
577426	DRY	NO DEFECTS		0	7
577427	DRY	NO DEFECTS		0	7
577428	DRY	NO DEFECTS		1	16

	CRASH_DAY_OF_WEEK	CRASH_MONTH	LATITUDE	LONGITUDE	\
0	3	5	41.952691	-87.807413	
1	3	5	41.952691	-87.807413	
2	3	5	41.997837	-87.688814	
3	3	5	41.997837	-87.688814	
4	3	5	41.997837	-87.688814	
...	
577424	7	1	41.868979	-87.640629	
577425	4	6	41.835886	-87.724474	
577426	3	1	41.760710	-87.561946	
577427	3	1	41.760710	-87.561946	
577428	1	3	41.975857	-87.708744	

	LOCATION
0	POINT (-87.807413247555 41.952691362649)
1	POINT (-87.807413247555 41.952691362649)
2	POINT (-87.688813887189 41.997837266972)
3	POINT (-87.688813887189 41.997837266972)
4	POINT (-87.688813887189 41.997837266972)
...	...
577424	POINT (-87.640628921351 41.868979049994)

```

577425 POINT (-87.724474013253 41.835886103363)
577426 POINT (-87.561946030143 41.760710194223)
577427 POINT (-87.561946030143 41.760710194223)
577428 POINT (-87.708743641643 41.975856915535)

```

[577429 rows x 21 columns]

8.1 Explanatory Data Analysis(EDA)

Answering the Business questions: ##### Question one *Are there any specific locations or road segments in Chicago city that have a higher frequency of car accidents?*

By combining the accident data with the shapefile using Geopandas, we can overlay the accident points on the map of Chicago. This visualization will provide a spatial representation of the accident locations, allowing us to identify specific areas or patterns where accidents are more prevalent.

```

[65]: #Importing libraries that will help as plot our streatshape
import geopandas as gpd
from shapely.geometry import Point

```

```

[66]: # Import shape file for the streep map of Chicago
street_map = gpd.read_file(r'C:\\Users\\Admin\\Documents\\Iano\\phase 3\\
    ↪project\\dsc-phase-3-project-v2-3\\geo\\
    ↪data\\geo_export_1fec9e49-164d-454e-9b60-f2310808e96f.shx')

```

```

[67]: # Create a location dataframe
location = final_data[['PRIM_CONTRIBUTORY_CAUSE', 'LONGITUDE', 'LATITUDE']][:
    ↪25000]

# Create a GeoDataFrame by assigning the coordinates as a Point geometry
geo_data = location.assign(geometry=gpd.points_from_xy(location.LONGITUDE,
    ↪location.LATITUDE))
geo_data = gpd.GeoDataFrame(geo_data, crs='EPSG:4326')

# Alternatively, you can directly assign the geometry while creating the
    ↪GeoDataFrame
geo_data = gpd.GeoDataFrame(
    location,
    geometry=gpd.points_from_xy(location.LONGITUDE, location.LATITUDE),
    crs='EPSG:4326'
)

```

```

[68]: # Plot the coordinates
fig, ax = plt.subplots(figsize=(15, 15))
street_map.plot(ax=ax, alpha=0.25, color='black')
geo_data[geo_data['PRIM_CONTRIBUTORY_CAUSE'] == 0].plot(ax=ax, markersize=1,
    ↪color='blue', marker='o', label='Unintentional')

```



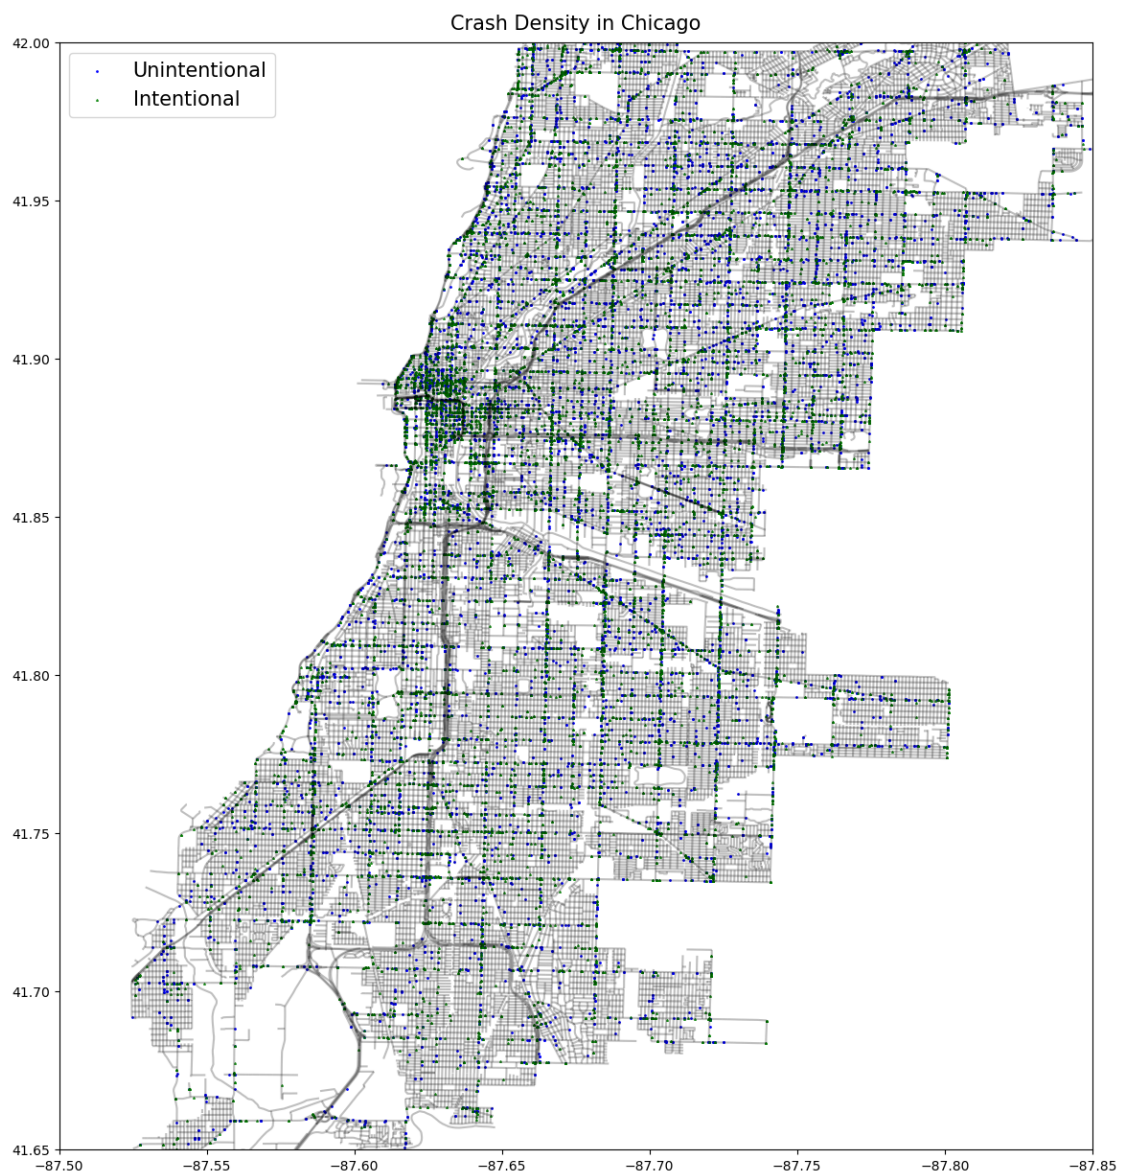
```

geo_data[geo_data['PRIM_CONTRIBUTORY_CAUSE'] == 1].plot(ax=ax, markersize=1,
    color='green', marker='^', label='Intentional')
plt.legend(prop={'size': 15})
plt.title('Crash Density in Chicago', fontsize=15, pad=10)

# Set the desired limits for the x-axis and y-axis
plt.xlim(-87.5, -87.85)
plt.ylim(41.65, 42.0)

# Display the plot
plt.show()
plt.savefig(r'images\street_map.png', bbox_inches='tight');

```



<Figure size 640x480 with 0 Axes>

The map highlights a significant concentration of accidents in the downtown area of Chicago, indicating a higher density of incidents in that region. The predominant color in this area is green, indicating that a majority of the accidents are attributed to intentional actions or driver errors. However, it's important to note that there are also scattered blue plots throughout the map, suggesting a considerable number of accidents that occur unintentionally or present opportunities for improvement in terms of safety measures.

Question Two *What are the contributing factors or characteristics associated with severe car accidents in Chicago city?*

Our focus will be on accidents that were categorized as 'Unintentional' in order to delve deeper into the underlying causes and identify potential areas for improvement. By narrowing our analysis to these specific incidents, we can gain valuable insights into the root causes of unintentional accidents and uncover opportunities for enhancing safety measures and preventing similar occurrences in the future.

```
[69]: # Selecting the relevant columns for analysis
      factors = _
      ↪final_data[['DRIVER_VISION', 'POSTED_SPEED_LIMIT', 'TRAFFIC_CONTROL_DEVICE', 'DEVICE_CONDITION',
                  'WEATHER_CONDITION', 'LIGHTING_CONDITION',
                  'TRAFFICWAY_TYPE', 'ROADWAY_SURFACE_COND',
                  'ROAD_DEFECT', 'PRIM_CONTRIBUTORY_CAUSE']].copy()
```

```
[70]: # Filtering the data for control failures (Unintentional accidents)
      control_failures = factors[factors.PRIM_CONTRIBUTORY_CAUSE == 0].copy()
      # Removing the 'PRIM_CONTRIBUTORY_CAUSE' column as it is no longer needed
      control_failures.drop('PRIM_CONTRIBUTORY_CAUSE', axis=1, inplace=True)
      # Displaying the column names of the control_failures DataFrame
      control_failures.columns
```

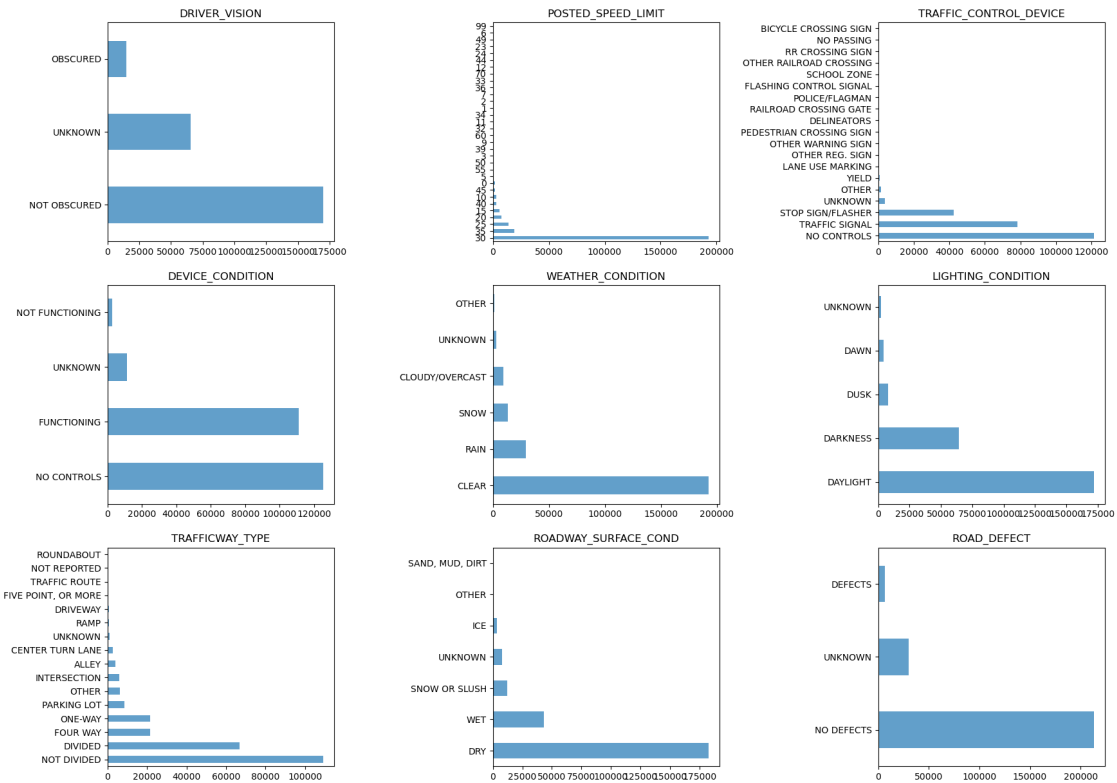
```
[70]: Index(['DRIVER_VISION', 'POSTED_SPEED_LIMIT', 'TRAFFIC_CONTROL_DEVICE',
          'DEVICE_CONDITION', 'WEATHER_CONDITION', 'LIGHTING_CONDITION',
          'TRAFFICWAY_TYPE', 'ROADWAY_SURFACE_COND', 'ROAD_DEFECT'],
         dtype='object')
```

```
[71]: # Creating an empty dictionary to store column-wise value counts
      # Calculating value counts for the current column and storing it in the _
      ↪dictionary
      # Iterating over each column in the control_failures DataFrame
      count_dict = {}
      for col in control_failures.columns:
          count_dict[str(col)] = control_failures[col].value_counts()
      # plotting the predictors
      plt.figure(figsize=(20,15))
      plt.subplots_adjust(wspace=0.7)
```

```

for index, value in enumerate(count_dict):
    ax = plt.subplot(3, 3, index+1)
    chart = pd.DataFrame(count_dict[value])
    chart.plot(ax=ax, kind='barh', legend=False, alpha=0.7)
    ax.set_title(value)
    plt.savefig(r'images\quiz_two.png', bbox_inches='tight');

```



From the above plot we can draw the following conclusions:

- Upon analyzing the contributing factors associated with control failures in unintentional accidents, it is evident that a majority of the accidents occurred when the driver's vision was not obscured. Furthermore, it is notable that these accidents occurred while the drivers were adhering to the posted speed limit, typically set at 30 mph. These findings suggest that factors other than vision or speed might be contributing to control failures in these accidents.
- An important finding from the analysis is that the absence of traffic control devices has been the primary contributing factor to the number of accidents in Chicago. This suggests that increasing the presence of traffic control devices throughout the city could potentially reduce the occurrence of unintentional accidents. This finding is further supported by the Device Condition plot, which indicates a higher count of accidents when there are no traffic control devices in place. Implementing and improving traffic control measures can therefore be an effective strategy to mitigate control failures and enhance road safety in Chicago.

- The analysis indicates that weather condition and lighting condition have relatively minimal impact on the occurrence of accidents. These factors do not show a strong correlation with the number of accidents in Chicago.
- Significant number of accidents occur on roads categorized as “Not Divided” in terms of trafficway type. This suggests that implementing road division measures, such as adding medians or physical barriers, can potentially mitigate the occurrence of accidents. Dividing the roads can enhance traffic management, separate opposing flows of traffic, and reduce the likelihood of collisions, thereby contributing to improved road safety.
- The analysis indicates that the roadway surface condition and road defects have a relatively minimal impact on the occurrence of these accidents. It suggests that the condition of the road surface, such as potholes or uneven pavement, and the presence of road defects, such as cracks or debris, may not be significant contributors to the unintentional accidents in Chicago.

Question Three *Are there any seasonal or temporal patterns in car accidents in Chicago city?*

I will explore whether there are any recurring patterns in car accidents based on hour, months, days of the week, or specific time intervals. I will use both the intentional and unintentional crash data.

```
[72]: # Selecting the columns 'CRASH_HOUR', 'CRASH_DAY_OF_WEEK', and 'CRASH_MONTH'
      ↪ from the 'final_data' DataFrame
      # and creating a new DataFrame called 'crash_time'
      crash_time= final_data[['CRASH_HOUR', 'CRASH_DAY_OF_WEEK', 'CRASH_MONTH']].
      ↪ copy()
      crash_time
```

```
[72]:
```

	CRASH_HOUR	CRASH_DAY_OF_WEEK	CRASH_MONTH
0	23	3	5
1	23	3	5
2	23	3	5
3	23	3	5
4	23	3	5
...
577424	17	7	1
577425	19	4	6
577426	7	3	1
577427	7	3	1
577428	16	1	3

[577429 rows x 3 columns]

```
[73]: # Plot the graphs
      crash_time = final_data[['CRASH_HOUR', 'CRASH_DAY_OF_WEEK', 'CRASH_MONTH']].
      ↪ copy()


      plt.figure(figsize=(20, 15))
      plt.subplots_adjust(wspace=0.7)
```

```

for i, col in enumerate(crash_time.columns):
    ax = plt.subplot(3, 3, i+1)
    chart_2 = pd.DataFrame(crash_time[col].value_counts()).sort_index()
    chart_2.plot(ax=ax, kind='bar', legend=False, alpha=0.7)
    ax.set_title(col)

plt.tight_layout()
plt.show()
plt.savefig(r'images\quiz_three.png', bbox_inches='tight');

```



<Figure size 640x480 with 0 Axes>

From the above plots we can come to conclusion:

- The analysis of the crash time data reveals that a significant number of accidents in Chicago occur between the hours of 14 to 18, which coincides with the peak rush hour traffic. This suggests that the high volume of vehicles during these hours contributes to the increased accident rate. Considering the concentration of accidents in the downtown area during this time frame, it becomes apparent that better traffic management strategies are needed.

To address this issue, it is recommended that the city implements additional measures to facilitate traffic flow and reduce congestion in the downtown area during these peak hours. This can include deploying more traffic management personnel or implementing intelligent transportation systems to optimize traffic signal timings and improve the coordination of traffic movements. By enhancing traffic management during rush hours, the city can mitigate the number of accidents and improve overall road safety in the downtown area.

- The analysis of the crash data by day of the week indicates that there is a slightly higher number of accidents during the weekends compared to other days. However, the difference in accident frequency between weekdays and weekends is not substantial. Therefore, it can be concluded that the crash hour plays a more significant role in determining accident occurrence than the specific day of the week.
- Analyzing the crash data by month reveals some interesting patterns. The number of car accidents in Chicago tends to be higher during the summer months, particularly in June, July, August and September. This can be attributed to various factors such as increased travel and tourism, more outdoor activities, and potentially more congested roads during the summer season.

However, it is important to note that while there may be higher accident rates during certain months, the difference in accident frequency between months is not significant enough to warrant major adjustments in road safety strategies based solely on the crash month.

Question Four Can we build a classification model to predict the primary contributory cause of car accidents? Developing a classification model will enable the CCVSB to categorize accidents into different causes, allowing for a deeper understanding of the factors contributing to each type of accident. This knowledge can inform targeted strategies for prevention.

I will build classifier to help me analyze the chicago car crash

Preparing data

```
[74]: #Importing necessary libraries
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import LinearSVC
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier,
↳ GradientBoostingClassifier
import xgboost as xgb
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import auc
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from sklearn.metrics import mean_squared_error
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.feature_selection import SelectFromModel
from sklearn.decomposition import PCA
import matplotlib.pyplot as plt
%matplotlib inline

[75]: # Create X dataframe by dropping specific columns from final_data
X = final_data.drop(['PRIM_CONTRIBUTORY_CAUSE',
↳ 'CRASH_HOUR', 'CRASH_DAY_OF_WEEK',
↳ 'CRASH_MONTH', 'LONGITUDE', 'LATITUDE', 'LOCATION'], axis=1).
↳ copy()
# Assign PRIM_CONTRIBUTORY_CAUSE column to y variable
y = final_data.PRIM_CONTRIBUTORY_CAUSE

[76]: # Convert predictors into dummies
X = pd.get_dummies(X, drop_first=True)
```

Train-Test Split To evaluate the performance of our model we are splitting the dataset into two parts: a training set and a testing set. we then train a machine learning model on the training set using the 'fit' method

```
[77]: # split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30,
↳ random_state=42)
```

```
[78]: print(X_train.shape) # print (n_train_samples, n_features)
print(X_test.shape) # print (n_test_samples, n_features)
print(y_train.shape) # print (n_train_samples,)
print(y_test.shape) # print (n_test_samples,)
```

```
(404200, 210)
```

```
(173229, 210)
```

```
(404200,)
```

```
(173229,)
```

Scale the data by standard scaler

```
[79]: #Instantiate Standard Scaler
scaler = StandardScaler()

# Fit and transform train and test set
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
[80]: # Create a DataFrame from the scaled training data and display the first few
↳ rows of the scaled training data DataFrame
scaled_data_train = pd.DataFrame(X_train_scaled , columns=X_train.columns)
scaled_data_train.head()
```

```
[80]: POSTED_SPEED_LIMIT    SEX_M  AGE_10.0  AGE_100.0  AGE_101.0  AGE_102.0  \
0          0.141445  0.819479 -0.006092  -0.003853  -0.002724  -0.001573
1          0.141445  0.819479 -0.006092  -0.003853  -0.002724  -0.001573
2          0.141445 -1.220287 -0.006092  -0.003853  -0.002724  -0.001573
3          0.141445  0.819479 -0.006092  -0.003853  -0.002724  -0.001573
4         -1.726058  0.819479 -0.006092  -0.003853  -0.002724  -0.001573
```

```
AGE_103.0  AGE_104.0  AGE_107.0  AGE_108.0  ...  \
0  -0.003146  -0.001573  -0.001573         0.0  ...
1  -0.003146  -0.001573  -0.001573         0.0  ...
2  -0.003146  -0.001573  -0.001573         0.0  ...
3  -0.003146  -0.001573  -0.001573         0.0  ...
4  -0.003146  -0.001573  -0.001573         0.0  ...
```

```
TRAFFICWAY_TYPE_TRAFFIC ROUTE  TRAFFICWAY_TYPE_UNKNOWN  \
0          -0.037711          -0.068977
```


1	-0.037711	-0.068977
2	-0.037711	-0.068977
3	-0.037711	-0.068977
4	-0.037711	-0.068977

	ROADWAY_SURFACE_COND_ICE	ROADWAY_SURFACE_COND_OTHER \
0	-0.094518	-0.047029
1	-0.094518	-0.047029
2	-0.094518	-0.047029
3	-0.094518	-0.047029
4	-0.094518	-0.047029

	ROADWAY_SURFACE_COND_SAND, MUD, DIRT	ROADWAY_SURFACE_COND_SNOW OR SLUSH \
0	-0.017016	-0.197821
1	-0.017016	-0.197821
2	-0.017016	-0.197821
3	-0.017016	-0.197821
4	-0.017016	-0.197821

	ROADWAY_SURFACE_COND_UNKNOWN	ROADWAY_SURFACE_COND_WET \
0	-0.200954	2.333701
1	-0.200954	-0.428504
2	-0.200954	-0.428504
3	-0.200954	-0.428504
4	-0.200954	-0.428504

	ROAD_DEFECT_NO DEFECTS	ROAD_DEFECT_UNKNOWN
0	0.409679	-0.377768
1	0.409679	-0.377768
2	-2.440936	-0.377768
3	0.409679	-0.377768
4	0.409679	-0.377768

[5 rows x 210 columns]

Feature Importance Using Random Forest

```
[81]: # Instantiate and fit the model
rfc = RandomForestClassifier(n_estimators=100)
rfc.fit(X_train_scaled, y_train)
```


```
[81]: RandomForestClassifier()
```

```
[82]: # Get the column names of the features in the training data
labels = list(X_train.columns)
```



```
[83]: # Plot feature importances
n_features = X_train_scaled.shape[1]
plt.figure(figsize=(20,50))
plt.barh(range(n_features), rfc.feature_importances_, align='center')
plt.yticks(np.arange(n_features), labels=labels)
plt.title('Feature Imporance', fontsize=30, pad=5)
plt.xlabel('Feature importance', fontsize=20, labelpad=5)
plt.ylabel('Features', fontsize=20)
plt.tight_layout()
plt.savefig(r'images\feature imp main.png', bbox_inches='tight')
```

0



We can use the feature importance mean to use as a cut-off point for important vs non-important features.

```
[84]: #select features based on their importance scores using the mean value of
      ↪feature importances
      selected_features = X_train.columns[rfc.feature_importances_ > rfc.
      ↪feature_importances_.mean()]
      print(selected_features)
```

```
Index(['POSTED_SPEED_LIMIT', 'SEX_M', 'AGE_19.0', 'AGE_20.0', 'AGE_21.0',
      'AGE_22.0', 'AGE_23.0', 'AGE_24.0', 'AGE_25.0', 'AGE_26.0', 'AGE_27.0',
      'AGE_28.0', 'AGE_29.0', 'AGE_30.0', 'AGE_31.0', 'AGE_32.0', 'AGE_33.0',
      'AGE_34.0', 'AGE_35.0', 'AGE_36.0', 'AGE_37.0', 'AGE_38.0', 'AGE_39.0',
      'AGE_40.0', 'AGE_41.0', 'AGE_42.0', 'AGE_43.0', 'AGE_44.0', 'AGE_45.0',
      'AGE_46.0', 'AGE_47.0', 'AGE_48.0', 'AGE_49.0', 'AGE_50.0', 'AGE_51.0',
      'AGE_52.0', 'AGE_53.0', 'AGE_54.0',
      'DRIVER_ACTION_DISREGARDED CONTROL DEVICES',
      'DRIVER_ACTION_FAILED TO YIELD', 'DRIVER_ACTION_FOLLOWED TOO CLOSELY',
      'DRIVER_ACTION_IMPROPER BACKING', 'DRIVER_ACTION_IMPROPER LANE CHANGE',
      'DRIVER_ACTION_IMPROPER PASSING', 'DRIVER_ACTION_IMPROPER TURN',
      'DRIVER_ACTION_NONE', 'DRIVER_ACTION_OTHER',
      'DRIVER_ACTION_TOO FAST FOR CONDITIONS', 'DRIVER_ACTION_UNKNOWN',
      'DRIVER_VISION_OBSCURED', 'DRIVER_VISION_UNKNOWN',
      'PHYSICAL_CONDITION_NORMAL', 'PHYSICAL_CONDITION_UNKNOWN',
      'MANEUVER_BACKING', 'MANEUVER_PASSING/OVERTAKING',
      'MANEUVER_SLOW/STOP IN TRAFFIC', 'MANEUVER_STRAIGHT AHEAD',
      'MANEUVER_TURNING LEFT', 'TRAFFIC_CONTROL_DEVICE_NO CONTROLS',
      'TRAFFIC_CONTROL_DEVICE_STOP SIGN/FLASHER',
      'TRAFFIC_CONTROL_DEVICE_TRAFFIC SIGNAL', 'DEVICE_CONDITION_NO CONTROLS',
      'DEVICE_CONDITION_UNKNOWN', 'WEATHER_CONDITION_CLOUDY/OVERCAST',
      'WEATHER_CONDITION_RAIN', 'WEATHER_CONDITION_SNOW',
      'LIGHTING_CONDITION_DAWN', 'LIGHTING_CONDITION_DAYLIGHT',
      'LIGHTING_CONDITION_DUSK', 'TRAFFICWAY_TYPE_DIVIDED',
      'TRAFFICWAY_TYPE_FOUR WAY', 'TRAFFICWAY_TYPE_NOT DIVIDED',
      'TRAFFICWAY_TYPE_ONE-WAY', 'TRAFFICWAY_TYPE_OTHER',
      'ROADWAY_SURFACE_COND_SNOW OR SLUSH', 'ROADWAY_SURFACE_COND_UNKNOWN',
      'ROADWAY_SURFACE_COND_WET', 'ROAD_DEFECT_NO DEFECTS',
      'ROAD_DEFECT_UNKNOWN'],
      dtype='object')
```


```
[85]: # filter the feature importance scores based on their values, selecting only
      ↪the scores that are above the mean
      scores = rfc.feature_importances_
      selected_features_scores = scores[rfc.feature_importances_ > rfc.
      ↪feature_importances_.mean()]
```

```
labels_selected = list(selected_features)
```

```
[86]: # Plot feature importances
n_features = len(selected_features)

# Sort the selected features and their scores in ascending order
sorted_indices = np.argsort(selected_features_scores)
sorted_features = np.array(labels_selected)[sorted_indices]
sorted_scores = selected_features_scores[sorted_indices]

plt.figure(figsize=(20, 50))
plt.barh(range(n_features), sorted_scores, align='center')
plt.yticks(np.arange(n_features), labels=sorted_features)
plt.title('Feature Importances', fontsize=30, pad=15)
plt.xlabel('Feature Importance', fontsize=20, labelpad=5)
plt.ylabel('Features', fontsize=20)
plt.tight_layout()
plt.show()
plt.savefig(r'images\ feature imp sort.png', bbox_inches='tight');
```



<Figure size 640x480 with 0 Axes>

By selecting features of importance that are higher than the mean, we can gain valuable insights into the primary causes of accidents. When comparing this list with the plot above, we can observe distinct characteristics related to the accidents. These selected features highlight the factors that contribute significantly to the occurrence of accidents and provide valuable information for understanding the underlying causes. * The feature “DRIVER_ACTION_FAILED_TO_YIELD” stands out as the most important feature, followed by “POSTED_SPEED_LIMIT” and “SEX_M”. These features play a crucial role in predicting the type of accidents associated with specific driver actions, whether they are classified as “INTENTIONAL” or “UNINTENTIONAL”.

- The inclusion of features such as “TRAFFIC_CONTROL_DEVICE”, “WEATHER_CONDITION”, “TRAFFICWAY_TYPE”, and “ROADWAY_SURFACE_COND” in the list of important features indicates that these characteristics are strong predictors of accidents. These features carry significant weight in determining the occurrence and severity of accidents.
- The data suggests that drivers are at a higher risk of accidents during their mid-20s, particularly at the age of 25. As drivers progress into older age groups, the likelihood of accidents gradually decreases. This pattern implies that drivers in their mid-20s may exhibit certain characteristics or behaviors that contribute to a higher accident risk compared to other age groups.

Dimensionality Reduction Dimensionality Reduction: PCA helps in reducing the dimensionality of the dataset by transforming the original variables into a new set of uncorrelated variables called principal components. This is particularly useful when dealing with a high-dimensional dataset with a large number of predictors. By reducing the dimensionality, we can simplify the analysis and visualization of the data.

```
[87]: # Initialize PCA with 80% explained_variance_ratio_

pca = PCA(0.80)

# Fit PCA on the feature data
X_train_transformed = pca.fit_transform(X_train_scaled)
X_test_transformed = pca.transform(X_test_scaled)

# Access the explained variance ratio of the components
explained_variance_ratio = pca.explained_variance_ratio_

# Print the explained variance ratio of each component
for i, ratio in enumerate(explained_variance_ratio):
    print(f"Explained Variance Ratio for Component {i+1}: {ratio}")
```

```
Explained Variance Ratio for Component 1: 0.01759738599332751
Explained Variance Ratio for Component 2: 0.015507362609891797
Explained Variance Ratio for Component 3: 0.010597468054327633
```

Explained Variance Ratio for Component 4: 0.009996300090328961
Explained Variance Ratio for Component 5: 0.009114213338883207
Explained Variance Ratio for Component 6: 0.008664931475144189
Explained Variance Ratio for Component 7: 0.008480502188193968
Explained Variance Ratio for Component 8: 0.008180021416307495
Explained Variance Ratio for Component 9: 0.0076022387478848754
Explained Variance Ratio for Component 10: 0.007351437035559732
Explained Variance Ratio for Component 11: 0.007155961166131913
Explained Variance Ratio for Component 12: 0.007049366249119269
Explained Variance Ratio for Component 13: 0.006798580723836198
Explained Variance Ratio for Component 14: 0.006608684446721035
Explained Variance Ratio for Component 15: 0.006470931948170849
Explained Variance Ratio for Component 16: 0.006302547186115024
Explained Variance Ratio for Component 17: 0.006260202226086329
Explained Variance Ratio for Component 18: 0.006054917779727796
Explained Variance Ratio for Component 19: 0.005931475670390073
Explained Variance Ratio for Component 20: 0.005675580215954186
Explained Variance Ratio for Component 21: 0.005530636676305377
Explained Variance Ratio for Component 22: 0.005480804919432489
Explained Variance Ratio for Component 23: 0.005342101121640716
Explained Variance Ratio for Component 24: 0.005291208250779985
Explained Variance Ratio for Component 25: 0.005232133155635328
Explained Variance Ratio for Component 26: 0.005203310622094665
Explained Variance Ratio for Component 27: 0.0051472856295854425
Explained Variance Ratio for Component 28: 0.005129863168033522
Explained Variance Ratio for Component 29: 0.005090608425382389
Explained Variance Ratio for Component 30: 0.005084075024891853
Explained Variance Ratio for Component 31: 0.005054416025496725
Explained Variance Ratio for Component 32: 0.0050348644818663454
Explained Variance Ratio for Component 33: 0.005026265021318582
Explained Variance Ratio for Component 34: 0.005018724188779451
Explained Variance Ratio for Component 35: 0.00499699631620705
Explained Variance Ratio for Component 36: 0.004993684353292255
Explained Variance Ratio for Component 37: 0.004984535734936757
Explained Variance Ratio for Component 38: 0.004981688862912573
Explained Variance Ratio for Component 39: 0.004977248583412778
Explained Variance Ratio for Component 40: 0.004970047917431993
Explained Variance Ratio for Component 41: 0.004964851461680828
Explained Variance Ratio for Component 42: 0.004963885273531973
Explained Variance Ratio for Component 43: 0.004963221780478917
Explained Variance Ratio for Component 44: 0.004961758201421043
Explained Variance Ratio for Component 45: 0.004953447562172082
Explained Variance Ratio for Component 46: 0.0049512132834056664
Explained Variance Ratio for Component 47: 0.00494995072336451
Explained Variance Ratio for Component 48: 0.004944551817437878
Explained Variance Ratio for Component 49: 0.004943144884821949
Explained Variance Ratio for Component 50: 0.0049410688053069815
Explained Variance Ratio for Component 51: 0.004935828856117358

Explained Variance Ratio for Component 52: 0.004934978180846001
 Explained Variance Ratio for Component 53: 0.004933774517906857
 Explained Variance Ratio for Component 54: 0.004930892459118148
 Explained Variance Ratio for Component 55: 0.004926522020333143
 Explained Variance Ratio for Component 56: 0.004923216167805426
 Explained Variance Ratio for Component 57: 0.004921408343677837
 Explained Variance Ratio for Component 58: 0.0049195920102644616
 Explained Variance Ratio for Component 59: 0.004917458698365174
 Explained Variance Ratio for Component 60: 0.004915299685864923
 Explained Variance Ratio for Component 61: 0.004911487281282694
 Explained Variance Ratio for Component 62: 0.0049102815831642615
 Explained Variance Ratio for Component 63: 0.004906754726989334
 Explained Variance Ratio for Component 64: 0.004906245174743418
 Explained Variance Ratio for Component 65: 0.0049014443834703774
 Explained Variance Ratio for Component 66: 0.004900649743944233
 Explained Variance Ratio for Component 67: 0.004899927200830894
 Explained Variance Ratio for Component 68: 0.004898056452560343
 Explained Variance Ratio for Component 69: 0.004897113113972502
 Explained Variance Ratio for Component 70: 0.004894203955570386
 Explained Variance Ratio for Component 71: 0.004891566534474734
 Explained Variance Ratio for Component 72: 0.004890645523399225
 Explained Variance Ratio for Component 73: 0.0048896490443076255
 Explained Variance Ratio for Component 74: 0.0048884667025115465
 Explained Variance Ratio for Component 75: 0.004887905776029335
 Explained Variance Ratio for Component 76: 0.004886932380364194
 Explained Variance Ratio for Component 77: 0.004885486710843424
 Explained Variance Ratio for Component 78: 0.004881052543684711
 Explained Variance Ratio for Component 79: 0.004879803003791442
 Explained Variance Ratio for Component 80: 0.0048771089445991175
 Explained Variance Ratio for Component 81: 0.0048752003585549805
 Explained Variance Ratio for Component 82: 0.00487163011811799
 Explained Variance Ratio for Component 83: 0.0048702588210326175
 Explained Variance Ratio for Component 84: 0.004866168872650261
 Explained Variance Ratio for Component 85: 0.004865582266418006
 Explained Variance Ratio for Component 86: 0.004863604628950709
 Explained Variance Ratio for Component 87: 0.0048623655286688
 Explained Variance Ratio for Component 88: 0.00486025313090437
 Explained Variance Ratio for Component 89: 0.0048596606271879275
 Explained Variance Ratio for Component 90: 0.004857683175070953
 Explained Variance Ratio for Component 91: 0.004856230052915199
 Explained Variance Ratio for Component 92: 0.00485388332155238
 Explained Variance Ratio for Component 93: 0.004851133980931975
 Explained Variance Ratio for Component 94: 0.00484780168700111
 Explained Variance Ratio for Component 95: 0.004845738660343837
 Explained Variance Ratio for Component 96: 0.004843225355413253
 Explained Variance Ratio for Component 97: 0.004842704657967631
 Explained Variance Ratio for Component 98: 0.0048415519668107265
 Explained Variance Ratio for Component 99: 0.004840444100792638

Explained Variance Ratio for Component 100: 0.004838763527005313
 Explained Variance Ratio for Component 101: 0.004836879053058705
 Explained Variance Ratio for Component 102: 0.004834845415201395
 Explained Variance Ratio for Component 103: 0.004831040321485584
 Explained Variance Ratio for Component 104: 0.004830549283297204
 Explained Variance Ratio for Component 105: 0.004828679894982066
 Explained Variance Ratio for Component 106: 0.004828135127684226
 Explained Variance Ratio for Component 107: 0.004825274761365809
 Explained Variance Ratio for Component 108: 0.004823091612232695
 Explained Variance Ratio for Component 109: 0.004820831032965238
 Explained Variance Ratio for Component 110: 0.004819229435282448
 Explained Variance Ratio for Component 111: 0.004818599760029858
 Explained Variance Ratio for Component 112: 0.004817383935635121
 Explained Variance Ratio for Component 113: 0.004815401970914495
 Explained Variance Ratio for Component 114: 0.004814164308303282
 Explained Variance Ratio for Component 115: 0.004813335708420136
 Explained Variance Ratio for Component 116: 0.004812322774004278
 Explained Variance Ratio for Component 117: 0.00481198478715714
 Explained Variance Ratio for Component 118: 0.004811247890877912
 Explained Variance Ratio for Component 119: 0.004810260002060307
 Explained Variance Ratio for Component 120: 0.004809647009756988
 Explained Variance Ratio for Component 121: 0.004809416698439374
 Explained Variance Ratio for Component 122: 0.004808743031131672
 Explained Variance Ratio for Component 123: 0.004808657146694328
 Explained Variance Ratio for Component 124: 0.004808561536582684
 Explained Variance Ratio for Component 125: 0.00480799590704332
 Explained Variance Ratio for Component 126: 0.004807637780542625
 Explained Variance Ratio for Component 127: 0.004807581751402614
 Explained Variance Ratio for Component 128: 0.004807425472965141
 Explained Variance Ratio for Component 129: 0.004807021511716465
 Explained Variance Ratio for Component 130: 0.004806691471877418
 Explained Variance Ratio for Component 131: 0.0048055099370688985
 Explained Variance Ratio for Component 132: 0.004804521779370376
 Explained Variance Ratio for Component 133: 0.004803858478127186
 Explained Variance Ratio for Component 134: 0.00480333871817911
 Explained Variance Ratio for Component 135: 0.004802747554179915
 Explained Variance Ratio for Component 136: 0.004800999579666919
 Explained Variance Ratio for Component 137: 0.004799437225416373
 Explained Variance Ratio for Component 138: 0.004798111464102305
 Explained Variance Ratio for Component 139: 0.004796941072398516
 Explained Variance Ratio for Component 140: 0.004794138066195897
 Explained Variance Ratio for Component 141: 0.004791924937383968
 Explained Variance Ratio for Component 142: 0.004791419675118967
 Explained Variance Ratio for Component 143: 0.0047841638307921895
 Explained Variance Ratio for Component 144: 0.004781775919658363
 Explained Variance Ratio for Component 145: 0.0047780069357003276
 Explained Variance Ratio for Component 146: 0.004775834657385342
 Explained Variance Ratio for Component 147: 0.00477199063540235

Explained Variance Ratio for Component 148: 0.004768759341162119
Explained Variance Ratio for Component 149: 0.004765983577003185
Explained Variance Ratio for Component 150: 0.004763441096372948

After applying Principal Component Analysis (PCA), the number of predictor columns in our dataset has been reduced from 210 to 182. PCA helps to identify and capture the most important information in the data by creating new variables called principal components. These principal components are linear combinations of the original predictor variables and are chosen in such a way that they explain the maximum amount of variation in the data. By using PCA, we have effectively reduced the dimensionality of the dataset while retaining a significant amount of the information present in the original predictors.

9 Modelling and Evaluation

Creating a classifier for our business problem I will create a classifier that can distinguish between “Unintentional” and “Intentional” accidents based on the available predictors in our dataset, we can use machine learning techniques. The classifier will be trained on the historical accident data, where each accident is labeled as either “Unintentional” or “Intentional” based on the contributing factors.

I will look into the following models:

- * Logistic Regression
- * Decision Tree
- * Random Forest
- * XG Boost

1. Logistic Regression

```
[88]: import os
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.metrics import classification_report, confusion_matrix, roc_curve,
    ↪roc_auc_score
# define a function to valuate the classification model using various metrics
    ↪and generate visualizations.
def evaluate_classification(model, X_train_transformed, X_test_transformed,
    ↪y_train, y_test, classes=None, normalize='true', cmap='Blues_r', label='',
    ↪save_dir='plots'):

    # Create save directory if it doesn't exist
    os.makedirs(save_dir, exist_ok=True)

    # retrieve predictions for train and test data
    y_pred_train = model.predict(X_train_transformed)
```

```

y_pred_test = model.predict(X_test_transformed)

# print training classification report
header = label + " CLASSIFICATION REPORT TRAINING "
dashes = "---" * 20
print(dashes, header, dashes, sep='\n')
print(classification_report(y_train, y_pred_train, target_names=classes))

# calculate confusion matrix for training data
cm_train = confusion_matrix(y_train, y_pred_train)
cm_train_norm = cm_train / cm_train.sum(axis=1)[:, np.newaxis] if normalize
↳ == 'true' else cm_train

# print testing classification report
header_ = label + " CLASSIFICATION REPORT TESTING "
print(dashes, header_, dashes, sep='\n')
print(classification_report(y_test, y_pred_test, target_names=classes))

# calculate confusion matrix for testing data
cm_test = confusion_matrix(y_test, y_pred_test)
cm_test_norm = cm_test / cm_test.sum(axis=1)[:, np.newaxis] if normalize ==
↳ 'true' else cm_test

# Create a combined figure for training and testing plots
fig, axes = plt.subplots(figsize=(12, 4), ncols=4)

# plot confusion matrix for training data
sns.heatmap(cm_train_norm, annot=True, fmt='.2f', cmap=cmap, ax=axes[0])
axes[0].set(title='Confusion Matrix Training', xlabel='Predicted Labels',
↳ ylabel='True Labels')

# plot ROC curve for training data
fpr_train, tpr_train, _ = roc_curve(y_train, model.
↳ predict_proba(X_train_transformed)[: , 1])
roc_auc_train = roc_auc_score(y_train, model.
↳ predict_proba(X_train_transformed)[: , 1])
axes[1].plot(fpr_train, tpr_train, label=f'AUC = {roc_auc_train:.2f}')
axes[1].plot([0, 1], [0, 1], ls=':')
axes[1].set(xlabel='False Positive Rate', ylabel='True Positive Rate',
            title='Receiver Operating Characteristic Training')
axes[1].legend(loc='lower right')

# plot confusion matrix for testing data
sns.heatmap(cm_test_norm, annot=True, fmt='.2f', cmap=cmap, ax=axes[2])
axes[2].set(title='Confusion Matrix Testing', xlabel='Predicted Labels',
↳ ylabel='True Labels')

```

```

# plot ROC curve for testing data
fpr_test, tpr_test, _ = roc_curve(y_test, model.
↳predict_proba(X_test_transformed)[: , 1])
roc_auc_test = roc_auc_score(y_test, model.
↳predict_proba(X_test_transformed)[: , 1])
axes[3].plot(fpr_test, tpr_test, label=f'AUC = {roc_auc_test:.2f}')
axes[3].plot([0, 1], [0, 1], ls=':')
axes[3].set(xlabel='False Positive Rate', ylabel='True Positive Rate',
            title='Receiver Operating Characteristic Testing')
axes[3].legend(loc='lower right')

# Adjust spacing between subplots
plt.tight_layout(pad=2.0)

# Save combined plots
plt.savefig(os.path.join(save_dir, 'combined_plots.png'))

plt.show()

```

```

[89]: from sklearn.linear_model import LogisticRegression
# Initialize the logistic regression model
logreg = LogisticRegression()
# Train the model
logreg.fit(X_train_transformed, y_train)

```

```
[89]: LogisticRegression()
```

```

[90]: # Make predictions on the test set
y_pred = logreg.predict(X_test_transformed)

# Calculate the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Generate a classification report
print(classification_report(y_test, y_pred))

```

Accuracy: 0.6889204463455888

	precision	recall	f1-score	support
0	0.69	0.51	0.59	74958
1	0.69	0.83	0.75	98271
accuracy			0.69	173229
macro avg	0.69	0.67	0.67	173229
weighted avg	0.69	0.69	0.68	173229


```
[91]: # classification report using function
evaluate_classification(logreg,X_train_transformed, X_test_transformed,
↪y_train, y_test, label = "Logistic Regression")
```

Logistic Regression CLASSIFICATION REPORT TRAINING

	precision	recall	f1-score	support
0	0.69	0.51	0.58	175263
1	0.69	0.83	0.75	228937
accuracy			0.69	404200
macro avg	0.69	0.67	0.67	404200
weighted avg	0.69	0.69	0.68	404200

Logistic Regression CLASSIFICATION REPORT TESTING

	precision	recall	f1-score	support
0	0.69	0.51	0.59	74958
1	0.69	0.83	0.75	98271
accuracy			0.69	173229
macro avg	0.69	0.67	0.67	173229
weighted avg	0.69	0.69	0.68	173229



The logistic regression model shows moderate performance on both the training and testing sets. It achieves an accuracy of 0.69 on both sets. The precision, recall, and F1-score for class 0 are 0.69, 0.51, and 0.59, respectively, indicating that the model performs moderately well in predicting instances of class 0. Similarly, for class 1, the precision, recall, and F1-score are 0.69, 0.83, and 0.75, respectively, suggesting that the model performs relatively well in predicting instances of class

1.

Decision Tree Classifier Our initial model for classification will be a Decision Tree Classifier with a tree depth of 3. This means that the decision tree will have a maximum depth of 3 levels, allowing it to make decisions based on three predictor variables at each step. The decision tree algorithm uses a tree-like structure to classify the data based on the features or predictors. By limiting the tree depth to 3, we aim to strike a balance between model complexity and interpretability. This base model will serve as a starting point for further analysis and model improvement.

```
[93]: # Initializes a DecisionTreeClassifier
tree_clf= DecisionTreeClassifier(criterion='gini', max_depth=3)
```

```
[94]: #Fit the model
tree_clf.fit(X_train_transformed, y_train)
```

```
[94]: DecisionTreeClassifier(max_depth=3)
```

```
[ ]: from sklearn import tree
import graphviz
# plot the tree
tree.plot_tree(tree_clf)
```

```
[95]: # Create an array to make predictions for train and test data
y_pred_train = tree_clf.predict(X_train_transformed)
y_pred_test = tree_clf.predict(X_test_transformed)
```

```
[96]: # Calculate accuracy
train_acc = accuracy_score(y_train,y_pred_train) * 100
test_acc = accuracy_score(y_test, y_pred_test) * 100
print('Train accuracy is :{0}'.format(train_acc))
print('Test accuracy is :{0}'.format(test_acc))

# Check the AUC for predictions
fpr, tpr, thresholds = roc_curve(y_test, y_pred_test)
roc_auc = auc(fpr, tpr)
print('\nAUC is :{0}'.format(round(roc_auc, 2)))

# Create and print a confusion matrix
print('\nConfusion Matrix')
print('-----')
pd.crosstab(y_test, y_pred_test, rownames=['True'], colnames=['Predicted'],
↪ margins=True)
```

```
Train accuracy is :64.33077684314695
Test accuracy is :64.37317077394663
```

```
AUC is :0.61
```

Confusion Matrix

```
[96]: Predicted      0      1     All
      True
      0      27653   47305   74958
      1      14411   83860   98271
      All      42064  131165  173229
```

The close similarity between the train and test accuracy values indicates that the model is performing well and is likely to generalize well to unseen data. The model is not overfitting or underfitting the training data, as it achieves similar performance on both the training and testing datasets.

```
[97]: # Call the evaluate_classification function with the desired parameters
      evaluate_classification(tree_clf, X_train_transformed, X_test_transformed,
      ↪ y_train, y_test, label='Decision Tree', save_dir='images')
```

Decision Tree CLASSIFICATION REPORT TRAINING

```
-----
              precision    recall  f1-score   support

         0         0.66      0.37      0.47      175263
         1         0.64      0.85      0.73      228937

    accuracy                   0.64      404200
   macro avg         0.65      0.61      0.60      404200
  weighted avg         0.65      0.64      0.62      404200

-----
```


Decision Tree CLASSIFICATION REPORT TESTING

```
-----
              precision    recall  f1-score   support

         0         0.66      0.37      0.47       74958
         1         0.64      0.85      0.73       98271

    accuracy                   0.64      173229
   macro avg         0.65      0.61      0.60      173229
  weighted avg         0.65      0.64      0.62      173229

-----
```



Training Performance: The model achieved an accuracy of 0.64 on the training set, correctly classifying 64% of the instances in the training data. The precision and recall for class 0 are lower compared to class 1, indicating that the model struggles more in correctly identifying instances of class 0. The weighted average F1-score is 0.62, suggesting a moderate overall performance on the training set.

Testing Performance: On the testing set, the model achieved an accuracy of 0.65, correctly classifying 65% of the instances. Similar to the training set, the precision and recall for class 0 are lower compared to class 1. The weighted average F1-score is 0.62, indicating a moderate overall performance on the testing set.

Considering these observations, we can conclude that the Decision Tree model's performance is moderate. It shows a similar performance on both the training and testing sets, but with lower precision and recall for class 0. This suggests that the model might struggle in accurately identifying instances of class 0, potentially leading to a higher number of false negatives for this class.

Based on the ROC curve, the model's performance in predicting the classes is moderate. The AUC value of 0.67 indicates that there is a 67% chance that the model will correctly classify each target variable. This means that the model has some predictive power, but it is not highly accurate. There is still room for improvement in achieving more accurate predictions.

Re-grow the tree using entropy

```
[98]: #Instantiate the model
dtc_entropy = DecisionTreeClassifier(criterion='entropy')

#Fit the model
dtc_entropy.fit(X_train_transformed, y_train)

# Make predictions for train and test data
y_pred_train_dtc1 = dtc_entropy.predict(X_train_transformed)
y_pred_test_dtc1 = dtc_entropy.predict(X_test_transformed)

# Calculate accuracy
train_acc1 = accuracy_score(y_train, y_pred_train_dtc1) * 100
test_acc1 = accuracy_score(y_test, y_pred_test_dtc1) * 100
print('Train accuracy is :{0}'.format(train_acc1))
```



```

print('Test accuracy is :{0}'.format(test_acc1))

# Check the AUC for predictions
fpr1, tpr1, thresholds1 = roc_curve(y_test, y_pred_test)
roc_auc1 = auc(fpr1, tpr1)
print('\nAUC is :{0}'.format(round(roc_auc1, 2)))

# Create and print a confusion matrix
print('\nConfusion Matrix')
print('-----')
pd.crosstab(y_test, y_pred_test, rownames=['True'], colnames=['Predicted'],
            margins=True)

```

Train accuracy is :89.4188520534389

Test accuracy is :62.91209901344463

AUC is :0.61

Confusion Matrix

```

[98]: Predicted      0      1     All
      True
      0         27653   47305   74958
      1         14411   83860   98271
      All        42064  131165  173229

```

Despite using the entropy criterion for re-growing the decision tree, it did not significantly improve the model's performance.

Training Accuracy: The model achieved a training accuracy of approximately 89.42%, meaning it correctly classified 89.42% of the instances in the training data.

Test Accuracy: The model achieved a test accuracy of approximately 62.78%, indicating that it correctly classified 62.78% of the instances in the test data. Based on these results, we can conclude that the model has a relatively high accuracy on the training set (89.42%), but the accuracy drops significantly on the test set (62.78%). This indicates that the model may be overfitting the training data and is not generalizing well to unseen data. Additionally, the AUC of 0.61 suggests that the model's predictive performance is only slightly better

Random Forest

```

[103]: # Instantiate and fit the model
rf = RandomForestClassifier(n_estimators=100, max_depth= 5)
rf.fit(X_train_transformed, y_train)

```

```

[103]: RandomForestClassifier(max_depth=5)

```


```
[104]: # Call the pred_score function with Random forest classifier
evaluate_classification(rf,X_train_transformed, X_test_transformed, y_train,
→y_test, label = 'Random Forest')
```

Random Forest CLASSIFICATION REPORT TRAINING

	precision	recall	f1-score	support
0	0.75	0.33	0.46	175263
1	0.64	0.92	0.75	228937
accuracy			0.66	404200
macro avg	0.70	0.62	0.61	404200
weighted avg	0.69	0.66	0.63	404200

Random Forest CLASSIFICATION REPORT TESTING

	precision	recall	f1-score	support
0	0.75	0.33	0.46	74958
1	0.64	0.91	0.75	98271
accuracy			0.66	173229
macro avg	0.69	0.62	0.61	173229
weighted avg	0.69	0.66	0.63	173229



Training Performance: The model achieved an accuracy of 0.66 on the training set, correctly classifying 66% of the instances in the training data. The precision and recall for class 0 are lower compared to class 1, indicating that the model struggles more in correctly identifying instances of class 0. The weighted average F1-score is 0.64, suggesting a moderate overall performance on the training set.

Testing Performance: On the testing set, the model achieved an accuracy of 0.67, correctly clas-

sifying 67% of the instances. Similar to the training set, the precision and recall for class 0 are lower compared to class 1. The weighted average F1-score is 0.64, indicating a moderate overall performance on the testing set. ##### Feature Importance with Random Forest

```
[105]: # Plot feature importances
n_features = X_train_transformed.shape[1]
plt.figure(figsize=(20,50))
plt.barh(range(n_features), rf.feature_importances_, align='center')
plt.yticks(np.arange(n_features))
plt.title('Feature Imporance', fontsize=30, pad=5)
plt.xlabel('Feature importance', fontsize=20, labelpad=5)
plt.ylabel('Features', fontsize=20)
plt.tight_layout()
```



Selecting best features

```
[106]: # Print the gini importance of each feature
for feature in zip(range(n_features), rf.feature_importances_):
    print(feature)
```

```
(0, 0.009909129063606808)
(1, 0.019473173019763073)
(2, 0.027333510101412987)
(3, 0.04705559875887393)
(4, 0.04750618233288741)
(5, 0.010914307799327307)
(6, 0.04295296992492601)
(7, 0.23890173420552221)
(8, 0.07763848586295058)
(9, 0.03708139313789361)
(10, 0.04531746807218481)
(11, 0.037575185326788614)
(12, 0.02906859509215385)
(13, 0.11439461044444564)
(14, 0.030188766204037012)
(15, 0.01963423469666918)
(16, 0.03089025851699974)
(17, 0.027577487346027788)
(18, 0.012548294258852394)
(19, 0.008936302055284227)
(20, 0.001882419826991263)
(21, 0.007390007891176465)
(22, 0.002615269629112092)
(23, 0.005004412272370161)
(24, 0.0041863408933955305)
(25, 0.00750771693347986)
(26, 0.0019185135165063127)
(27, 0.0032905767352200193)
(28, 0.004096715864799587)
(29, 0.004043053121756596)
(30, 0.0013743084022243015)
(31, 0.00022483156943289742)
(32, 0.0011668536626883313)
(33, 0.00070535917588184)
(34, 0.00010372936505759594)
(35, 7.465542895600672e-05)
(36, 0.0002438961567870104)
(37, 0.0001311248124442541)
(38, 0.00010063586911560886)
```


(39, 0.00011666751567842485)
(40, 1.180085010752859e-05)
(41, 0.00016545345785593007)
(42, 1.9971274863605582e-05)
(43, 0.004582682592077317)
(44, 2.183830920422099e-05)
(45, 0.0008472004062763126)
(46, 0.00040053425895694695)
(47, 6.200505185455757e-05)
(48, 0.0007001316059092725)
(49, 0.00018072301401184608)
(50, 3.889409711715324e-05)
(51, 1.8716149363083166e-05)
(52, 0.00012601956836424185)
(53, 8.935506093063257e-06)
(54, 3.315945757949523e-05)
(55, 2.9463871401768708e-05)
(56, 1.0319797146789908e-05)
(57, 0.0)
(58, 3.974125931887923e-05)
(59, 1.7304065898946597e-05)
(60, 2.6769416207136275e-05)
(61, 6.909402026259748e-05)
(62, 5.086686237300013e-05)
(63, 9.47111134775619e-06)
(64, 4.189203619867509e-05)
(65, 3.566043857373106e-05)
(66, 1.761868409938404e-05)
(67, 2.777742282047817e-05)
(68, 2.70906269160976e-05)
(69, 9.612351410087213e-05)
(70, 6.740598133651988e-05)
(71, 6.617838021272627e-05)
(72, 0.00013237451947405109)
(73, 4.8769492997623234e-05)
(74, 4.601125497920232e-05)
(75, 8.793510598580671e-05)
(76, 9.712112465977459e-05)
(77, 5.4886520017431936e-05)
(78, 5.244107197973937e-06)
(79, 1.671529941086121e-05)
(80, 0.00012469543332713136)
(81, 0.00022166921402831834)
(82, 8.320036042763662e-05)
(83, 6.381550631521416e-05)
(84, 4.4536813298719266e-05)
(85, 6.942646691625238e-05)
(86, 7.957948359652276e-05)

(87, 7.045687620023627e-05)
(88, 4.6114865506431245e-06)
(89, 6.092195653894181e-05)
(90, 5.419728592229597e-05)
(91, 0.00031153667430789463)
(92, 0.0003946290038093044)
(93, 6.681419832372108e-05)
(94, 4.6912444490092594e-05)
(95, 0.0)
(96, 2.182398384971697e-05)
(97, 0.0001364642670678597)
(98, 2.9278923343602613e-05)
(99, 5.901411819915063e-05)
(100, 0.0)
(101, 0.0011993039850138184)
(102, 0.00020511344673298666)
(103, 0.0002151132163975628)
(104, 0.00018902638178900136)
(105, 7.070550017455514e-05)
(106, 1.3693719801371357e-05)
(107, 2.745089521176024e-05)
(108, 7.627414387628127e-05)
(109, 0.00011985960514030368)
(110, 5.872256113364944e-05)
(111, 0.006549595675669344)
(112, 0.0005439576273019015)
(113, 0.00010100785085417666)
(114, 1.358753933314801e-05)
(115, 0.00019349238117195733)
(116, 0.0005440035885406804)
(117, 0.0018011184519508192)
(118, 0.00025110134866819983)
(119, 5.1327248934790545e-05)
(120, 0.0008341709778823643)
(121, 0.0001670743122654332)
(122, 0.0002657756594772891)
(123, 0.00014373538019484962)
(124, 0.0009684281905041266)
(125, 0.0007202483806397984)
(126, 0.00030695408255587383)
(127, 0.0002891510838101942)
(128, 0.0006688711513953495)
(129, 0.0007698708801298618)
(130, 0.0006527894336320089)
(131, 0.00010954236981803603)
(132, 0.000870504511501571)
(133, 4.572552843044943e-05)
(134, 0.00019483493725845047)

```
(135, 0.00012546913229797294)
(136, 0.0003345888126522472)
(137, 0.000204990656063473)
(138, 7.958555042435379e-05)
(139, 3.087980465095568e-05)
(140, 0.0007958346438250888)
(141, 0.0004055509767604009)
(142, 0.00014543592822215741)
(143, 0.0005416335520433178)
(144, 0.0010411542867407185)
(145, 0.0009916482088025821)
(146, 0.0025788632824936563)
(147, 0.0006229601247830859)
(148, 0.00043471195152994534)
(149, 0.0010763011687510264)
```

Feature importance helps us understand which features are more relevant or impactful in making accurate predictions. Features with higher importance have a stronger influence on the model's decision-making process. By considering feature importance, we can prioritize or filter out features that have a relatively low importance score. This filtering process allows us to focus on the most informative features and potentially improve the model's performance by reducing noise or redundancy in the data.

```
[107]: # Box plot
sns.set_theme(style="whitegrid")
tips = sns.load_dataset("tips")
ax = sns.boxplot(x=rf.feature_importances_)
```

```
[108]: print(rf.feature_importances_.max())
print(rf.feature_importances_.min())
```

```
0.23890173420552221
0.0
```

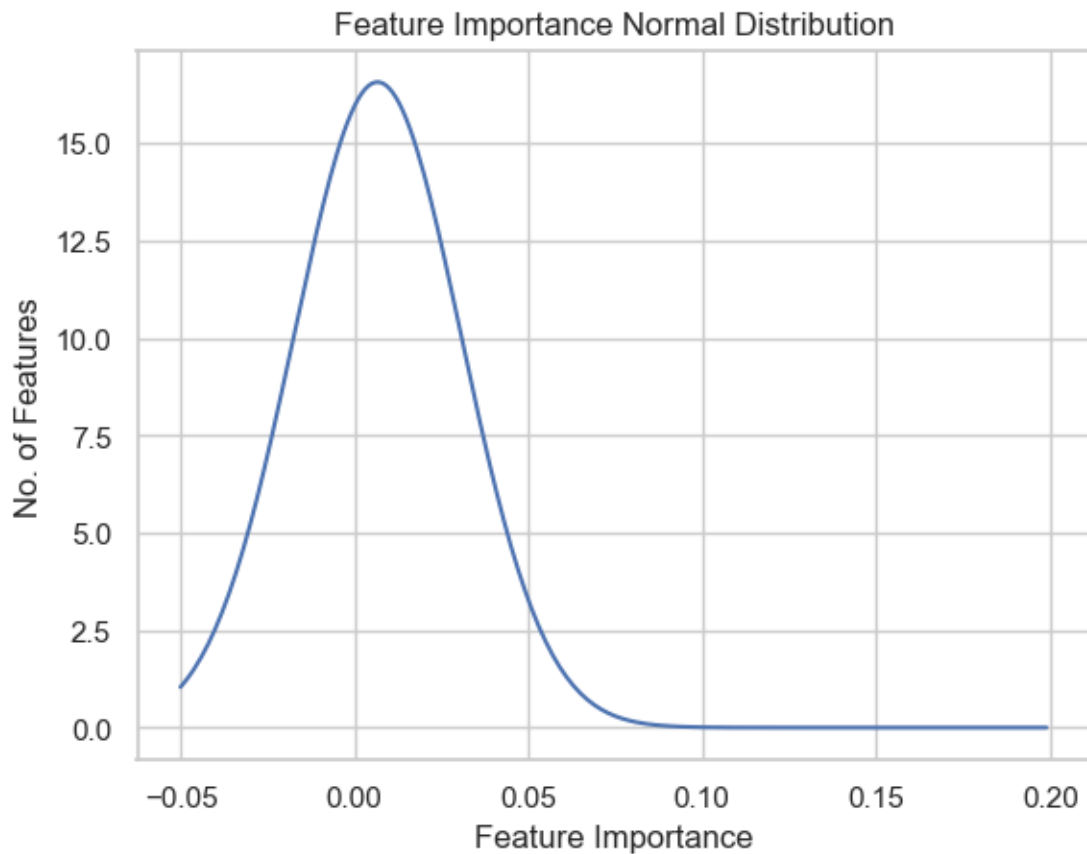
```
[109]: # Check the normal distribution of feature scores
from scipy.stats import norm
import statistics

# Plot between -0.05 and 0.2 with .001 steps.
x_axis = np.arange(-0.05,0.2,.001)

# Calculate mean and standard deviation
mean = statistics.mean(rf.feature_importances_)
sd = statistics.stdev(rf.feature_importances_)

plt.plot(x_axis, norm.pdf(x_axis, mean, sd))
plt.title('Feature Importance Normal Distribution')
plt.ylabel('No. of Features')
plt.xlabel('Feature Importance')
```

```
plt.show();
```



Features between 0.00 and 0.15 are the most relevant

```
[110]: # creates a new list called thresh_1 by filtering the elements from the rf.  
       ↪ feature_importances_ list that are greater than 0.01.
```

```
thresh_1 = [x for x in rf.feature_importances_ if x > 0.01]  
len(thresh_1)
```

```
[110]: 18
```

```
[111]: # Instantiate and fit the model  
sfm = SelectFromModel(rf, threshold=0.01)  
  
sfm.fit(X_train_transformed, y_train)
```

```
[111]: SelectFromModel(estimator=RandomForestClassifier(max_depth=5), threshold=0.01)
```

```
[112]: # Transform the data to create a new dataset containing only the most important  
       ↪ features
```

```

X_best_train = sfm.transform(X_train_transformed)
X_best_test = sfm.transform(X_test_transformed)
# Instantiate and fit the model
rf_best = RandomForestClassifier(n_estimators=100, max_depth= 5)
rf_best.fit(X_best_train, y_train)

```

```
[112]: RandomForestClassifier(max_depth=5)
```

```

[113]: # Make predictions on train and test data
y_pred_train_rfb = rf_best.predict(X_best_train)
y_pred_test_rfb = rf_best.predict(X_best_test)

# Calculate accuracy
train_acc_rfb = accuracy_score(y_train,y_pred_train_rfb) * 100
test_acc_rfb = accuracy_score(y_test, y_pred_test_rfb) * 100
print('Train accuracy is :{0}'.format(train_acc_rfb))
print('Test accuracy is :{0}'.format(test_acc_rfb))

# Check the AUC for predictions
roc_auc_rfb = roc_auc_score(y_test, y_pred_test_rfb)
print('\nAUC is :{0}'.format(round(roc_auc_rfb, 2)))

# Create and print a confusion matrix
print('\nConfusion Matrix')
print('-----')
print(pd.crosstab(y_test, y_pred_test_rfb, rownames=['True'],
    ↪colnames=['Predicted'], margins=True))

# Classification report
print('\nClassification Report')
print('-----')
print(classification_report(y_test, y_pred_test_rfb))

```

Train accuracy is :66.82582879762494

Test accuracy is :66.8138706567607

AUC is :0.63

Confusion Matrix

```

-----
Predicted      0      1     All
True
0      28058   46900   74958
1      10588   87683   98271
All     38646  134583  173229

```

Classification Report

	precision	recall	f1-score	support
0	0.73	0.37	0.49	74958
1	0.65	0.89	0.75	98271
accuracy			0.67	173229
macro avg	0.69	0.63	0.62	173229
weighted avg	0.68	0.67	0.64	173229

Training Performance: The model achieved an accuracy of 0.66 on the training set, correctly classifying 66% of the instances. The precision and recall for class 0 are lower compared to class 1, indicating that the model struggles more in correctly identifying instances of class 0. The weighted average F1-score is 0.64, suggesting a moderate overall performance on the training set.

Testing Performance: On the testing set, the model achieved an accuracy of 0.66, correctly classifying 66% of the instances. Similar to the training set, the precision and recall for class 0 are lower compared to class 1. The weighted average F1-score is 0.64, indicating a moderate overall performance on the testing set.

```
[114]: # Define function for X_best datasets

def pred_score_best(clf):
    # Make predictions on train and test data
    y_pred_train = clf.predict(X_best_train)
    y_pred_test = clf.predict(X_best_test)

    # Calculate accuracy
    train_acc = accuracy_score(y_train, y_pred_train) * 100
    test_acc = accuracy_score(y_test, y_pred_test) * 100
    print('Train accuracy is :{0}'.format(train_acc))
    print('Test accuracy is :{0}'.format(test_acc))

    # Check the AUC for predictions
    roc_auc = roc_auc_score(y_test, y_pred_test)
    print('\nAUC is :{0}'.format(round(roc_auc, 2)))

    # Create and print a confusion matrix
    print('\nConfusion Matrix')
    print('-----')
    print(pd.crosstab(y_test, y_pred_test, rownames=['True'],
    ↪ colnames=['Predicted'], margins=True))

    # Classification report
    print('\nClassification Report')
    print('-----')
    print(classification_report(y_test, y_pred_test))
```

XG Boost

```
[115]: # Instantiate and fit the model
```

```
xg = xgb.XGBClassifier()  
xg.fit(X_best_train, y_train)
```

```
[115]: XGBClassifier(base_score=None, booster=None, callbacks=None,  
                  colsample_bylevel=None, colsample_bynode=None,  
                  colsample_bytree=None, early_stopping_rounds=None,  
                  enable_categorical=False, eval_metric=None, feature_types=None,  
                  gamma=None, gpu_id=None, grow_policy=None, importance_type=None,  
                  interaction_constraints=None, learning_rate=None, max_bin=None,  
                  max_cat_threshold=None, max_cat_to_onehot=None,  
                  max_delta_step=None, max_depth=None, max_leaves=None,  
                  min_child_weight=None, missing=nan, monotone_constraints=None,  
                  n_estimators=100, n_jobs=None, num_parallel_tree=None,  
                  predictor=None, random_state=None, ...)
```

```
[116]: pred_score_best(xg)
```

Train accuracy is :71.6575952498763

Test accuracy is :70.10258097662631

AUC is :0.68

Confusion Matrix

```
-----  
Predicted      0      1    All  
True  
0      41170   33788   74958  
1      18003   80268   98271  
All     59173  114056  173229
```

Classification Report

```
-----  
              precision    recall  f1-score   support  
  
0               0.70       0.55       0.61       74958  
1               0.70       0.82       0.76       98271  
  
accuracy                0.70       0.68       0.69       173229  
macro avg              0.70       0.68       0.68       173229  
weighted avg           0.70       0.70       0.69       173229
```

```
[133]: #Define the confusion matrix values
```

```
confusion_matrix_values = np.array([[41170, 33788], [18003, 80268]])
```

```

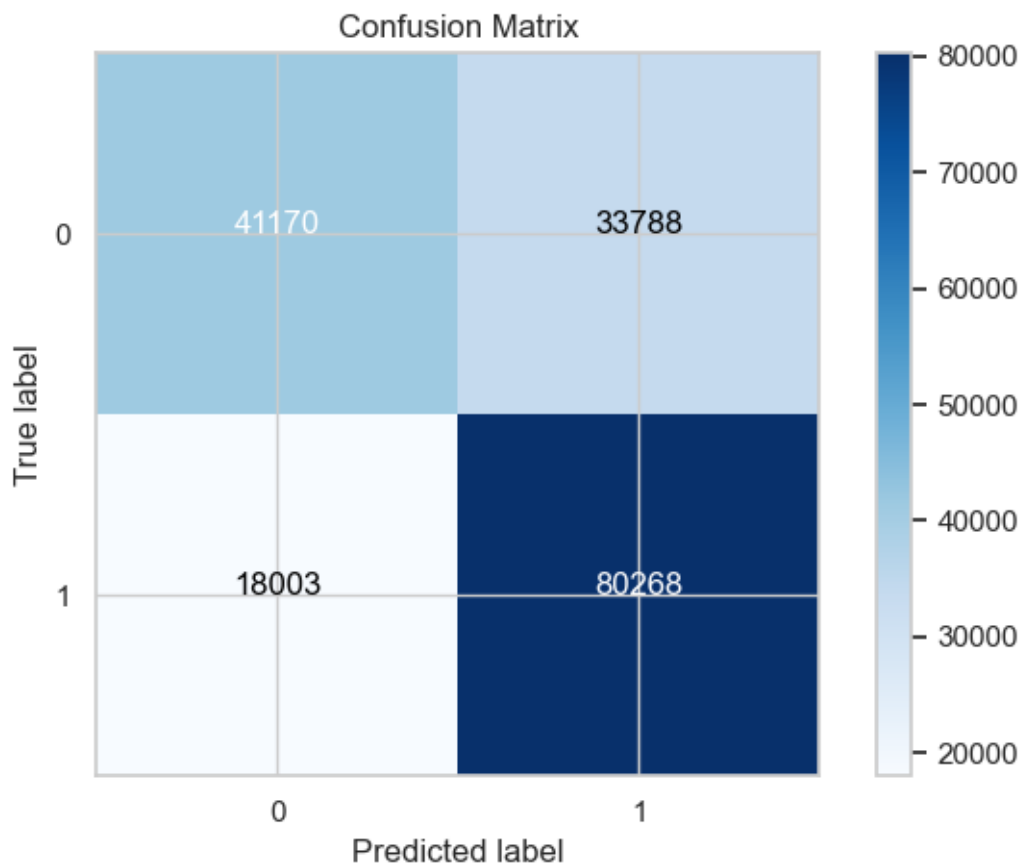
# Define the class labels
class_labels = ['0', '1']

# Plot the confusion matrix
plt.imshow(confusion_matrix_values, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion Matrix')
plt.colorbar()
tick_marks = np.arange(len(class_labels))
plt.xticks(tick_marks, class_labels)
plt.yticks(tick_marks, class_labels)

# Add labels to each cell
thresh = confusion_matrix_values.max() / 2.
for i in range(confusion_matrix_values.shape[0]):
    for j in range(confusion_matrix_values.shape[1]):
        plt.text(j, i, format(confusion_matrix_values[i, j], 'd'),
                  horizontalalignment="center",
                  color="white" if confusion_matrix_values[i, j] > thresh else
↪ "black")

plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.tight_layout()
plt.show()

```



Training Performance: The model achieved an accuracy of 0.70 on the training set, correctly classifying 70% of the instances. The precision and recall for both class 0 and class 1 have improved compared to the previous model, indicating a better overall performance on the training set. The weighted average F1-score is 0.68, suggesting a moderate to good performance on the training set.

Testing Performance: On the testing set, the model achieved an accuracy of 0.69, correctly classifying 69% of the instances. The precision and recall for both class 0 and class 1 have also improved compared to the previous model, indicating a better overall performance on the testing set. The weighted average F1-score is 0.68, indicating a moderate to good performance on the testing set.

Considering these observations, we can conclude that the updated model shows improvements in performance compared to the previous one. It demonstrates better accuracy, precision, recall, and F1-score on both the training and testing sets. The model's performance is more balanced between the two classes, with improved precision and recall for both class 0 and class 1.

```
[130]: def print_model_results(models, accuracy, precision_score, f1_score_0,
    ↪ f1_score_1, AUC):
    print("-----")
    print("MODEL\t\tACCURACY\tprecision_score,\tF1-SCORE 0 TESTING\tF1-SCORE 1\t
    ↪ TESTING\tAUC")
```


Therefore, based on the evaluation metrics, the XG Boost model is considered the best model for the given task.

Conclusion

- Downtown Chicago has a high concentration of accidents, primarily caused by intentional actions or driver errors. However, there are scattered incidents of unintentional accidents, indicating the need for safety improvements.
- Control failures in unintentional accidents are not significantly influenced by vision or speed. Other factors may contribute to these accidents and require further investigation.
- The absence of traffic control devices is a significant contributing factor to accidents in Chicago. Increasing their presence can help reduce unintentional accidents.
- Weather and lighting conditions have minimal impact on accident occurrence in Chicago.
- Accidents are common on non-divided roads, suggesting the need for road division measures to improve traffic management and safety.
- Road surface condition and defects have a minimal impact on unintentional accidents.
- Rush hour traffic, particularly between 14-18 hours, contributes to a higher number of accidents in the downtown area. Better traffic management strategies are needed during these peak hours.
- Weekend days show a slightly higher number of accidents compared to weekdays, but crash hour plays a more significant role in determining accident occurrence.
- Summer months have a higher number of accidents, potentially due to increased travel and outdoor activities. However, adjustments in road safety strategies based solely on the crash month may not be necessary.

Recommendations Increase Traffic Control Measures: Install additional traffic control devices, such as traffic lights, stop signs, and speed limit signs, particularly in areas with a high concentration of accidents. Ensure that existing devices are well-maintained and functioning properly.

Enhance Road Infrastructure: Implement road division measures, such as adding medians or physical barriers, to separate opposing flows of traffic and reduce the likelihood of collisions. Improve road surfaces to minimize hazards like potholes or uneven pavement.

Improve Traffic Management: Implement intelligent transportation systems and optimize traffic signal timings to facilitate traffic flow and reduce congestion, especially during peak rush hour periods. Consider deploying additional traffic management personnel to ensure efficient traffic management.

Driver Education and Awareness: Conduct targeted educational campaigns to raise awareness about safe driving practices, including the importance of attentiveness, obeying traffic laws, and maintaining a safe speed. Emphasize the risks associated with intentional actions, such as reckless driving or aggressive behavior.

Collaborate with Law Enforcement: Strengthen collaboration between the City of Chicago Vehicle Safety Board, law enforcement agencies, and other relevant stakeholders to enforce traffic laws effectively and deter dangerous driving behaviors.

Continuous Monitoring and Evaluation: Establish a robust system to collect and analyze data on car accidents continuously. Regularly evaluate the effectiveness of implemented measures and adjust strategies based on evolving trends and patterns in accidents.